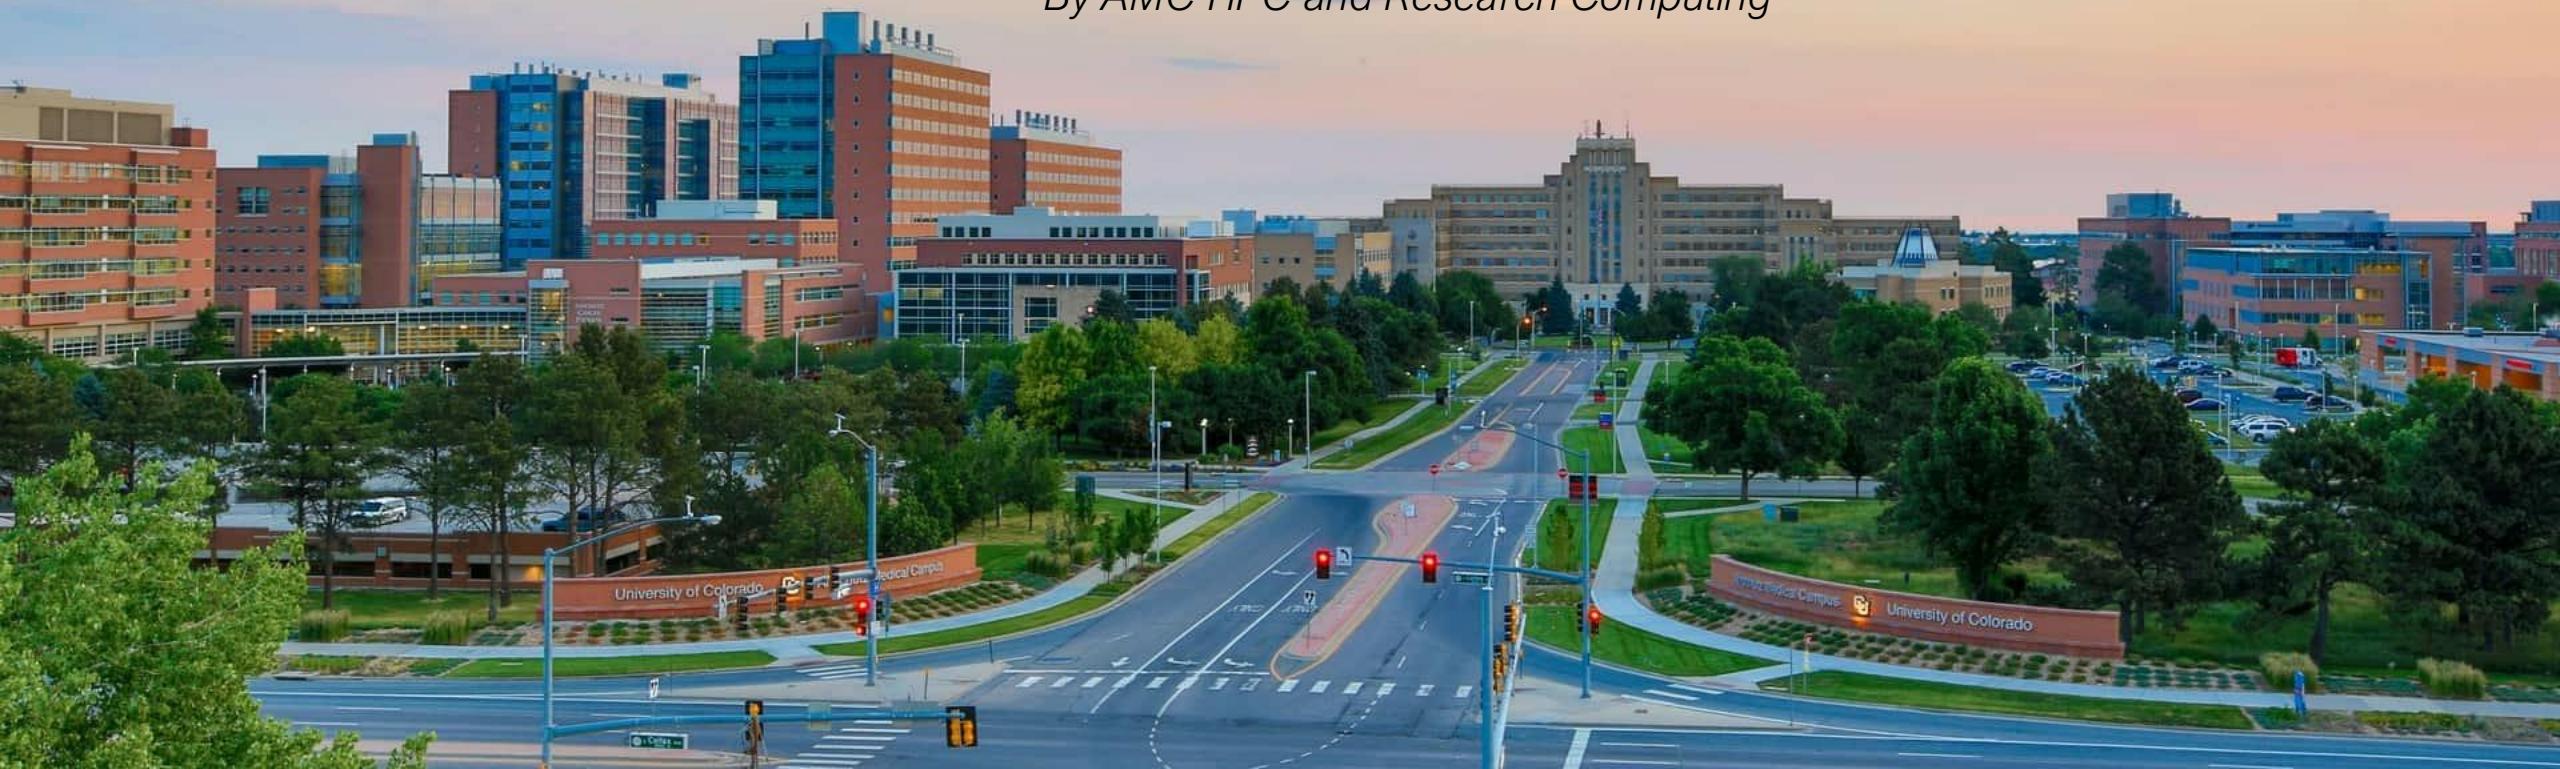




University of Colorado **Anschutz Medical Campus**

Alpine Best Practices Workshop

By AMC HPC and Research Computing



If you do not currently have an active Alpine account you can request one for free!

<https://cuanschutz.edu/hpc>

High Performance Computing (HPC)

Category: Computer and Device Support, High Performance Computing, Productivity and Business

Audience: Faculty, Researchers, Staff, Students

OIT and the School of Medicine (SOM) are partnering with CU Boulder Research Computing (CURC) to provide compute resources for researchers on the CU Anschutz campus.



Request Access

- You must have a valid cuanschutz.edu email address (we do not take ucdenver.edu email addresses; please request a cuanschutz.edu from your departmental administration team)
- You must have your ACCESS/XSEDE ID. If you do not have one, you can create one for free at:
<https://identity.access-ci.org/new-user>



University of Colorado
Anschutz Medical Campus

Reminder! Alpine and PetaLibrary cannot be used for highly sensitive data sets. Not sure about your data set? Please ask us!



- No HIPAA or Protected Health Information (PHI)
- No General Data Protection Regulation (GDPR)
 - No CMMC
 - No FERPA
- No data that requires special security and compliance



FERPA

Family Educational
Rights & Privacy Act



EUROPEAN GENOME-PHENOME ARCHIVE



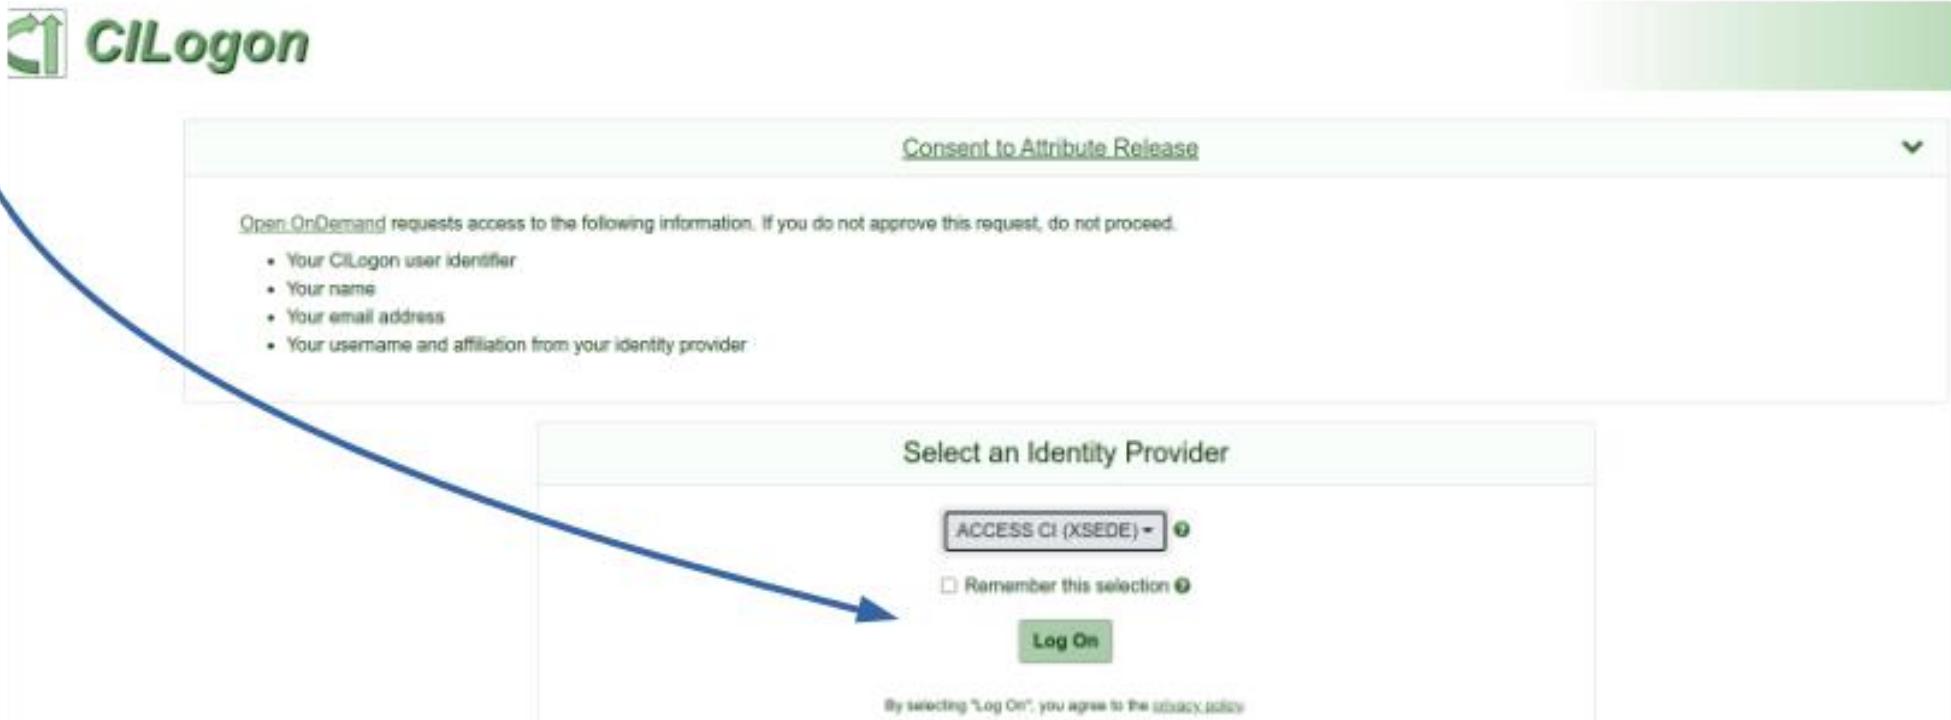
University of Colorado
Anschutz Medical Campus

All data on Alpine must be 100% de-identified and please make sure it is compliant with your IRB, DUA, data repository requirements, and/or grant funding agency

A refresher on how to log into Alpine

1 Go to <https://ondemand-rmacc.rc.colorado.edu>

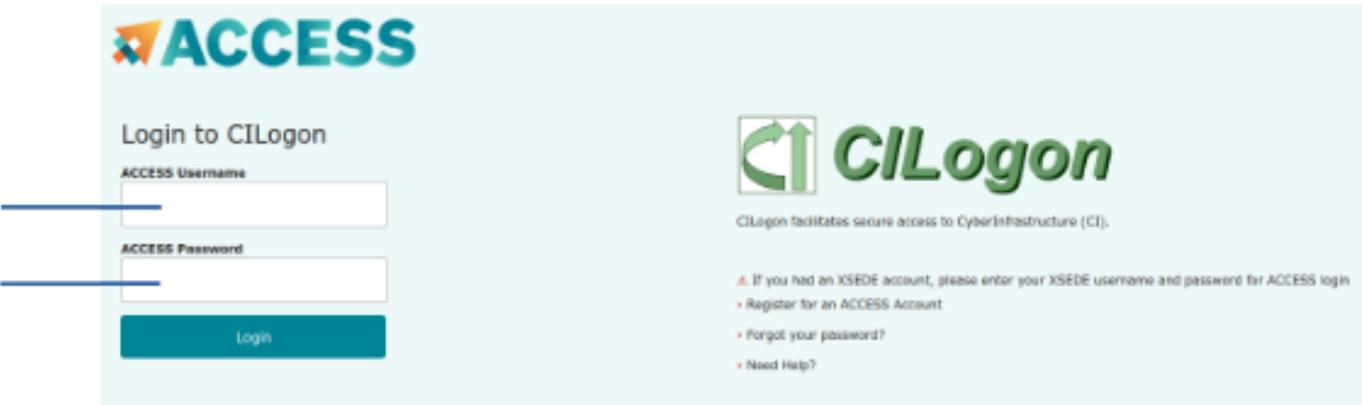
2 It should redirect you to CILogon which is how you authenticate your Alpine session. Make sure you select "ACCESS CI (XSEDE)" as your identity provider and then press "Log On"



A refresher on how to log into Alpine

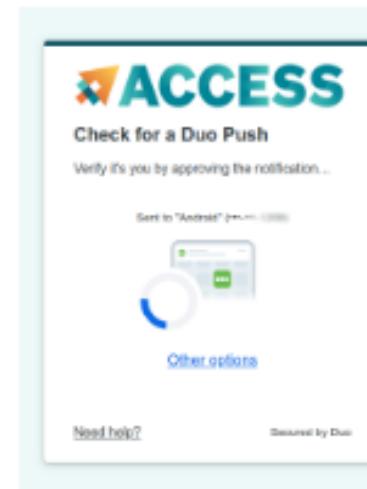
- 3 Next it will take you to this page where you will put in your ACCESS ID and ACCESS password and press Login. This is NOT your CU Anschutz ID!!

ACCESS ID and
ACCESS
password not CU
Anschutz
credentials!!



The screenshot shows the ACCESS login interface on the CILogon platform. It features a large 'ACCESS' logo at the top left. Below it is a 'Login to CILogon' form with two input fields: 'ACCESS Username' and 'ACCESS Password', each with its own placeholder text ('Enter your ACCESS ID here' and 'Enter your ACCESS password here'). A teal 'Login' button is positioned below the password field. To the right of the form is the 'CILogon' logo, which consists of a stylized green 'C' and 'I' icon followed by the word 'CILogon' in a bold, green, sans-serif font. Below the logo, the text 'CILogon facilitates secure access to CyberInfrastructure (CI)' is displayed.

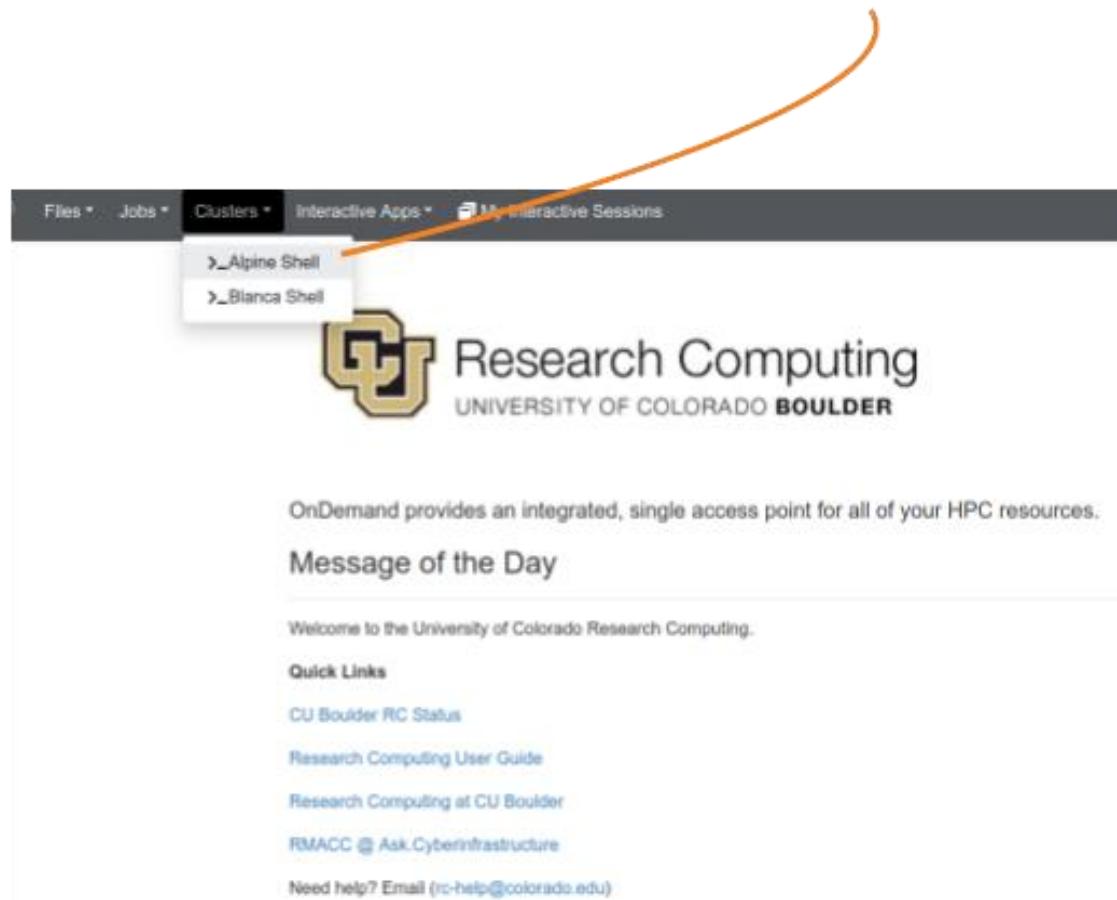
- 4 This will prompt a DUO MFA push to your phone or whichever way you have DUO set up on your phone to authenticate. Accept the push sent to your device.



A refresher on how to log into Alpine

5

This will take to the official CURC page. Let's select the Alpine shell.



6

This will log you into the head node of Alpine. It will **always default to your home directory**.



University of Colorado
Anschutz Medical Campus

Users can apply for ssh access to Alpine and we will notify applicants to attend one of our monthly meetings to configure your account

The benefits of having ssh access:

- Logging into Alpine can be done on the terminal via ssh; no need for the Open onDemand portal
- Access to transfer protocols such as rsync, scp, and sftp are available to you and therefore you can bypass globus and any globus related issues



University of Colorado
Anschutz Medical Campus

Anschutz ssh-beta testing request Agreement

Disclaimer

At the moment, there is no set up for users to automatically enroll themselves for ssh access for Alpine. This is partly because the enrollment process is still in beta testing due to a few bugs we are fixing on the back end.

However, we are ready to slowly on-board users who have requested this feature, especially for those of who need this to make their workflows more efficient. Please carefully read and follow the written instructions below as they are very time sensitive.

Once we send you an invite to register for ssh (email title will be "Invitation to join RMACC") , you must accept the invite within 7 days, otherwise your invite expires and we will have to restart the process all over. Before we send invites, make sure to enter your first name, last name and email in order to facilitate the logistical process. Finally, you will have to confirm and sign the agreement at the bottom of this page. This will indicate that you are ready for the invite and that we can send that out to you.



University of Colorado
Anschutz Medical Campus

https://ucdenverdata.formstack.com/forms/alpine_ssh_request_form

If you do not have ssh access there are many other methods you can use to transfer data

- Globus for large data transfers

<https://curc.readthedocs.io/en/latest/compute/data-transfer.html#data-transfer>

- Open OnDemand for small data transfers

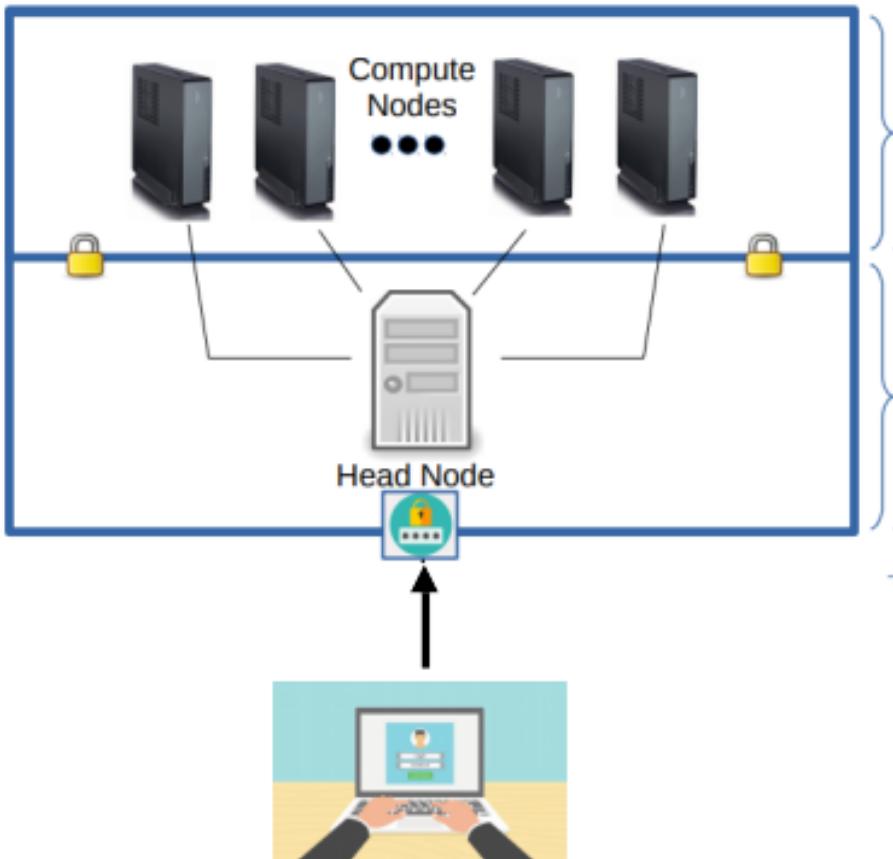
- Rclone on OneDrive

<https://curc.readthedocs.io/en/latest/storage/petalibrary/onedrive.html#using-rclone>

Please keep in mind, sometimes our networks become overwhelmed on globus. In that case using ssh or Rclone to transfer data will help alleviate these bottlenecks so you can continue to transfer data regardless if the globus endpoint is being bogged down.



Head node vs Compute Node: never run software, installs, or analysis on a head node!



The compute nodes are where all the power of the HPC is harnessed. This is where all the compute nodes are located. On Alpine, this is where your CPU nodes (amilan, amem, ahub, acompile) and GPU nodes are located (aa100, ami100, al40, atesting_a100, atesting_mi100)

Sometimes referred to as the login node. Its main purpose is to take your job requests and find compute resources for you to use.

This node is the master or brain of the HPC.

Compute Node

Any really long (> few minutes) and/or large (>500MB of memory) programs or calls needs to be run on compute nodes.

The compute nodes wait for the head node to assign them a job or task.

Software compiling and installations

Head/Login Node

The head node is very small and limited in ability. Its main function is to work with the scheduler and delegate tasks and jobs to the compute nodes that users want to utilize

Strictly used to navigate the Alpine file system with basic linux commands or submit jobs to the scheduler using sbatch

Rule of Thumb

Always request a compute node for running any scripts or software related tasks, *including compiling and installing local software*



How can I tell if I am on the head node or a compute node?

```
[tobr0001_amc@login-ci2 brunetti@xsede.org]$ ]
```

You can:

- submit jobs using sbatch,
- use basic Linux commands (`mv`, `cd`, `rm`, `ls`, etc...),
- use slurm commands to check job status and efficiency (`seff`, `seff-array`, `jobstats`, `sacct`, `sstat`)
- Use slurm commands to request interactive use of compute resources (`acompile`, `sinteractive`)
- Use text editors (`nano`, `emacs`, `vi/vim`) that are very low memory (i.e. do not open large (>500 Mb) files)

Anything that says `login-` tells you are on the login/head node

Please do not:

- run software
- execute code or scripts
- install custom software

You are on a compute node if...

```
[tobr0001_amc@c3gpu-c2-u29 brunetti@xsede.org]$ ]
```

Anything that says `gpu-` tells you are on the inside of a GPU node in an interactive session

It is safe to run any software, installs, programs, code here.

```
[tobr0001_amc@c3cpu-a2-u32-2 brunetti@xsede.org]$ ]
```

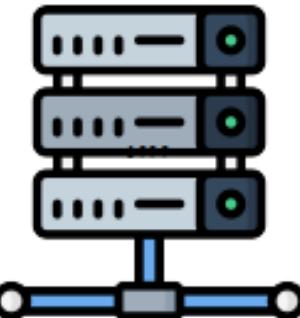
Anything that says `cpu-` tells you are on the inside of a CPU node in an interactive session

It is safe to run any software, installs, programs, code here.



Job efficiency matters – it costs the campus available service units

Although, you as users, do not have a personal financial cost associated for using Alpine (with the exception of PetaLibrary), there is a cost to reserving resources. The currency we use are called **service units (SUs)**.



Compute requested*



SUs charged

Every time you request a set amount of compute resources, that costs us a portion of SUs we have available to the Anschutz Alpine community. We have a finite amount of SUs available.

IMPORTANT NOTE

We are charged, from our finite pool of SUs, **based on the resource reservation requested, regardless if you use those resources efficiently or not**. Therefore, we want to make sure that when you request large amounts of resources you are using those resources efficiently so that we don't drain our pool of community SUs unnecessarily, resulting in everyone's jobs having longer wait times.



The SU CPU currency has the following **approximate** conversion rate

As of March 2025, we have 5,195 total SUs available per hour for our entire campus to use which equates to roughly 124,680 SUs available in a given 24 hour period.

CPU nodes	cost
amilan CPU nodes (partition = amilan)	1 SU/core/hour
High memory CPU nodes (partition = amem)	4 SUs/core/hour

Example calculation

One typical amilan CPU node can have up to 64 cores, so if you request a single amilan CPU node and all 64 cores and your job runs for a full 24 hours we are charged:

$$1 \text{ amilan node} \times 64 \text{ cores} \times 1 \text{ SU per core hour} \times 24 \text{ hours} = 1,536 \text{ SUs charged for your job}$$



On Alpine, `seff` allow you to check the efficiency of CPU based jobs once they are complete

```
[brunetti@xsede.org@login-ci2 brunetti@xsede.org]$ seff 10873469
```

Your job ID

```
Job ID: 10873469
Cluster: alpine
User/Group: tobr0001_amc/brunettipgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 27
CPU Utilized: 22:54:40
CPU Efficiency: 5.89% of 16-05:14:06 core-walltime
Job Wall-clock time: 14:24:58
Memory Utilized: 73.04 GB
Memory Efficiency: 73.04% of 100.00 GB
```

Goal: ensure either CPU efficiency **or** memory efficiency is at least 70% or better



University of Colorado
Anschutz Medical Campus

What did job ID 10873469 cost us in SUs?

```
Job ID: 10873469
Cluster: alpine
User/Group: tobr0001_amc/brunettipgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 27
CPU Utilized: 22:54:40
CPU Efficiency: 5.89% of 16-05:14:06 core-walltime
Job Wall-clock time: 14:24:58
Memory Utilized: 73.04 GB
Memory Efficiency: 73.04% of 100.00 GB
```

We know we used 27 cores

We know we used it for about
14 hours and 24 minutes

However, what we don't know is whether this was an amilan node or an amem high memory node. Luckily, we can use `sacct` to grab the script associated with this job ID.

```
[brunetti@xsede.org@login-ci2 brunetti@xsede.org]$ sacct -B -j 10873469
```

Print the batch script
associated with the
listed jobID

Parameter to tell
SLURM the following
value is a jobID

Your job ID



University of Colorado
Anschutz Medical Campus

What did job ID 10873469 cost us in SUs?

I requested amilan nodes, so now we can calculate the SUs.

```
Batch Script for 10873469
-----
#!/bin/bash

#SBATCH --nodes=1 # use one node
#SBATCH --time=5-00:00:00 # 5 days
#SBATCH --account=amc-general # normal, amc, long, mem (use mem when using the amem partition)
#SBATCH --partition=amilan # amilan, ami100, aa100, amem, amc
#SBATCH --qos=long
#SBATCH --ntasks=25 # total processes/threads
#SBATCH --job-name=cellranger
#SBATCH --output=step1_cellranger_LI047_gex_out_01062025_%J.log
#SBATCH --mem=100G # suffix K,M,G,T can be used with default to M
#SBATCH --mail-user=tonya.brunetti@cuanschutz.edu
#SBATCH --mail-type=END
#SBATCH --error=step1_cellranger_LI047_gex_out_01062025_%J.err

cellranger="/projects/$USER/software/cellranger-7.0.1/cellranger"
id="step1_gex_output"
inputcsv="/scratch/alpine/tobr0001_amc/[REDACTED]/cite_vdjB_hto_LI047_noLIBRA_01022024/step1_input_csv_cellranger_multi.csv"

${cellranger} multi --id=${id} --csv=${inputcsv} --disable-ui
```

$$1 \text{ amilan node} \times [27 \text{ cores}] \times 1 \text{ SU per core hour} \times [14.5 \text{ hours}] = \sim 391.5 \text{ SUs charged for the job}$$

Notice the cores come from the `seff` command, not necessarily what we requested which was 25 – why?

We are only charged for the time we use for the job to complete. Notice even though I requested a 5 day run, the job finished in 14.5 hours so we are charged for 14.5 hours



Our core architecture is tied directly to shared RAM, so SLURM takes the ceiling of the requested resources

<https://curc.readthedocs.io/en/latest/clusters/alpine/alpine-hardware.html#partitions>

Partition	Description	# of nodes	cores/node	RAM/core (GB)	Billing_weight/core	Default/Max Walltime	Resource Limits
amilan	AMD Milan (default)	386	32 or 48 or 64	3.75	1	24H, 7D	see qos table
ami100	GPU-enabled (3x AMD MI100)	8	64	3.75	6.1 ³	24H, 7D	15 GPUs across all jobs
aa100	GPU-enabled (3x NVIDIA A100) ⁴	12	64	3.75	6.1 ³	24H, 7D	21 GPUs across all jobs
al40	GPU-enabled (3x NVIDIA L40) ⁴	3	64	3.75	6.1 ³	24H, 7D	6 GPUs across all jobs
amem ¹	High-memory	24	48 or 64 or 128	16 ²	4.0	4H, 7D	128 cores across all jobs

I requested 25 cores and 100 GB of RAM in my script, however, each core on an amilan node has 3.75 GB of RAM associated with it.

Calculation #1: 25 cores requested x 3.75 GB RAM = 93.75GB of RAM available,
but I requested 100GB, so SLURM adds the minimum number of cores to reach 100GB of RAM.

Calculation #2: 100GB RAM/3.75 = ~26.667, so if SLURM does me a favor and bumps me up to 27 cores to achieve my max memory. Note, that SLURM will kill my job however, if I exceed 100G of max RAM.



Let's take a look at the efficiency of job ID 4381567

```
[brunetti@xsede.org@login-ci2 brunetti@xsede.org]$ seff 4381567
```

Not good!
CPU efficiency **or**
memory efficiency is
at not at least 70%
or better.

**We need to
optimize resources**
before submitting
this job again in the
future

```
Job ID: 4381567
Cluster: alpine
User/Group: tobr0001_amc/brunettipgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 14
CPU Utilized: 3-16:02:39
CPU Efficiency: 14.46% of 25-08:44:22 core-walltime
Job Wall-clock time: 1-19:28:53
Memory Utilized: 87.38 MB
Memory Efficiency: 0.17% of 50.00 GB
```

used 14 cores

Used resources
for 1 day, 19
hours, and 28
minutes

Let's check the script and see which nodes partitions I requested and what resources I requested

```
[brunetti@xsede.org@login-ci2 brunetti@xsede.org]$ sacct -B -j 4381567
```



University of Colorado
Anschutz Medical Campus

Let's take a look at the efficiency of job ID 4381567

I requested amilan nodes, so now we can calculate the SUs.

```
Batch Script for 4381567
-----
#!/bin/bash
#SBATCH --nodes=1 # use one node
#SBATCH --time=72:00:00
#SBATCH --account=amc-general # normal, amc, long, mem (use mem when using the amem partition)
#SBATCH --partition=amilan # amilan, ami100, aa100, amem, amc
#SBATCH --qos=long
#SBATCH --ntasks=5 # total processes/threads
#SBATCH --job-name=demuxTB
#SBATCH --mem=50G # suffix K,M,G,T can be used with default to M
#SBATCH --mail-user=tonya.brunetti@cuanschutz.edu
#SBATCH --mail-type=END
#SBATCH --output=ALAW-AS049-multiplexGroup_1-2_%J.log
#SBATCH --error=ALAW-AS049-multiplexGroup_1-2_%J.err
```

1 amilan node x [14 cores] x 1 SU per core hour x [43.5 hours] = ~609 SUs charged for the job

SLURM charged me for 14 cores instead of the 5 I requested because the memory I requested could not be achieved with just 5 cores:

Calculation #1: 5 cores requested x 3.75 GB RAM = 18.75 GB of RAM available,
but I requested 50GB, so SLURM adds the minimum number of cores to reach 50GB of RAM.

Calculation #2: 50GB RAM/3.75 = ~13.333, so it SLURM does me a favor and bumps me up to 14 cores to achieve my max memory. Note, that
SLURM will kill my job however, if I exceed 50GB of max RAM.



University of Colorado
Anschutz Medical Campus

How can I make better use of resources and use less SUs to essentially complete the same job while reducing resource waste?

```
Job ID: 4381567
Cluster: alpine
User/Group: tobr0001_amc/brunettipgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 14
CPU Utilized: 3-16:02:39
CPU Efficiency: 14.46% of 25-08:44:22 core-walltime
Job Wall-clock time: 1-19:28:53
Memory Utilized: 87.38 MB
Memory Efficiency: 0.17% of 50.00 GB
```



```
#SBATCH --ntasks=4
#SBATCH --mem =1G
```

A better optimized use of resources that likely **reduces the cost of SUs consumed by over 50%!**



Here, I only used 14.46% of CPUs efficiently, so technically I could have dropped this to 3 CPUs:

$14 \text{ cores} * .1446 = 2.02 \text{ cores}$; however let's pad it a bit to account for inefficiencies to either 3 or 4 cores

$50 \text{ GB} * 0.0017 = \sim 0.085 \text{ GB used}$; well 1 core has 3.75 GB so we would have been fine on using just a single core.



University of Colorado
Anschutz Medical Campus

Arrays should never be used without running a representative small batch of samples first, in order to optimize efficiency of resources

Arrays are an easy way to parallelize upwards of thousands of jobs in a script; however, you can imagine that when slurm submits a 1000 jobs, if the job is already incredibly inefficient, then that now gets compounded over the number of jobs you submitted.

For job arrays, you can get a summarized view of all jobs in that array using:

```
seff-array yourJobID
```



Arrays should never be used without running a representative small batch of samples first, in order to optimize efficiency of resources

seff-array yourJobID

In this example, this array has 900 jobs associated with it and by running the above command on this jobID, you can see a summary of efficiency across all 900 jobs. Just make sure all jobs have completed! Any actively running jobs will not provide an accurate estimation of efficiency.

Goal: ensure either CPU efficiency or memory efficiency is at least 70% or better across at least 70-80% of all jobs submitted under that job array

```
Requested CPUs: 12 cores on 1 node(s)
Requested Memory: 45G
Requested Time: 07:00:00
-----
Job Status
COMPLETED: 900
-----
Finished Job Statistics
(excludes pending, running, and cancelled jobs)
Average CPU Efficiency 11.83%
Average Memory Usage 8.02G
Average Run-time 1146.37s
```

Here you can see what was requested for each job in the array: 12 CPU cores with a max memory of 45G. There were 900 jobs in this array therefore, **this request was made 900 separate times**

Here you can see that on average across all 900 jobs, **11.83% of CPUs were efficiently used**, and $8.02G/45G$ requested = **17.82% of available requested memory was efficiently used**. **Neither one of these is at least 70%**



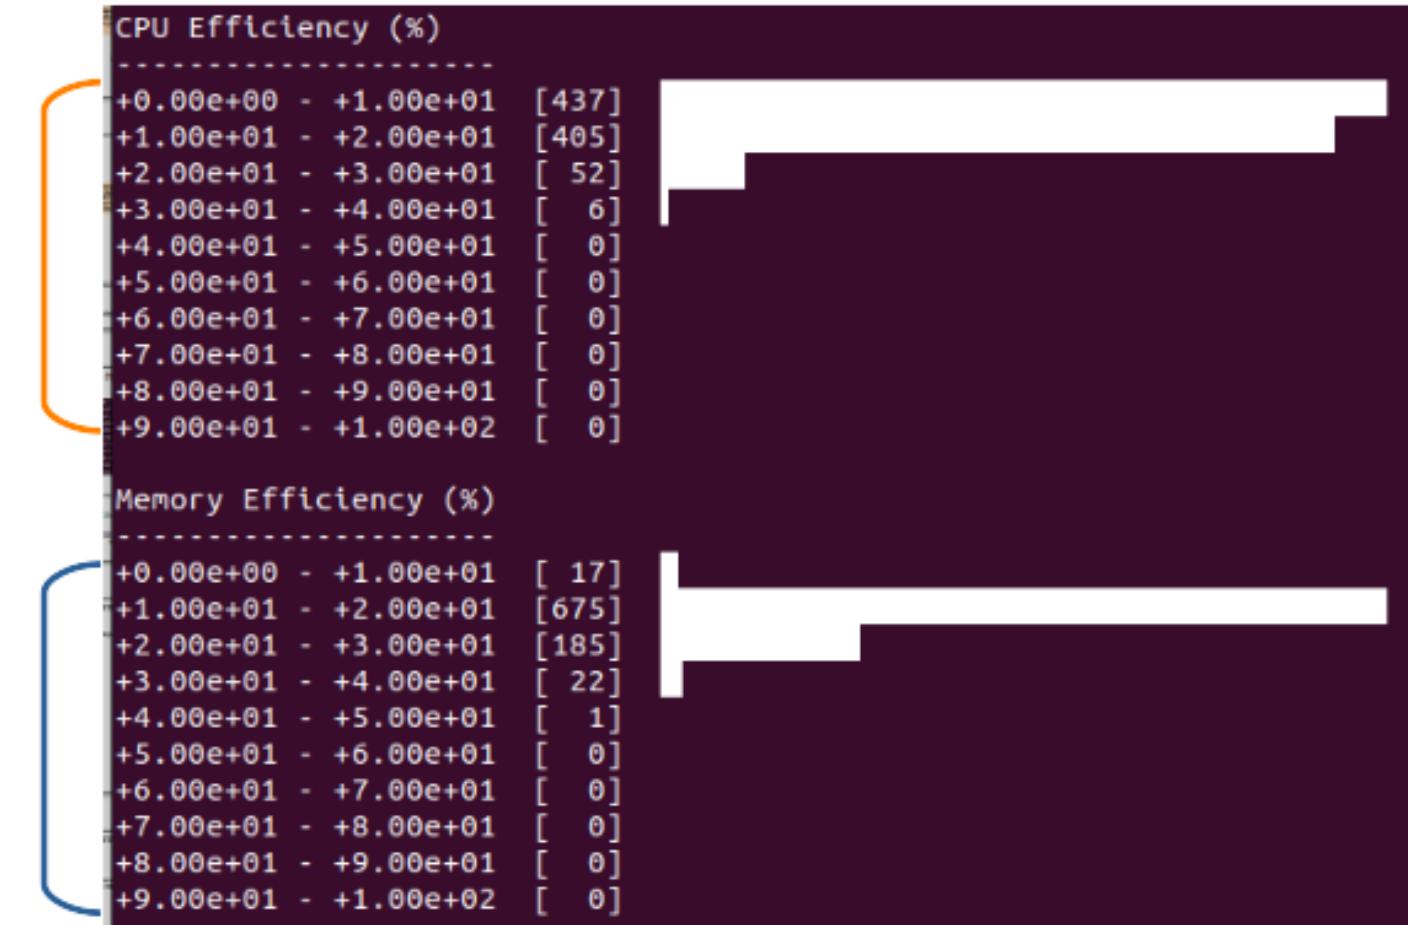
`seff-array` also provides a more granular detailed visual breakdown of CPU and memory efficiency across all completed jobs in your array

Not a single job had a CPU efficiency greater than 30-40%!

$842/900 = 93.56\%$ of jobs never exceed 20% CPU efficiency

Not a single job had a memory efficiency greater than 40-50%!

$877/900 = 97.44\%$ of jobs never exceeding 30% memory efficiency



The problem here is we are charged for the amount we reserved even though we did not need that amount

This was on an amilan node so the charge is 1 SU/1 core hour with 45G of memory with the average run time being ~20 minutes. This means we were charged roughly the following SUs per job:

$$1 \text{ amilan node} \times 12 \text{ cores} \times 1 \text{ SU per core hour} \times 0.3 \text{ hours} = \sim 3.6 \text{ SUs charged for the job}$$

However, we requested this 900 times:

$$\sim 3.6 \text{ SUs per job} \times 900 \text{ jobs} = 3,240 \text{ SUs charged}$$

Based on our `seff-array` results we know that > 90% of the jobs could have been run within with 1/5 the requested CPUs and ~1/3 the amount of memory requested.

What could we have done to make our SUs go farther and our jobs are more efficient?



We should have scaled down our resources before submitting 900 jobs in an array

Not a single job had a CPU efficiency greater than 30-40%!

$842/900 = 93.56\%$ of jobs never exceed 20% CPU efficiency



Requested 12 cores but 20% of this would have covered 93.56% perfectly fine:

$$12 \text{ cores} * .20 \text{ efficiency} = 2.4 \text{ cores}$$

The nearest ceiling core is 3.



Take the max of these 2 values

Original Version (not correct)

Charged ~3,240 SUs for 900 jobs inefficiently used

```
#SBATCH --ntasks=12  
#SBATCH --mem=45G
```



Not a single job had a memory efficiency greater than 40-50%!

$877/900 = 97.44\%$ of jobs never exceed 30% memory efficiency



Requested 45 GB of RAM but 30% of this would have covered 97.44% perfectly fine:

$$45\text{GB} * .30 \text{ efficiency} = 13.5 \text{ GB of RAM}$$

$13.5\text{GB}/3.75\text{GB per core} = 3.6$ cores needed to achieve this RAM.

The nearest ceiling core is 4.



Corrected Version

```
#SBATCH --ntasks=4  
#SBATCH --mem=15G
```

$4 \text{ cores} \times 3.75\text{GB per core} = 15\text{G RAM}$

Charged ~1,080 SUs for the same 900 jobs efficiently used



University of Colorado
Anschutz Medical Campus

Best Practices for Using Arrays on Alpine

1) ALWAYS run a small subset of samples in an array first before scaling to your full dataset.

Had the previously job just tested 20-40 samples first, the number of wasted resources would have been drastically reduced.

2) After your test run above, get efficiency statistics using `seff-array` and adjust your resource SBATCH directives accordingly.

3) Request resources based on what covers 70-80% of use cases, not the max! It is more efficient to run jobs mostly efficient and re-run the small number of job failures separately at a different resource amount.



Document your error and log files with unique Job IDs

The benefits of knowing what job ID are tied to error and log outputs is incredibly useful and makes troubleshooting malfunctioning code much easier

```
Batch Script for 10873469
-----
#!/bin/bash

#SBATCH --nodes=1 # use one node
#SBATCH --time=5-00:00:00 #5 days
#SBATCH --account=amc-general # normal, amc, long, mem (use mem when using the amem partition)
#SBATCH --partition=amilan # amilan, ami100, aa100, amem, amc
#SBATCH --qos=long
#SBATCH --ntasks=25 # total processes/threads
#SBATCH --job-name=cellranger
#SBATCH --output=step1_cellranger_LI047_gex_out_01062025_%J.log
#SBATCH --mem=100G # suffix K,M,G,T can be used with default to M
#SBATCH --mail-user=tonya.brunetti@cuanschutz.edu
#SBATCH --mail-type=END
#SBATCH --error=step1_cellranger_LI047_gex_out_01062025_%J.err
```

`%J` will automatically substitute the job ID into the error and log outputs.

This means you can always grab the script submitted with this job by using,

```
sacct -B -j yourJobID
```

as we have seen before, and that can help our team troubleshoot where the problems in your script may be if you reach out to us when a job fails



University of Colorado
Anschutz Medical Campus

“How do I know what resources my job needs?”

```
#!/bin/bash

#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --partition=amilan
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
#  #SBATCH --mem=10G

## module load

echo "Job started on $(hostname) at $(date)"

# Say we want to sequence 1,000,000,000 numbers, but can start small with 1,000,000 first to see what that costs!
# Simple computation: simulate CPU load with a loop
echo "Running a simple computation..."
for i in $(seq 1 1000000); do
    result=$((result + i))
done

# Print the result
echo "Computation finished. Result: $result"

echo "Job completed at $(date)"
```



“How do I know what resources my job needs?”

```
[msherren@xsede.org@login-ci1 best_practices_scripts]$ seff 12535812
Job ID: 12535812
Cluster: alpine
User/Group: msherren@xsede.org/msherrenpgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 4
CPU Utilized: 00:00:03
CPU Efficiency: 8.33% of 00:00:36 core-walltime
Job Wall-clock time: 00:00:09
Memory Utilized: 308.00 KB
Memory Efficiency: 0.00% of 15.00 GB
```

“How do I know what resources my job needs?”

```
[msherren@xsede.org@login-cil best_practices_scripts]$ seff 12535812
Job ID: 12535812
Cluster: alpine
User/Group: msherren@xsede.org/msherrenpgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 4
CPU Utilized: 00:00:03
CPU Efficiency: 8.33% of 00:00:36 core-walltime
Job Wall-clock time: 00:00:09
Memory Utilized: 308.00 KB
Memory Efficiency: 0.00% of 15.00 GB
```

“How do I know what resources my job needs?”

```
[msherren@xsede.org@login-cil best_practices_scripts]$ seff 12535812
Job ID: 12535812
Cluster: alpine
User/Group: msherren@xsede.org/msherrenpgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 4
CPU Utilized: 00:00:03
CPU Efficiency: 8.33% of 00:00:36 core-walltime
Job Wall-clock time: 00:00:09
Memory Utilized: 308.00 KB
Memory Efficiency: 0.00% of 15.00 GB
```



- From seff we learn:
- CPU usage and efficiency
 - Memory usage and efficiency
 - Wall time

Diagnosing a job that doesn't finish correctly

```
1  #!/bin/bash
2
3
4  #SBATCH --job-name=job_scaling
5  #SBATCH --output=dump.%j.out
6  #SBATCH --account=amc-general
7  #SBATCH --nodes=1
8  #SBATCH --partition=amilan
9  #SBATCH --time=00:05:00
10 #SBATCH --ntasks=4
11 #  #SBATCH --mem=10G
12
13
14 ## module load
15
16 echo "Job started on $(hostname) at $(date)"
17
18 echo "Running a simple computation..."
19 # The loop below is intentionally designed to run indefinitely/until resources are exhausted
20 for i in ${seq 1 1000000000000}; do
21     arr=(${seq 1 1000}) # This array will take up memory with each iteration
22
23     # Simulate heavy computation by simply doing nothing but waste CPU
24     # (Note: this could be replaced with a non-useful computation)
25     result=$((result + i)) # We are updating an uninitialized variable `result`
26
27     # This loop will never end and will continue using CPU and memory.
28 done
29
30 # Print the result
31 echo "Computation finished. Result: $result"
32
33 [msherren@xsede.org@login-c11 best_practices_scripts]$ sbatch broken_for_loop_example.sh
Submitted batch job 12535838
[msherren@xsede.org@login-c11 best_practices_scripts]$ seff 12535838
Job ID: 12535838
Cluster: alpine
User/Group: msherren@xsede.org/msherrenpgrp@xsede.org
State: FAILED (exit code 2)
```



Diagnosing a job that doesn't finish correctly

With that Job ID you can:

- `seff <JobID>`
- Check the logs:
`dump.<JobID>, slurm-<JobID>.out,`
`error.<JobID>.out`

With module load
slurmtools:

- `sacct -j <JobID>`
- `sacct -u <username>`
- `jobstats`

See slides **40-56** For
more details about
these commands

- <https://curc.readthedocs.io/en/latest/compute/managing-resources.html>

```
1  #!/bin/bash
2
3
4  #SBATCH --job-name=job_scaling
5  #SBATCH --output=dump.%j.out
6  #SBATCH --account=amc-general
7  #SBATCH --nodes=1
8  #SBATCH --partition=amilan
9  #SBATCH --time=00:05:00
10 #SBATCH --ntasks=4
11 #SBATCH --mem=10G
12
13
14 ## module load
15
16 echo "Job started on $(hostname) at $(date)"
17
18 echo "Running a simple computation..."
19 # The loop below is intentionally designed to run indefinitely/until resources are exhausted
20 for i in ${seq 1 1000000000000}; do
21     arr=(${seq 1 1000}) # This array will take up memory with each iteration
22
23     # Simulate heavy computation by simply doing nothing but waste CPU
24     # (Note: this could be replaced with a non-useful computation)
25     result=$((result + i)) # We are updating an uninitialized variable `result`
26
27     # This loop will never end and will continue using CPU and memory.
28 done
29
30 # Print the result
31 echo "Computation finished. Result: $result"
32
33 echo "Job completed at $(date)"
```

```
[msherren@xsede.org@login-c11 best_practices_scripts]$ head dump.12535838.out
The job ID is: 12535838
Job started on c3cpu-c15-u17-1.rc.int.colorado.edu at Wed Mar 26 14:51:04 MDT 2025
Running a simple computation...
[msherren@xsede.org@login-c11 best_practices_scripts]$ head error.12535838.out
/var/spool/slurmd/job12535838/slurm_script: xrealloc: cannot allocate 18446744071562067968 bytes
```

CPU resource requirements for each QOS/Partition

#SBATCH --qos	Partitions	#SBATCH --time	#SBATCH --mem	RAM per core (GB)	Max cores per node
normal	amilan	Max: 24 hrs	Max: 240GB	3.75	32/48/64
long	amilan	Min: 24 hrs Max: 7 Days	Max: 240GB	3.75	32/48/64
mem	amem only	Max: 7 Days	Min: 240GB Max: 2TB	16	48/64/128



GPU resource requirements for each QOS/Partition

#SBATCH --qos	Partitions	#SBATCH --time	#SBATCH --mem	RAM per core (GB)	Max cores per node
normal	aa100,ami100, al40	Max: 24 hrs	Max: 240GB	3.75	32/48/64
long	aa100,ami100, al40	Min: 24 hrs Max: 7 Days	Max: 240GB	3.75	32/48/64

*To specify how many GPUs (up to 3 per node), remember to include the --gres flag:

#SBATCH --gres=gpu for a single GPU, #SBATCH --gres=gpu:2 for 2, and #SBATCH --gres=gpu:3 for 3

**Make sure your job is set up for multiple GPUs before requesting multiple!



Which of these 4 scripts will fail?

```
#!/bin/bash

#SBATCH --job-name=qos_batch_examples
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=normal
#SBATCH --partition=amilan
#SBATCH --mem=1000MB
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
```

```
#!/bin/bash

#SBATCH --job-name=qos_batch_examples
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=normal
#SBATCH --partition=amilan
#SBATCH --mem=10GB
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
```

```
#!/bin/bash

#SBATCH --job-name=qos_batch_examples
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=mem
#SBATCH --partition=amem
#SBATCH --mem=100GB
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
```

```
#!/bin/bash

#SBATCH --job-name=qos_batch_examples
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=mem
#SBATCH --partition=amem
#SBATCH --mem=500GB
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
```



Which of these 4 scripts will fail?

```
#!/bin/bash
```

```
#SBATCH --job-name=qos_batch_examples  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --qos=normal  
#SBATCH --partition=amilan  
#SBATCH --mem=1000MB  
#SBATCH --time=01:00:00  
#SBATCH --ntasks=4
```

```
#!/bin/bash
```

```
#SBATCH --job-name=qos_batch_examples  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --qos=mem  
#SBATCH --partition=amem  
#SBATCH --mem=100GB ←  
#SBATCH --time=01:00:00  
#SBATCH --ntasks=4
```

```
#!/bin/bash
```

```
#SBATCH --job-name=qos_batch_examples  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --qos=normal  
#SBATCH --partition=amilan  
#SBATCH --mem=10GB  
#SBATCH --time=01:00:00  
#SBATCH --ntasks=4
```

```
#!/bin/bash
```

```
#SBATCH --job-name=qos_batch_examples  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --qos=mem  
#SBATCH --partition=amem  
#SBATCH --mem=500GB  
#SBATCH --time=01:00:00  
#SBATCH --ntasks=4
```

This doesn't work because the partition has a minimum memory requirement of 240GB



University of Colorado
Anschutz Medical Campus

Which of these 4 scripts will fail?

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --partition=aa100
#SBATCH --time=02:00:00
#SBATCH --ntasks=50
#SBATCH --gres=gpu:2
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --time=4:00:00
#SBATCH --ntasks=8
#SBATCH --gres=gpu
#SBATCH --mem=20GB
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=long
#SBATCH --partition=al40
#SBATCH --time=7-00:00:00
#SBATCH --ntasks=16
#SBATCH --gres=gpu:3
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --qos=normal
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks=16
#SBATCH --time=48:00:00
#SBATCH --gres=gpu:3
```



Which of these 4 scripts will fail?

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --partition=aa100
#SBATCH --time=02:00:00
#SBATCH --ntasks=50
#SBATCH --gres=gpu:2
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --time=4:00:00
#SBATCH --ntasks=8
#SBATCH --gres=gpu
#SBATCH --mem=20GB
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --nodes=1
#SBATCH --qos=long
#SBATCH --partition=al40
#SBATCH --time=7-00:00:00
#SBATCH --ntasks=16
#SBATCH --gres=gpu:3
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling
#SBATCH --output=dump.%j.out
#SBATCH --account=amc-general
#SBATCH --qos=normal
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks=16
#SBATCH --time=48:00:00
#SBATCH --gres=gpu:3
```

This doesn't work
because the normal
qos has a maximum
time limit of 24 hours



University of Colorado
Anschutz Medical Campus

```
sbatch: error: Error: The normal QOS has a maximum wall time limit of 24 hours. Please adjust your requested time or choose an appropriate QOS.  
sbatch: error: Batch job submission failed: Unspecified error
```

Which of these 4 scripts will fail?

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --partition=aa100  
#SBATCH --time=02:00:00  
#SBATCH --ntasks=50  
#SBATCH --gres=gpu:2
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --qos=long  
#SBATCH --partition=al40  
#SBATCH --time=7-00:00:00  
#SBATCH --ntasks=16  
#SBATCH --gres=gpu:3
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --nodes=1  
#SBATCH --time=4:00:00  
#SBATCH --ntasks=8  
#SBATCH --gres=gpu  
#SBATCH --mem=20GB
```

```
#!/bin/bash
```

```
#SBATCH --job-name=job_scaling  
#SBATCH --output=dump.%j.out  
#SBATCH --account=amc-general  
#SBATCH --qos=long  
#SBATCH --partition=ami100  
#SBATCH --nodes=1  
#SBATCH --ntasks=16  
#SBATCH --time=20:00:00  
#SBATCH --gres=gpu:3
```

This also does not work! Requesting long with less than 24 hours will queue for longer and waste resources



University of Colorado
Anschutz Medical Campus

squeue

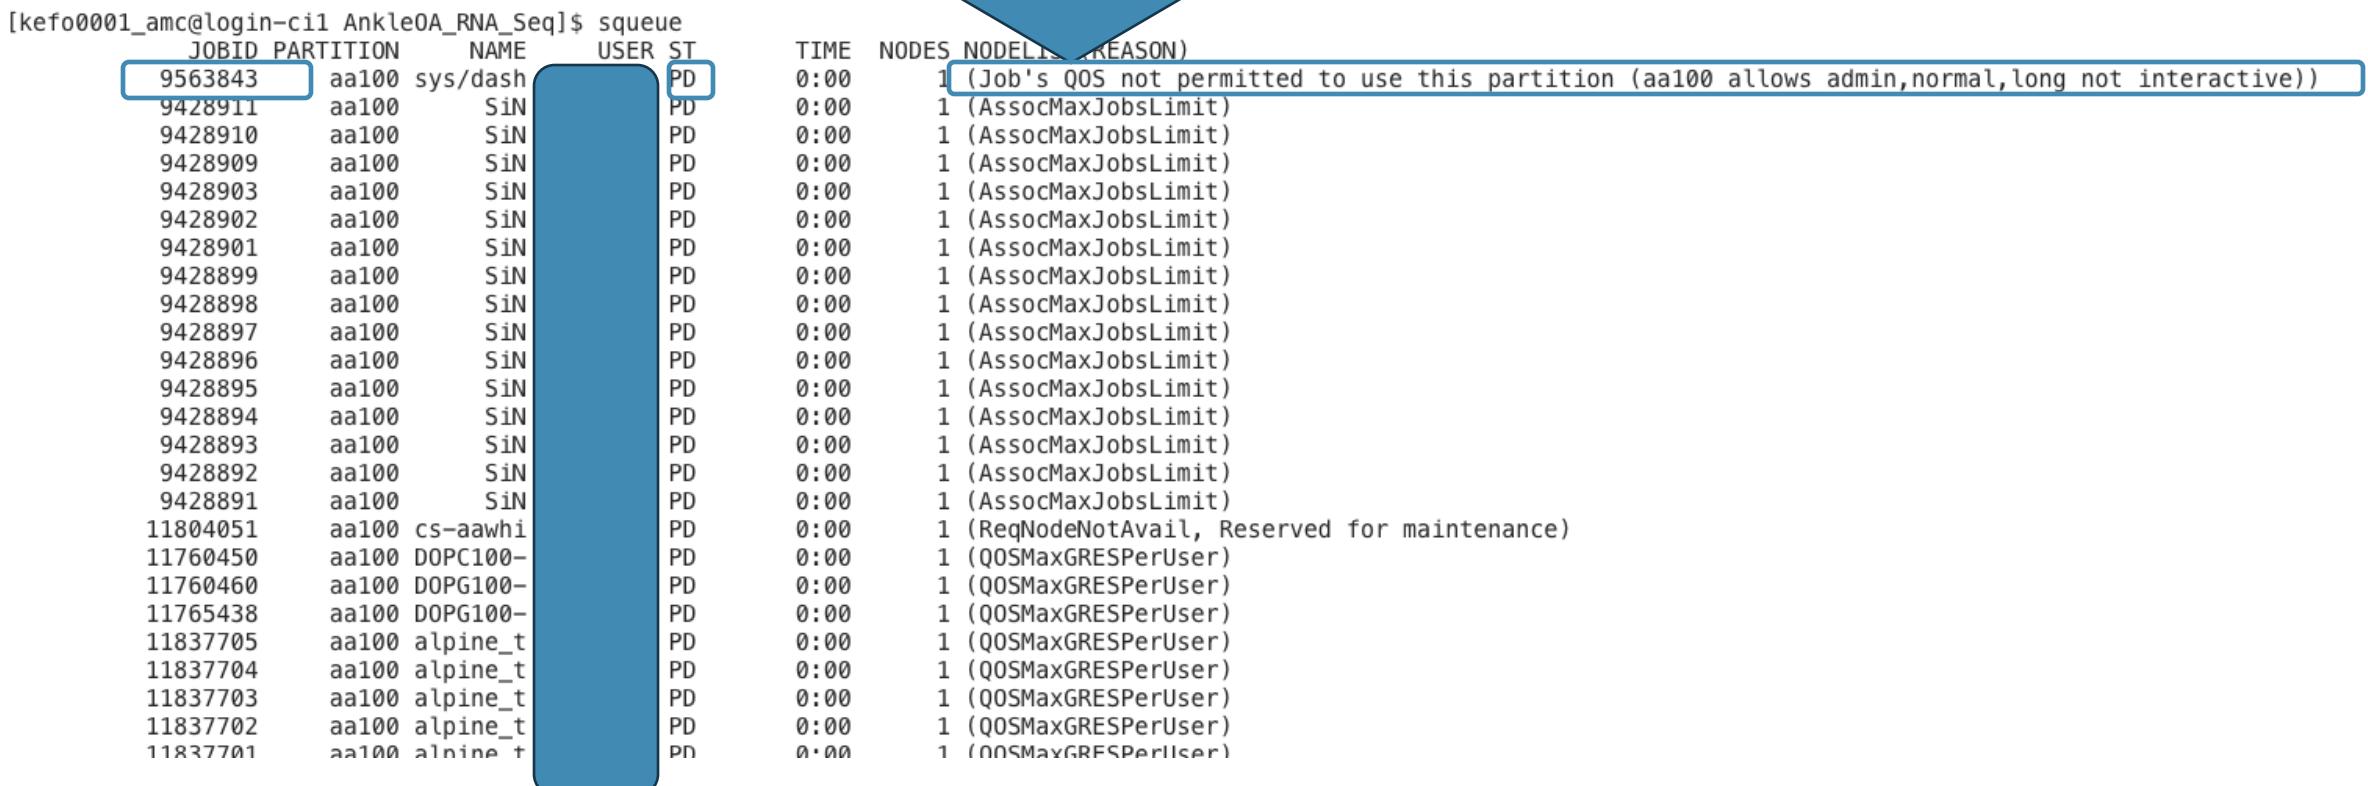
- Allows to see jobs located in the slurm scheduling queue.

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq1$ squeue
   JOBID PARTITION      NAME   ST    TIME  NODES NODELIST(REASON)
      9563843      aa100  sys/dash  PD   0:00      1 (Job's QOS not permitted to use this partition (aa100 allows admin,normal,long not interactive))
      9428911      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428910      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428909      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428903      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428902      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428901      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428899      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428898      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428897      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428896      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428895      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428894      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428893      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428892      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
      9428891      aa100       SiN  PD   0:00      1 (AssocMaxJobsLimit)
     11804051      aa100  cs-aawhi  PD   0:00      1 (ReqNodeNotAvail, Reserved for maintenance)
     11760450      aa100  DOPC100-
     11760460      aa100  DOPG100-
     11765438      aa100  DOPG100-
     11837705      aa100  alpine_t  PD   0:00      1 (QOSMaxGRESPerUser)
     11837704      aa100  alpine_t  PD   0:00      1 (QOSMaxGRESPerUser)
     11837703      aa100  alpine_t  PD   0:00      1 (QOSMaxGRESPerUser)
     11837702      aa100  alpine_t  PD   0:00      1 (QOSMaxGRESPerUser)
     11837701      aa100  alpine_t  PD   0:00      1 (QOSMaxGRESPerUser)
```



squeue

- Job 9563843 is pending because it is trying to use aa100 as an interactive partition for #SBATCH --qos



JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST	REASON
9563843	aa100	sys/dash		PD	0:00	1	(Job's QOS not permitted to use this partition (aa100 allows admin,normal,long not interactive))	
9428911	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428910	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428909	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428903	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428902	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428901	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428899	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428898	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428897	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428896	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428895	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428894	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428893	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428892	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
9428891	aa100	SiN		PD	0:00	1	(AssocMaxJobsLimit)	
11804051	aa100	cs-aawhi		PD	0:00	1	(ReqNodeNotAvail, Reserved for maintenance)	
11760450	aa100	DOPC100-		PD	0:00	1	(QOSMaxGRESPerUser)	
11760460	aa100	DOPG100-		PD	0:00	1	(QOSMaxGRESPerUser)	
11765438	aa100	DOPG100-		PD	0:00	1	(QOSMaxGRESPerUser)	
11837705	aa100	alpine_t		PD	0:00	1	(QOSMaxGRESPerUser)	
11837704	aa100	alpine_t		PD	0:00	1	(QOSMaxGRESPerUser)	
11837703	aa100	alpine_t		PD	0:00	1	(QOSMaxGRESPerUser)	
11837702	aa100	alpine_t		PD	0:00	1	(QOSMaxGRESPerUser)	
11837701	aa100	alpine_t		PD	0:00	1	(QOSMaxGRESPerUser)	



squeue

- Job 9428909 is requesting too many jobs!!

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue
   JOBID PARTITION      NAME     USER ST    TIME  NODES NODEL
 9563843  aa100  sys/dash  PD  0:00      1 (Job
 9428911  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428910  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428909  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit) ON)
 9428903  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428902  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428901  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428899  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428898  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428897  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428896  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428895  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428894  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428893  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428892  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428891  aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 11804051  aa100  cs-aawhi  PD  0:00      1 (ReqNodeNotAvail, Reserved for maintenance)
 11760450  aa100  DOPC100-
 11760460  aa100  DOPG100-
 11765438  aa100  DOPG100-
 11837705  aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837704  aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837703  aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837702  aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837701  aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
```



squeue

- Job 9428909 is requesting too many jobs!!
- Note that all jobs associated with that username will eventually run.

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue
   JOBID PARTITION      NAME     USER ST      TIME  NODES NODEL
 9563843    aa100  sys/dash  PD  0:00      1 (Job
 9428911    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428910    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428909    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428903    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428902    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428901    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428899    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428898    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428897    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428896    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428895    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428894    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428893    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428892    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 9428891    aa100      SiN  PD  0:00      1 (AssocMaxJobsLimit)
 11804051   aa100  cs-aawhi  PD  0:00      1 (ReqNodeNotAvail, Reserved for maintenance)
 11760450   aa100  DOPC100-
 11760460   aa100  DOPG100-
 11765438   aa100  DOPG100-
 11837705   aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837704   aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837703   aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837702   aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
 11837701   aa100  alpine_t  PD  0:00      1 (QOSMaxGRESPerUser)
```



squeue

- Job 11837705 probably requested too many gpus at once!!!
- In this case, the user needs to decrease their amount of gpus requested.

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue
   JOBID PARTITION      NAME     USER ST      TIME  NODES NODES REASON
      9563843    aa100  sys/dash  PD      0:00      1 (Job
      9428911    aa100      SiN  PD      0:00      1 (As
      9428910    aa100      SiN  PD      0:00      1 (As
      9428909    aa100      SiN  PD      0:00      1 (As
      9428903    aa100      SiN  PD      0:00      1 (As
      9428902    aa100      SiN  PD      0:00      1 (As
      9428901    aa100      SiN  PD      0:00      1 (As
      9428899    aa100      SiN  PD      0:00      1 (As
      9428898    aa100      SiN  PD      0:00      1 (As
      9428897    aa100      SiN  PD      0:00      1 (As
      9428896    aa100      SiN  PD      0:00      1 (As
      9428895    aa100      SiN  PD      0:00      1 (As
      9428894    aa100      SiN  PD      0:00      1 (As
      9428893    aa100      SiN  PD      0:00      1 (As
      9428892    aa100      SiN  PD      0:00      1 (As
      9428891    aa100      SiN  PD      0:00      1 (As
      11804051    aa100  cs-aawhi  PD      0:00      1 (Req
      11760450    aa100 DOPC100-
      11760460    aa100 DOPG100-
      11765438    aa100 DOPG100-
      11837705    aa100  alpine_t  PD      0:00      1 (QOSM
      11837704    aa100  alpine_t  PD      0:00      1 (QOSM
      11837703    aa100  alpine_t  PD      0:00      1 (QOSM
      11837702    aa100  alpine_t  PD      0:00      1 (QOSM
      11837701    aa100  alpine_t  PD      0:00      1 (QOSM

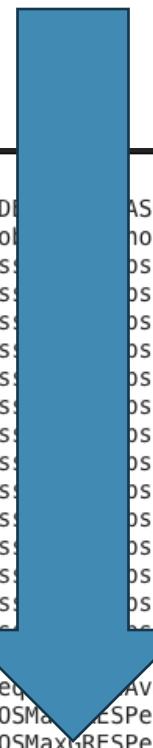
A large blue arrow points from the bottom of the slide towards the squeue command output, highlighting the error message for job 11837705.
```



squeue

- Users are limited to a maximum of 16, 15 and 6 total GPUs on aa100, ami100 and al40 partitions respectively

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue
   JOBID PARTITION      NAME     USER ST      TIME  NODES NODELIST(REASON)
      9563843    aa100  sys/dash  PD      0:00      1 (Job
      9428911    aa100      SiN  PD      0:00      1 (Ass
      9428910    aa100      SiN  PD      0:00      1 (Ass
      9428909    aa100      SiN  PD      0:00      1 (Ass
      9428903    aa100      SiN  PD      0:00      1 (Ass
      9428902    aa100      SiN  PD      0:00      1 (Ass
      9428901    aa100      SiN  PD      0:00      1 (Ass
      9428899    aa100      SiN  PD      0:00      1 (Ass
      9428898    aa100      SiN  PD      0:00      1 (Ass
      9428897    aa100      SiN  PD      0:00      1 (Ass
      9428896    aa100      SiN  PD      0:00      1 (Ass
      9428895    aa100      SiN  PD      0:00      1 (Ass
      9428894    aa100      SiN  PD      0:00      1 (Ass
      9428893    aa100      SiN  PD      0:00      1 (Ass
      9428892    aa100      SiN  PD      0:00      1 (Ass
      9428891    aa100      SiN  PD      0:00      1 (Ass
      11804051    aa100  cs-aawhi  PD      0:00      1 (Req
      11760450    aa100  DOPC100-
      11760460    aa100  DOPG100-
      11765438    aa100  DOPG100-
      11837705    aa100  alpine_t  PD      0:00      1 (QOSMaxGRESPerUser)
      11837704    aa100  alpine_t  PD      0:00      1 (QOSMaxGRESPerUser)
      11837703    aa100  alpine_t  PD      0:00      1 (QOSMaxGRESPerUser)
      11837702    aa100  alpine_t  PD      0:00      1 (QOSMaxGRESPerUser)
      11837701    aa100  alpine_t  PD      0:00      1 (QOSMaxGRESPerUser)


```

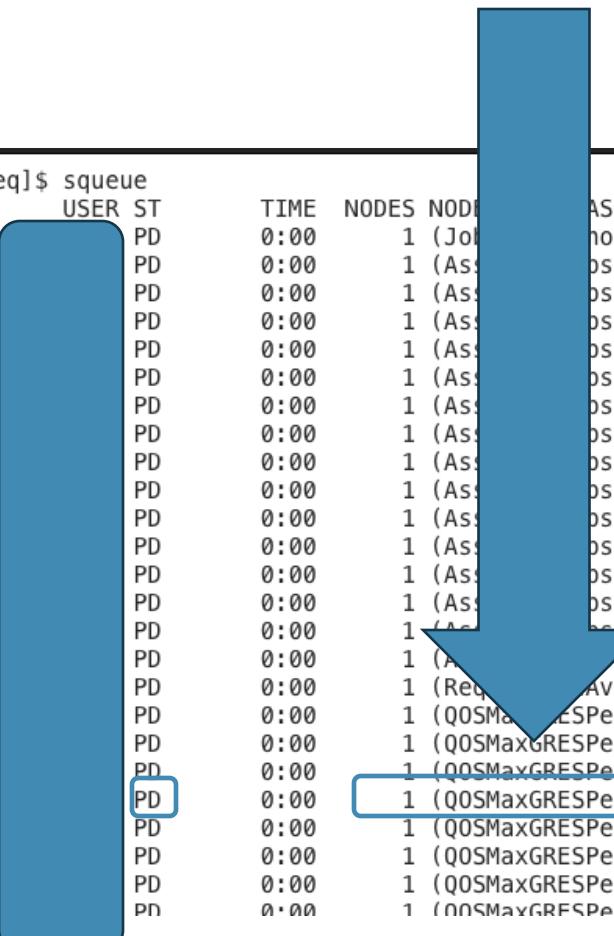


squeue

- A maximum of 22 total gpus per users.

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue
   JOBID PARTITION      NAME     USER ST    TIME  NODES NODES REAS
 9563843  aa100  sys/dash PD 0:00      1 (Job
 9428911  aa100      SiN    PD 0:00      1 (As
 9428910  aa100      SiN    PD 0:00      1 (As
 9428909  aa100      SiN    PD 0:00      1 (As
 9428903  aa100      SiN    PD 0:00      1 (As
 9428902  aa100      SiN    PD 0:00      1 (As
 9428901  aa100      SiN    PD 0:00      1 (As
 9428899  aa100      SiN    PD 0:00      1 (As
 9428898  aa100      SiN    PD 0:00      1 (As
 9428897  aa100      SiN    PD 0:00      1 (As
 9428896  aa100      SiN    PD 0:00      1 (As
 9428895  aa100      SiN    PD 0:00      1 (As
 9428894  aa100      SiN    PD 0:00      1 (As
 9428893  aa100      SiN    PD 0:00      1 (As
 9428892  aa100      SiN    PD 0:00      1 (As
 9428891  aa100      SiN    PD 0:00      1 (As
11804051  aa100  cs-aawhi PD 0:00      1 (Req
11760450  aa100  DOPC100-
11760460  aa100  DOPG100-
11765438  aa100  DOPG100-
11837705  aa100  alpine_t PD 0:00      1 (QOSM
11837704  aa100  alpine_t PD 0:00      1 (QOSM
11837703  aa100  alpine_t PD 0:00      1 (QOSM
11837702  aa100  alpine_t PD 0:00      1 (QOSM
11837701  aa100  alpine_t PD 0:00      1 (QOSM

```



squeue

- --account helps to list only jobs related Anschutz Medical Campus

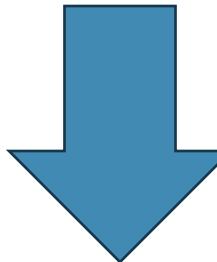


```
[kefo0001_amc@login-c11 AnkleOA_RNA_Seq]$ squeue --account=amc-general
  JOBID PARTITION   NAME   USER ST TIME  NODELIST(REASON)
11838139  acompile acompile      R 19:44  1 c3cpu-a2-u32-1
11838067  acompile acompile      R 25:38  1 c3cpu-a2-u32-1
11838097  acompile acompile      R 23:45  1 c3cpu-a2-u32-1
11838498  acompile acompile      R 2:30   1 c3cpu-a2-u32-1
11717781    amem dissect       PD 0:00   1 (Job's QOS not permitted to use this partition (amem allows admin,mem not normal))
11836668    amem ICM_2nd_       R 1:02:40  1 c3mem-a9-u22-1
11797601  amilan nf-NFCOR      PD 0:00   1 (QOSGrpNodeLimit)
11836259  amilan sys/dash      PD 0:00   1 (Job's QOS not permitted to use this partition (amilan allows admin,normal,long not interactive))
11836309  amilan kp_rf        PD 0:00   1 (QOSGrpNodeLimit)
11838140  amilan NKcellra     PD 0:00   1 (QOSGrpNodeLimit)
11838141  amilan NKcellra     PD 0:00   1 (QOSGrpNodeLimit)
11757411  amilan sbatch       R 4-19:54:23 1 c3cpu-c15-u3-4
11788871  amilan sarek        R 2-00:53:37 1 c3cpu-c11-u15-1
11838133  amilan cellrang     R 19:52   1 c3cpu-c15-u34-1
11837053  amilan 50b981f1     R 55:15   1 c3cpu-c13-u3-3
11838490  amilan ID.nPOD6     R 3:00    1 c3cpu-c9-u1-3
11838473  amilan ID.nPOD6     R 4:08    1 c3cpu-a5-u3-1
11837051  amilan snakemak     R 55:28   1 c3cpu-c15-u3-3
11837862_18 amilan run_rfmi    R 14:30   1 c3cpu-a7-u1-1
11837862_17 amilan run_rfmi    R 14:31   1 c3cpu-a2-u1-1
11837862_16 amilan run_rfmi    R 14:33   1 c3cpu-a7-u7-4
11837862_14 amilan run_rfmi    R 14:34   1 c3cpu-a5-u1-2
11837862_15 amilan run_rfmi    R 14:34   1 c3cpu-c15-u7-1
11837862_13 amilan run_rfmi    R 16:48   1 c3cpu-a7-u28-4
11837862_12 amilan run_rfmi    R 30:14   1 c3cpu-a7-u28-1
11837862_11 amilan run_rfmi    R 30:33   1 c3cpu-a5-u7-1
11836253  amilan sinterac     R 1:47:30  2 c3cpu-c15-u1-2,c3cpu-c15-u3-1
11836307  amilan sys/dash      R 1:25:53  1 c3cpu-c15-u3-1
11836249  amilan picard       R 1:49:23  1 c3cpu-c15-u1-3
11836680 atesting_ sinterac    R 59:29   1 c3gpu-c2-u13
```



squeue

- “`sbatch <slurm_jobscript>` to submit jobs

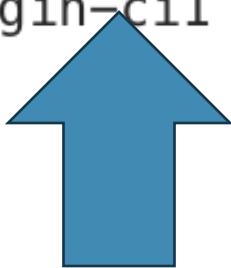


```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ sbatch scripts/dummy_script.sh
Submitted batch job 11838571
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ █
```



squeue

```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ sbatch scripts/dummy_script.sh
Submitted batch job 11838571
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ █
```

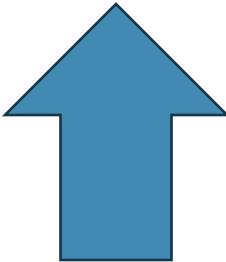


- A submitted job can be monitored with its jobID.



squeue

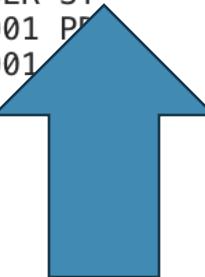
```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ sbatch scripts/dummy_script.sh
Submitted batch job 11838630
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue -u $USER --start
  JOBID PARTITION      NAME      USER ST          START_TIME   NODES SCHEDNODES
 11838630    amilan rsem_pre kefo0001 PD           N/A            1 (null)          NODELIST(REASON)
                                         (Priority)
```



- “--start” option may sometimes provide an estimate as to when the job will actually begin. In this case, it is not working. The job is still pending and we do not have an estimated start time

squeue

```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ squeue --me
JOBID PARTITION      NAME      USER ST    TIME   NODES NODELIST(REASON)
11838630    amilan rsem_pre kefo0001 PR    0:00     1 (Priority)
11838615    amilan rsem_pre kefo0001 PR    2:59     1 c3cpu-a7-u19-3
```



- “squeue –me” gives you information about your jobs only.



squeue

Host: login-ci1.rc.int.colorado.edu

```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ man squeue
```

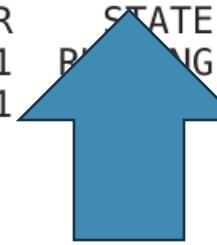


- To get more information about the command.



squeue

```
[kefo0001_amc@login-ci1 AnkleOA_RNA_Seq]$ squeue --me --long
Thu Feb 27 12:00:54 2025
  JOBID PARTITION      NAME      USER      STATE      TIME  TIME_LIMI  NODES NODELIST(REASON)
11838630    amilan rsem_pre kefo0001  RUNNING      1:24  3:00:00      1 c3cpu-a2-u21-1
11838615    amilan rsem_pre kefo0001  RUNNING      5:30  3:00:00      1 c3cpu-a7-u19-3
```



- --long give more detailed information about a running job.



jobstats

```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq] $ module load slurmtools  
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq] $ jobstats $USER 1  
job stats for user kefo0001_amc over past 1 days
```

jobid	jobname	partition	qos	account	cpus	state	start-date-time	elapsed	wait
11807474	acompile	acompile	compile	amc-	1	TIMEOUT	2025-02-26T11:02:12	01:00:28	0 hrs
11808633	acompile	acompile	compile	amc-	4	COMPLETE	2025-02-26T12:05:40	00:22:17	0 hrs
11808861	acompile	acompile	compile	amc-	4	COMPLETE	2025-02-26T12:27:58	00:03:19	0 hrs
11808893	acompile	acompile	compile	amc-	4	FAILED	2025-02-26T12:31:18	00:42:23	0 hrs
11810653	acompile	acompile	compile	amc-	4	COMPLETE	2025-02-26T13:13:44	02:37:01	0 hrs
11820224	acompile	acompile	compile	amc-	4	FAILED	2025-02-26T15:50:47	01:25:33	0 hrs
11823581	acompile	acompile	compile	amc-	4	COMPLETE	2025-02-26T17:16:52	00:42:26	0 hrs
11824523	acompile	acompile	compile	amc-	4	FAILED	2025-02-26T18:00:40	00:24:44	0 hrs
11824714	acompile	acompile	compile	amc-	4	FAILED	2025-02-26T18:25:28	04:15:57	0 hrs
11838571	rsem_pre	amilan	normal	amc-	30	CANCELLED	2025-02-27T11:52:06	00:01:11	0 hrs
11838615	rsem_pre	amilan	normal	amc-	30	CANCELLED	2025-02-27T11:55:24	00:06:09	0 hrs
11838630	rsem_pre	amilan	normal	amc-	30	COMPLETE	2025-02-27T11:59:30	00:44:34	0 hrs
11838957	sinterac	aa100	normal	amc-	5	CANCELLED	2025-02-27T12:16:13	00:00:02	0 hrs

- With jobstats, you can print informations about previous jobs such as start date, jobIDs etc ...



sacct

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
11810653	acompile	acompile	amc-gener+	4	COMPLETED	0:0
11810653.ex+	extern		amc-gener+	4	COMPLETED	0:0
11810653.0	bash		amc-gener+	4	COMPLETED	0:0



- With sacct you can give informations about past jobs



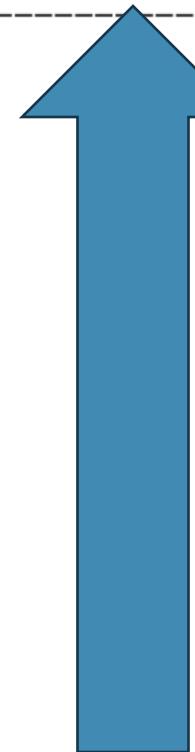
sacct

```
[kefo0001_amc@login-ci1 Ankle0A_RNA_Seq]$ sacct -B -j 11838630  
Batch Script for 11838630
```

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --partition=amilan
#SBATCH --time=03:00:00
#SBATCH --ntasks=30
##SBATCH --mem=100G
#SBATCH --job-name="rsem_prepare_and_run"
#SBATCH -o "rsem_star.%j.out"
#SBATCH -e "rsem_star.%j.err"
#SBATCH --account=amc-general
#SBATCH --qos=normal

#module load apptainer
```



- With sacct you can print the batch script that was used to submit the job



Parallel jobs

```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan is the CPU partition.
#SBATCH --job-name=first_mpirun     # Name of the job. You can check it with "squeue --me -l". Retrospectively you may
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j stands for the jobID number
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j stands for the jobID number as well
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal   # Quality of service. (e.g. Normal qos allows for max walltime of 24 hours)
#SBATCH --nodes=1      # Total number of nodes requested. Always use 1 node unless you are running MPI!!!
#SBATCH --ntasks=4      # Total number of cores requested
#SBATCH --mail-type=ALL # Begin, End and abort
#SBATCH --mail-to=kevin.fotso@cuanschutz.edu # Email of the user
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm

# Load the compiler module that Openmpi depends on.
# The compiler module allows to transfer source code to machine code
module load

#Load OpenMP module to achieve tasks parallelization
module load

# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun -n $SLURM_NTASKS TASKS hostname
```



- Only use 1 node unless you are using **MPI, gnu parallel, or sparkcluster!!!**



Parallel jobs

```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan is the CPU partition.
#SBATCH --job-name=first_mpirun    # Name of the job. You can check it with "squeue --me -l". Retrospectively you may
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j stands for the jobID number
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j stands for the jobID number as well
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal # Quality of service. (e.g. Normal qos allows for max walltime of 24 hours)
#SBATCH --nodes=1    # Total number of nodes requested. Always use 1 node unless you are running MPI!!!
#SBATCH --ntasks=4   # Total number of cores requested
#SBATCH --mail-type=ALL # Begin, End and abort
#SBATCH --mail-to=kevin.fotso@cuanschutz.edu # Email of the user
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm

# Load the compiler module that Openmpi depends on.
# The compiler module allows to transfer source code to machine code
module load

#Load OpenMP module to achieve tasks parallelization
module load

# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun -n $SLURM_NTASKS TASKS hostname
```

- **Warning:** If you call for 2 nodes but your software only uses 1, you will still waste core hours associated with using 2!!!



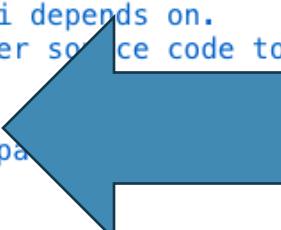
Parallel jobs

```
#!/bin/bash
```

```
#SBATCH --partition=amilan          # Name of the partition. Remember that amilan is the CPU partition.  
#SBATCH --job-name=first_mpirun    # Name of the job. You can check it with "squeue --me -l". Retrospectively you may  
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j stands for the jobID number  
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j stands for the jobID number as well  
#SBATCH --account=amc-general # Account for the campus  
#SBATCH --qos=normal # Quality of service. (e.g. Normal qos allows for max walltime of 24 hours)  
#SBATCH --nodes=1   # Total number of nodes requested. Always use 1 node unless you are running MPI!!!  
#SBATCH --ntasks=4   # Total number of cores requested  
#SBATCH --mail-type=ALL # Begin, End and abort  
#SBATCH --mail-user=kevin.fotso@cuanschutz.edu # Email of the user  
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm
```

```
# Load the compiler that Openmpi depends on.  
# The compiler allows to transfer source code to machine code  
module load gcc  
  
#Load OpenMPI to achieve tasks parallelization  
module load openmpi
```

```
# Run the mpirun command  
# $SLURM_NTASKS is the variable associated with the number of tasks requested.  
mpirun -n $SLURM_NTASKS hostname
```

- 
- gcc is a dependency for openmpi. If you tried to load openmpi alone, it would not have worked. “**module spider openmpi**”



Parallel jobs

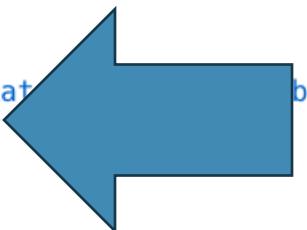
```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan is the CPU partition.
#SBATCH --job-name=first_mpirun    # Name of the job. You can check it with "squeue --me -l". Retrospectively you may
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j stands for the jobID number
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j stands for the jobID number as well
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal # Quality of service. (e.g. Normal qos allows for max walltime of 24 hours)
#SBATCH --nodes=1    # Total number of nodes requested. Always use 1 node unless you are running MPI!!!
#SBATCH --ntasks=4   # Total number of cores requested
#SBATCH --mail-type=ALL # Begin, End and abort
#SBATCH --mail-user=kevin.fotso@cuanschutz.edu # Email of the user
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm

# Load the compiler that Openmpi depends on.
# The compiler allows to transfer source code to machine code
module load gcc

#Load OpenMPI to achieve tasks parallelization
module load openmpi

# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun -n $SLURM_NTASKS hostname
```

- 
- Running MPI on 1 node with 4 cores.



Parallel jobs

```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan
#SBATCH --job-name=first_mpirun     # Name of the job. You can check it with "
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j s
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j st
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal # Quality of service. (e.g. Normal qos allows for max walltim
#SBATCH --nodes=2    # Total number of nodes
#SBATCH --ntasks-per-node=4   # Number of cores per node
#SBATCH --mail-type=ALL # Begin, End, and abort
#SBATCH --mail-user=kevin.fotso@cuanuschutz.edu # Email of the user
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm

# Load the compiler that Openmpi depends on.
# The compiler allows to transfer source code to machine code
module load gcc

#Load OpenMPI to achieve tasks parallelization
module load openmpi

# Get information about the node I am running on:
scontrol show hostname > $SLURM_SUBMIT_DIR/nodelist.txt

# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun --machinefile $SLURM_SUBMIT_DIR/nodelist.txt -n $SLURM_NTASKS hostname
```

Parallel jobs

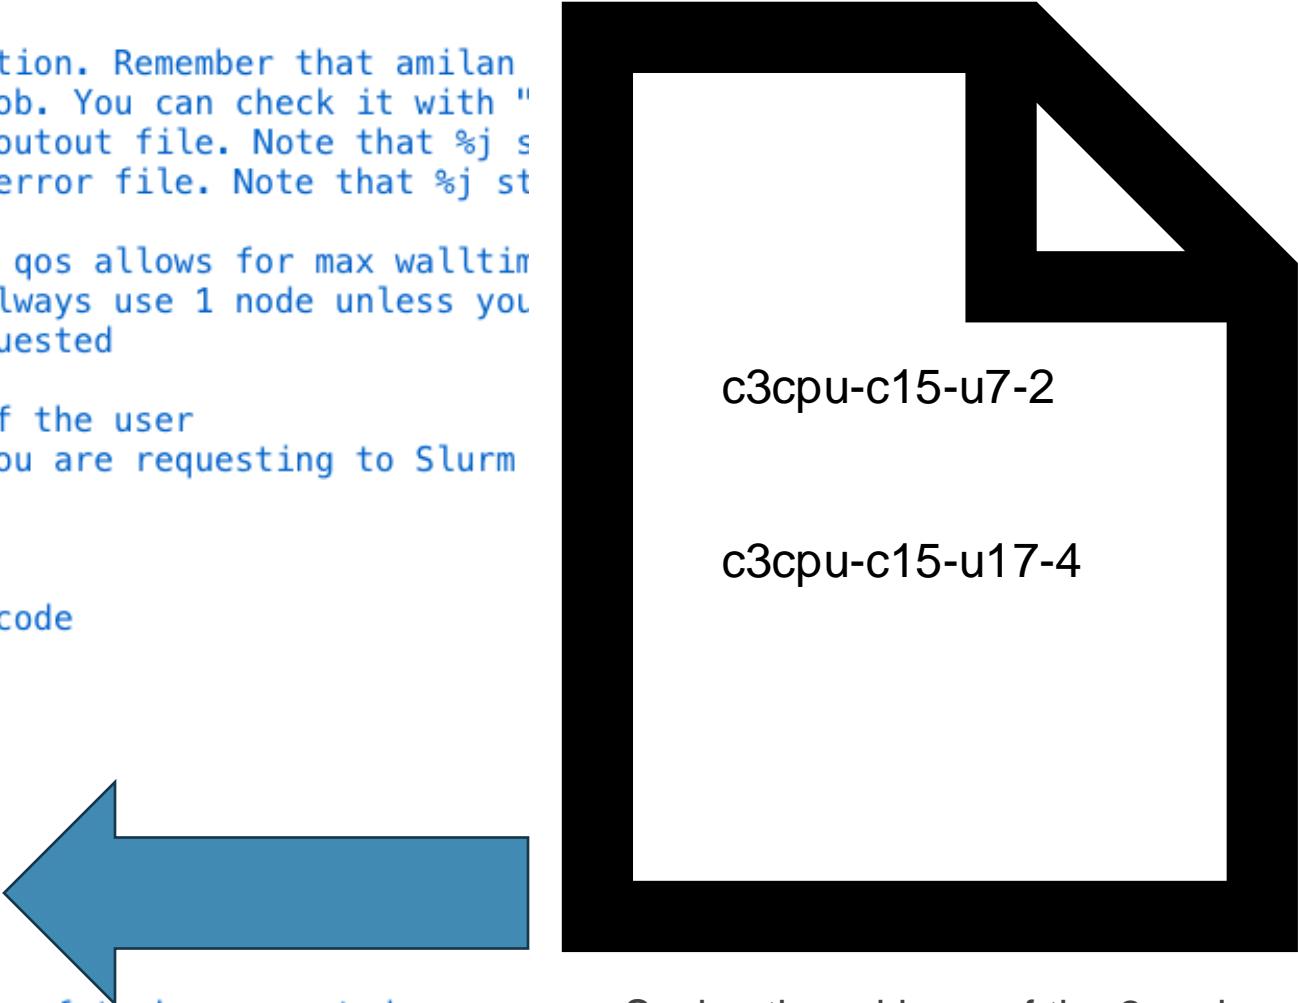
```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan
#SBATCH --job-name=first_mpirun      # Name of the job. You can check it with "
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j s
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j st
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal   # Quality of service. (e.g. Normal qos allows for max walltim
#SBATCH --nodes=2      # Total number of nodes requested. Always use 1 node unless you
#SBATCH --ntasks-per-node=4 # Total number of cores requested
#SBATCH --mail-type=ALL # Begin, End and abort
#SBATCH --mail-user=kevin.fotso@cuanschutz.edu # Email of the user
#SBATCH --time=00:00:01             # Walltime that you are requesting to Slurm

# Load the compiler that Openmpi depends on.
# The compiler allows to transfer source code to machine code
module load gcc

#Load OpenMPI to achieve tasks parallelization
module load openmpi

# Get information about the node I am running on:
scontrol show hostname > $SLURM_SUBMIT_DIR/nodelist.txt
```



```
# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun --machinefile $SLURM_SUBMIT_DIR/nodelist.txt -n $SLURM_NTASKS hostname
```

c3cpu-c15-u7-2

c3cpu-c15-u17-4

- Saving the address of the 2 nodes

Parallel jobs

```
#!/bin/bash

#SBATCH --partition=amilan          # Name of the partition. Remember that amilan
#SBATCH --job-name=first_mpirun      # Name of the job. You can check it with "
#SBATCH --output=first_mpirun-job.%j.out # Name of the output file. Note that %j s
#SBATCH --error=first_mpirun-job.%j.err # Name of the error file. Note that %j st
#SBATCH --account=amc-general # Account for the campus
#SBATCH --qos=normal   # Quality of service. (e.g. Normal qos allows for max walltim
#SBATCH --nodes=2      # Total number of nodes requested. Always use 1 node unless you
#SBATCH --ntasks-per-node=4 # Total number of cores requested
#SBATCH --mail-type=ALL # Begin, End and abort
#SBATCH --mail-user=kevin.fotso@cuanschutz.edu # Email of the user
#SBATCH --time=00:00:01           # Walltime that you are requesting to Slurm
#
# Load the compiler that Openmpi depends on.
# The compiler allows to transfer source code to machine code
module load gcc

#Load OpenMPI to achieve tasks parallelization
module load openmpi

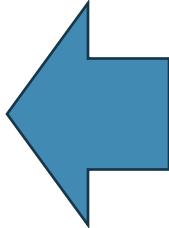
# Get information about the node I am running on:
scontrol show hostname > $SLURM_SUBMIT_DIR/nodelist.txt

# Run the mpirun command
# $SLURM_NTASKS is the variable associated with the number of tasks requested.
mpirun --machinefile $SLURM_SUBMIT_DIR/nodelist.txt -n $SLURM_NTASKS hostname
```

- Running on 2 nodes

Parallel jobs

```
3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
~
```

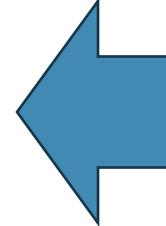


- Proof that MPI ran on 2 nodes.



Parallel jobs

```
3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
~
```

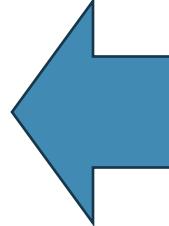


- Proof that MPI ran on 2 nodes.
- We have 4 tasks per node.
- For gnu parallel related applications please refer to this guide on our repository:

https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/Office-hours-presentation-files/GNU_parallel_presentation.pdf

Parallel jobs

```
3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u7-2.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
c3cpu-c15-u17-4.rc.int.colorado.edu
~
```



- Proof that MPI ran on 2 nodes.
- We have 4 tasks per node.
- For gnu parallel related applications please refer to this guide on our repository:

https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/Office-hours-presentation-files/GNU_parallel_presentation.pdf

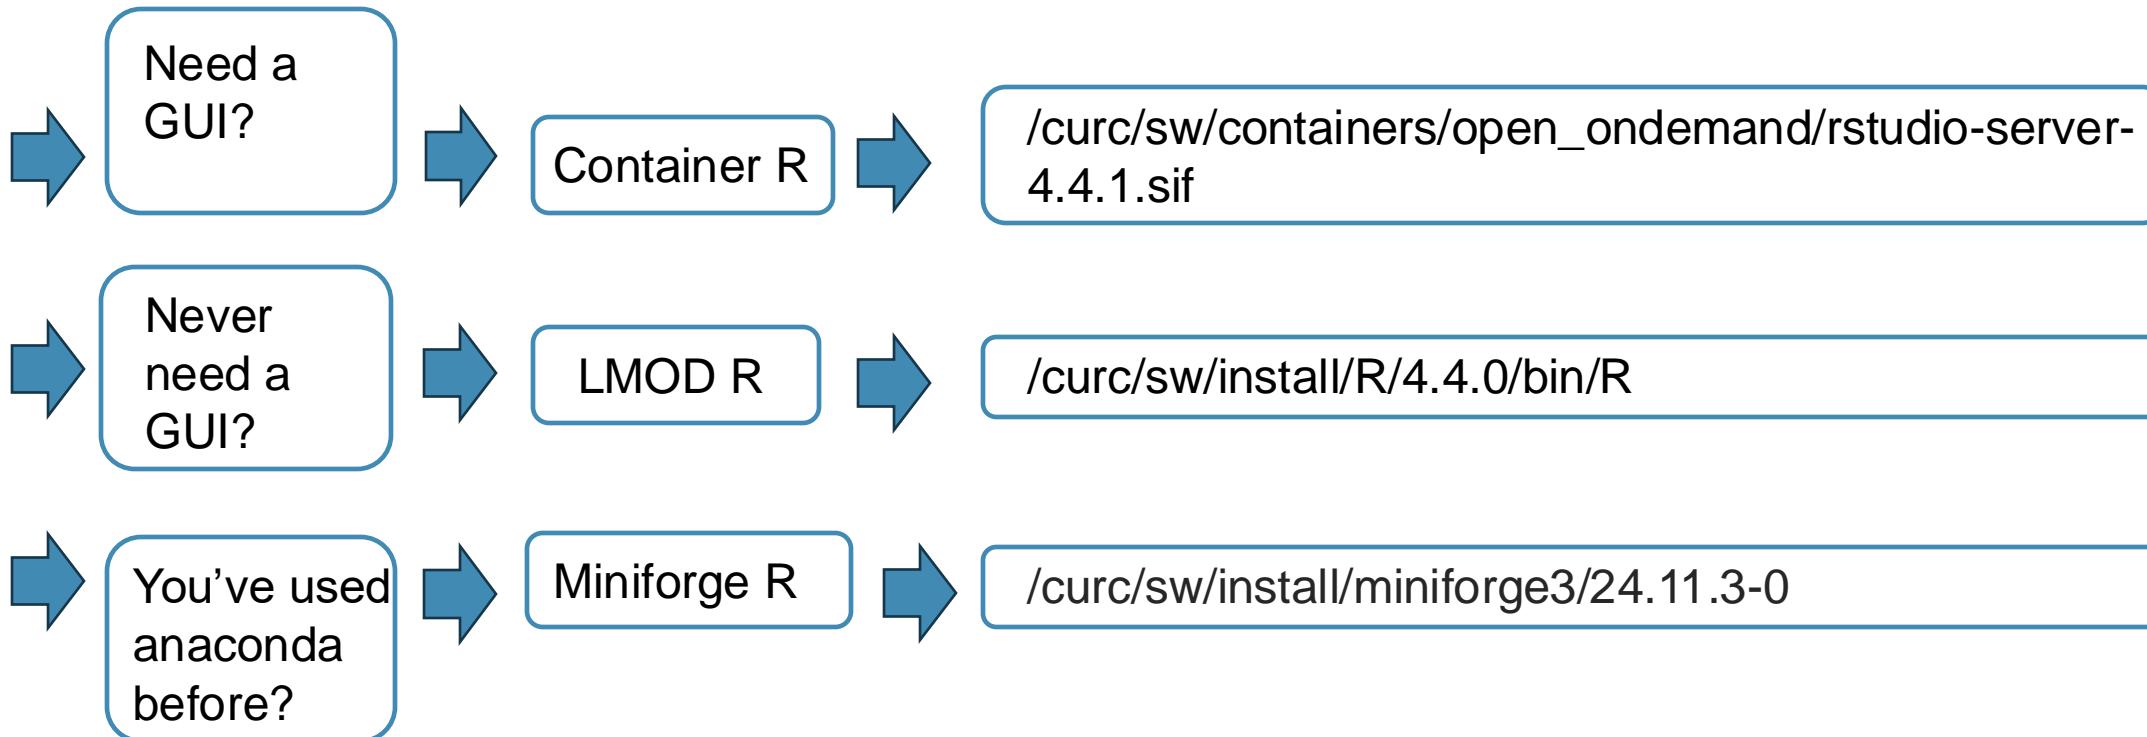
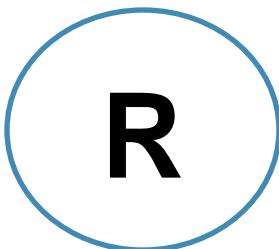
- For information on job arrays please refer to this guide:

https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/Job_array_demonstration.md



What are the 3 ways to use R?

There are 3 totally distinct ways to use R on Alpine.



Preliminary conditions for each way

There are 3 totally distinct ways to use R on Alpine

- Rstudio on Ondemand with versions **4.4.1** and **4.2.2**.
The version **4.4.1** is strongly recommended.
- Under the Alpine shell, there is an R version 4.4.0 that is accessible by using **LMOD**, the modern environment module system for HPC.
- One may decide to install R through **miniforge** or anaconda. The latest version they can install is R **4.4.2**



RStudio Server

RStudio Server

This app will launch RStudio Server, an IDE for R on Alpine.

Before utilizing this application, please see the [RStudio Server](#) and [Configuring Open OnDemand interactive applications](#) sections in our documentation. This documentation includes important information regarding quitting an RStudio session. For more information on installing dependencies required by R packages, please see the [Installing dependencies for RStudio Server](#) section in our documentation.

RStudio Version

Rstudio 2024.04.2, R 4.4.1

Configuration type

Preset configuration

Preset configuration

4 cores, 4 hours

Launch

* The RStudio Server session data for this session can be accessed under the [data root directory](#).



Currently Loaded Modules:

1) jdk/18.0.1.1 2) R/4.4.0

module load R/4.4.0

module list



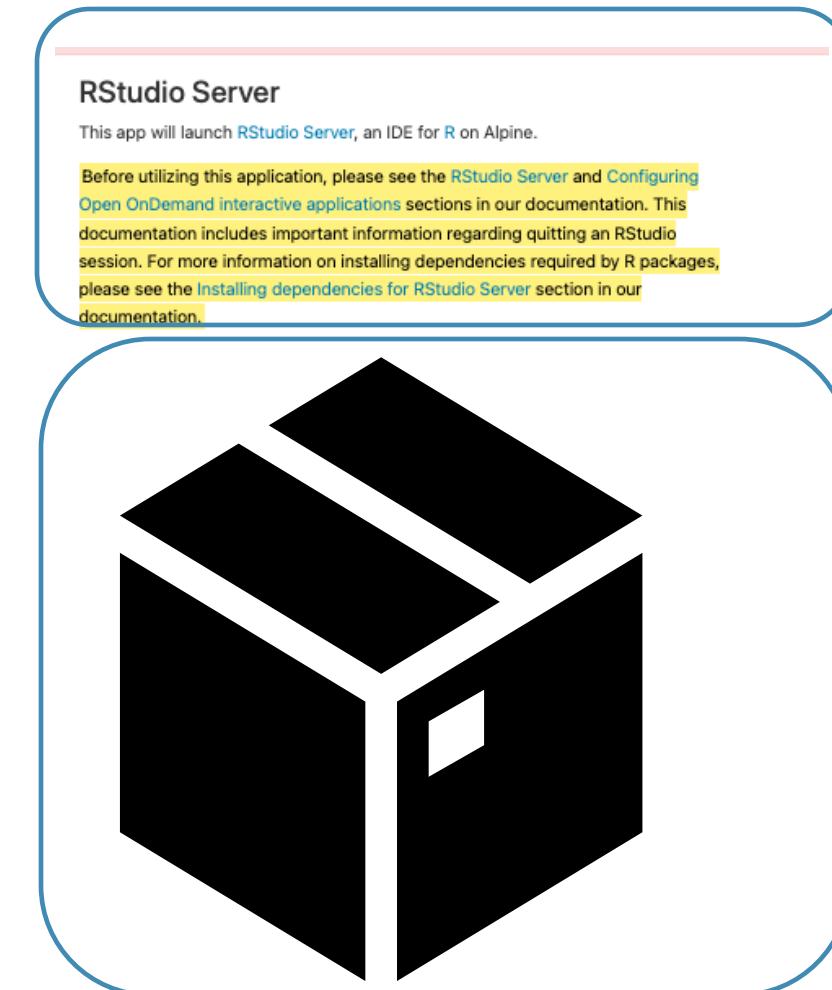
module load miniforge



I) Rstudio on Alpine

What is Ondemand Rstudio?

- It was built on top of an apptainer container image, thus, it runs on a different operating system than Alpine.
- A container is a tool that allows you to run different applications that were built on different operating systems on Alpine.
- The Rstudio .sif image is located at
/curc/sw/containers/open_ondemand/rstudio-server-4.4.1.sif



I) Rstudio on Alpine

Use Rstudio on Alpine if:

- If you are used to Rstudio outside of Alpine
- You will need access to a GUI.
- If you think that most of your pipelines will not require more than 16 cores or 60GB of RAM.
- **Note:** * If you are new to Rstudio on Alpine, please refer

to this guide: https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio_related_scripts

* If you wish to run your Rstudio contained environment as a slurm batch script, please refer to this guide: https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/Rstudio_on_Slurm_.md



University of Colorado
Anschutz Medical Campus

RStudio Server

This app will launch [RStudio Server](#), an IDE for R on Alpine.

Before utilizing this application, please see the [RStudio Server](#) and [Configuring Open OnDemand interactive applications](#) sections in our documentation. This documentation includes important information regarding quitting an RStudio session. For more information on installing dependencies required by R packages, please see the [Installing dependencies for RStudio Server](#) section in our documentation.

RStudio Version

Rstudio 2024.04.2, R 4.4.1

Configuration type

Custom configuration

Cluster

Alpine

Account

amc-general

Partition

ahub

QoS

interactive

Time

6

Number of cores

16

Reservation (default is None)

None

gres options (default is None)

None

Launch

* The RStudio Server session data for this session can be accessed under the [data root directory](#).

II) LMOD R

What is LMOD R?

- It is accessible through the use of **LMOD**, the modern environment module system for HPC.
- You need to be on a compute node interactively or you need to submit a slurm batch script.
- We only recommend to use the version **4.4.0** to install your packages!!!
- Note that some advance package installation might requires some dependencies. Please open a ticket at hpcsupport@cuanschutz.edu if you are stuck!

```
acompile --ntasks=4 --time=12:00:00
```

Compute
node

```
module avail R
```

Check for
LMOD R

```
R/3.6.3  
R/4.2.2  
R/4.4.0
```

```
-----  
module load R/4.4.0  
which R  
/curc/sw/install/R/4.4.0/bin/R
```

Load R/4.4.0
and confirm



University of Colorado
Anschutz Medical Campus

II) LMOD R

Use LMOD R if:

- You will run heavy job on Alpine, such as high memory computation jobs.
- You will not need to use a GUI interactively.
- You need to run distributed computations, such as with gnu parallel.

```
#!/bin/bash

#SBATCH --partition=amilan      #
#SBATCH --job-name=Rscript_test
#SBATCH --output=Rscript_test.%j.out   #
#SBATCH --error=Rscript_test.%j.err   #
#SBATCH --account=amc-general # Account
#SBATCH --qos=normal  # Quality of service
#SBATCH --nodes=1    # Total number of nodes
#SBATCH --ntasks=4    # Total number of tasks
#SBATCH --mail-type=ALL # Begin, End and Error
#SBATCH --mail-user=kevin.fotso@cuanschi.ch
#SBATCH --time=00:00:01

# Load the compiler that R depends on.
module load R/4.4.0

# Run the R command
Rscript R_test.
```



University of Colorado
Anschutz Medical Campus

III) Miniforge R

What is miniforge/anaconda R?

- It gives you the possibility to install many R related packages through miniforge.
- Note that your miniforge R located in your miniforge environment will be very distinct from LMOD R!!!



University of Colorado
Anschutz Medical Campus

III) Miniforge R

Use miniforge/anaconda R if:

- You are a frequent miniconda/anaconda user.
- If you have followed this guide already:

https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/miniforge_migration.md

- If you plan to share the environment you are using with the rest of your lab on Alpine. For example, one can export a miniforge environment as yml.
- Consider this r-irkernel guide if you plan on using Jupyterlab:
https://curc.readthedocs.io/en/latest/open_ondemand/jupyter_session.html

```
acompile --ntasks=4 --time=12:00:00
```

Compute
node

```
module load miniforge
```

Load
miniforge

```
conda create --name r_test_miniforge  
r_test_mini r-base r-essentials
```

Create
ENV



University of Colorado
Anschutz Medical Campus

III) Miniforge R

Use miniforge/anaconda R if:

- You are a frequent miniconda/anaconda user.
- If you have followed this guide already:

https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/miniforge_migration.md

- If you plan to share the environment you are using with the rest of your lab on Alpine. For example, one can export a miniforge environment as yml.
- Consider this r-irkernel guide if you plan on using Jupyterlab:
https://curc.readthedocs.io/en/latest/open_ondemand/jupyter_session.html

```
acompile --ntasks=4 --time=12:00:00
```

Compute
node

```
module load miniforge
```

Load
miniforge

```
conda create --name r_test_miniforge  
r_test_mini r-base r-essentials
```

Create
ENV

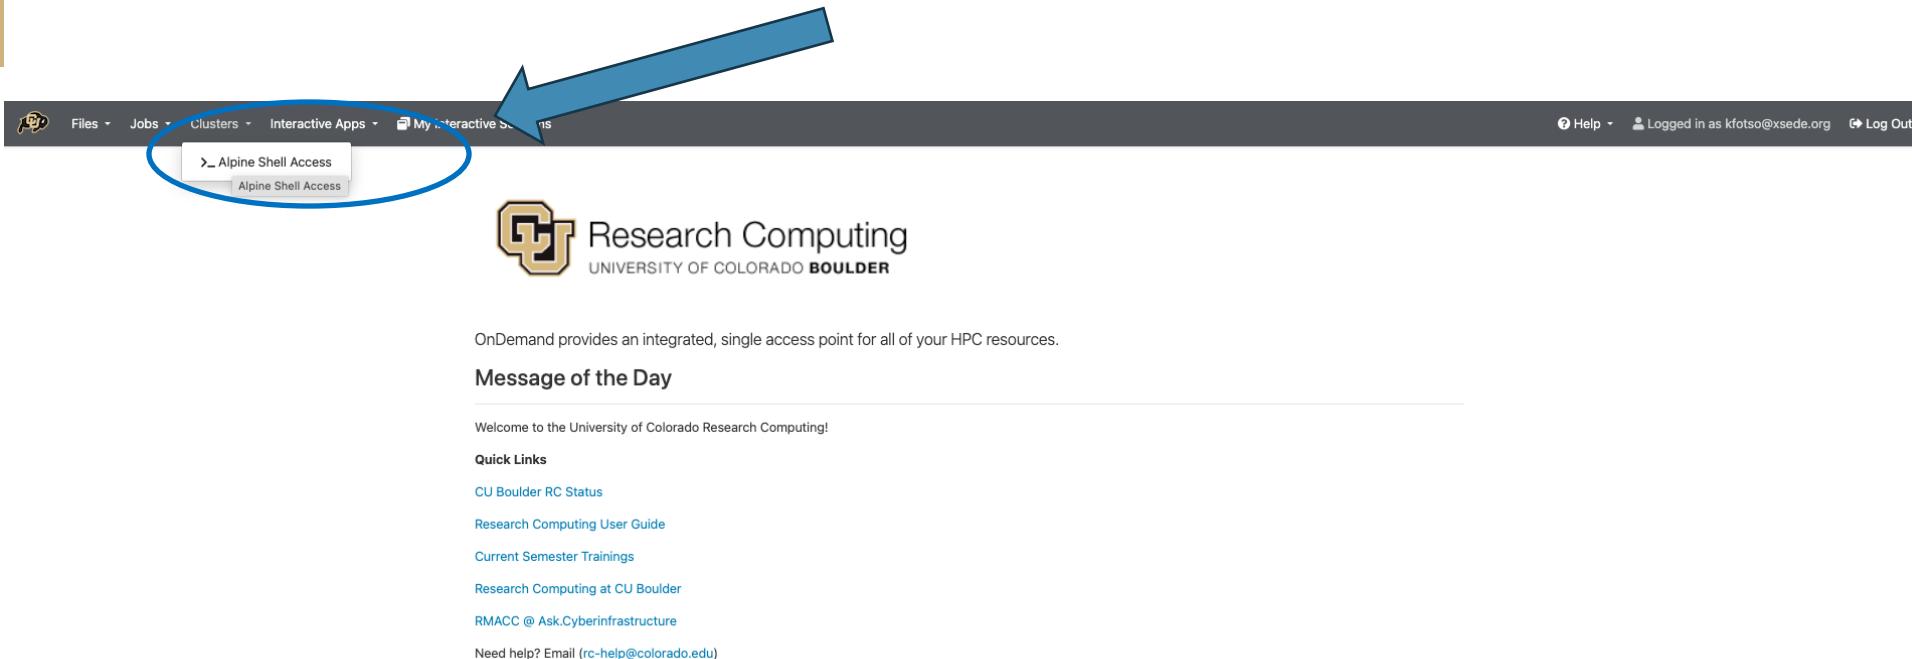
Packages installed under:
[/projects/\\$USER/software/anaconda/envs](/projects/$USER/software/anaconda/envs)



University of Colorado
Anschutz Medical Campus

How to monitor GPU jobs?

- Remember to be inside an Alpine shell!



The screenshot shows the OnDemand web interface. At the top, there is a navigation bar with links for Files, Jobs, Clusters, Interactive Apps, My Interactive Sessions, Help, Log in, and Log Out. A blue arrow points from the text "Remember to be inside an Alpine shell!" down to a button labeled "Alpine Shell Access" which is circled in blue. Below the navigation bar is the Research Computing logo for the University of Colorado Boulder. The main content area includes a message about integrated access to HPC resources, a "Message of the Day" section, and a "Quick Links" sidebar with various links like CU Boulder RC Status, Research Computing User Guide, and Current Semester Trainings.

OnDemand provides an integrated, single access point for all of your HPC resources.

Message of the Day

Welcome to the University of Colorado Research Computing!

Quick Links

- CU Boulder RC Status
- Research Computing User Guide
- Current Semester Trainings
- Research Computing at CU Boulder
- RMACC @ Ask.Cyberinfrastructure

Need help? Email (rc-help@colorado.edu)

How to monitor GPU jobs?

```
sede.org@login-ci1 ~]$ squeue --me --long
```

```
1 09:46:17 2025
```

JOBID	PARTITION	NAME	USER	STATE
12400310	atesting_	sinterac	kfotso@x	RUNNING

```
sede.org@login-ci1 ~]$ ssh c3gpu-c2-u13
```

TIME	TIME_LIMI	NODES	NODELIST(REASON)
0:54	1:00:00	1	c3gpu-c2-u13



- Use "squeue" to get the address of the gpu node where your job is running.



How to monitor GPU jobs?

```
sede.org@login-ci1 ~]$ squeue --me --long
```

```
1 09:46:17 2025
```

JOBID	PARTITION	NAME	USER	STATE
12400310	atesting_	sinterac	kfotso@x	RUNNING

```
sede.org@login-ci1 ~]$ ssh c3gpu-c2-u13
```

TIME	TIME_LIMI	NODES	NODELIST(REASON)
0:54	1:00:00	1	c3gpu-c2-u13

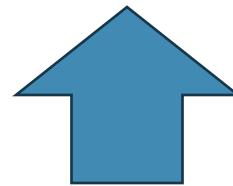


- Use the “**ssh**” command to access that node.
- “**ssh**” is a network protocol that means secure shell.



How to monitor GPU jobs?

```
3 ~]$ watch -n 1 nvidia-smi -l  
3 ~]$
```

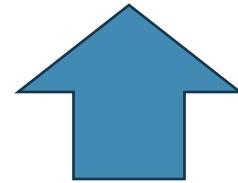


- Watch is a command that allows to constantly **refresh/update** the screen and the command being used



How to monitor GPU jobs?

```
3 ~]$ watch -n 1 nvidia-smi -l  
3 ~]$
```



- Our main command to monitor gpu jobs is nvidia-smi.



How to monitor GPU jobs?

Every 1.0s: nvidia-smi -l

Fri Mar 21 09:47:51 2025

NVIDIA-SMI 550.90.12				Driver Version: 550.90.12		CUDA Version: 12.4					
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC				
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.				
0	NVIDIA A100-PCIE-40GB	On	00000000:21:00.0	Off	0	0	0				
N/A	31C	P0	52W / 250W	75MiB / 40960MiB	N/A	Default	Enabled				
1	NVIDIA A100-PCIE-40GB	On	00000000:81:00.0	Off	0	0	0				
N/A	28C	P0	34W / 250W	75MiB / 40960MiB	N/A	Default	Enabled				
2	NVIDIA A100-PCIE-40GB	On	00000000:E2:00.0	Off	0	0	0				
N/A	29C	P0	34W / 250W	75MiB / 40960MiB	N/A	Default	Enabled				
MIG devices:											
GPU	GI	CI	MIG	Memory-Usage	Vol	Shared					
ID	ID	Dev		BAR1-Usage	SM	CE	ENC	DEC	OFA	JPG	
ECC											
0	1	0	0	38MiB / 19968MiB 0MiB / 32767MiB	42	0	3	0	2	0	0
0	2	0	1	38MiB / 19968MiB 0MiB / 32767MiB	42	0	3	0	2	0	0
1	1	0	0	38MiB / 19968MiB 0MiB / 32767MiB	42	0	3	0	2	0	0

- GPU memory utilization
- Consider installing and using **nvitop** for a more detailed analysis, including for the kernel:

<https://github.com/XuehaiPan/nvitop>



University of Colorado
Anschutz Medical Campus

Need further assistance or resources?

- Visit our github page that has many useful tutorials and workshops geared towards CU Anschutz users:
<https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation>
- Visit the Alpine CURC documentation:
https://curc.readthedocs.io/en/latest/getting_started/navigating_docs.html
- Email us at rc-help@colorado.edu to open a ticket





University of Colorado **Anschutz Medical Campus**

THANK YOU