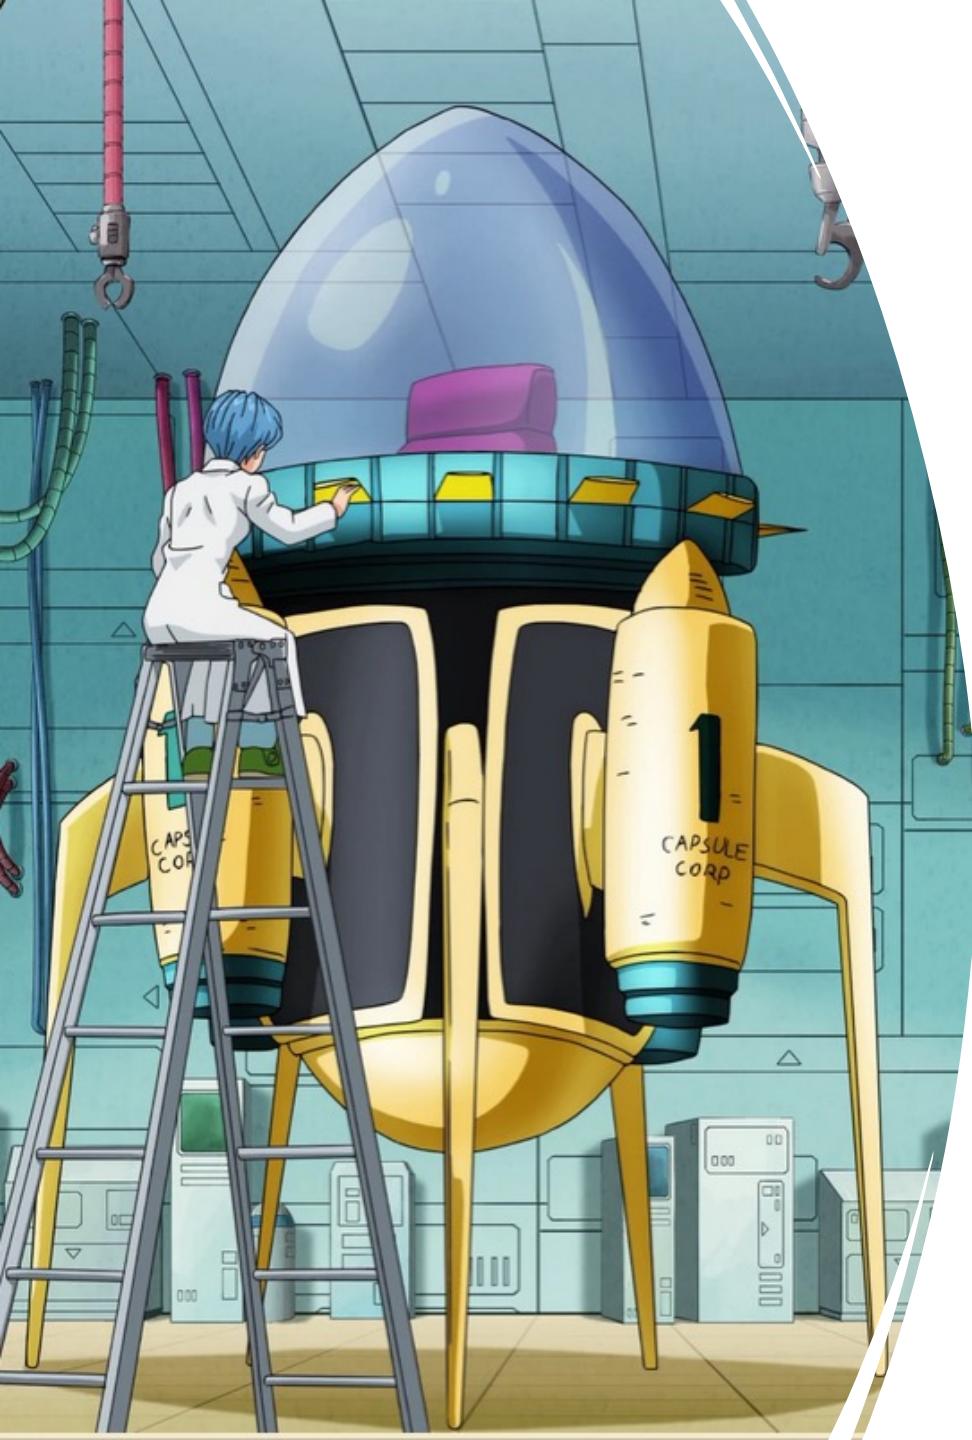


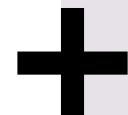
# Introduction to Alpine

Kevin Fotso





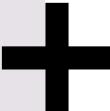
# I- Part 1) Elements of computing



# Parts of the processing chip (1)

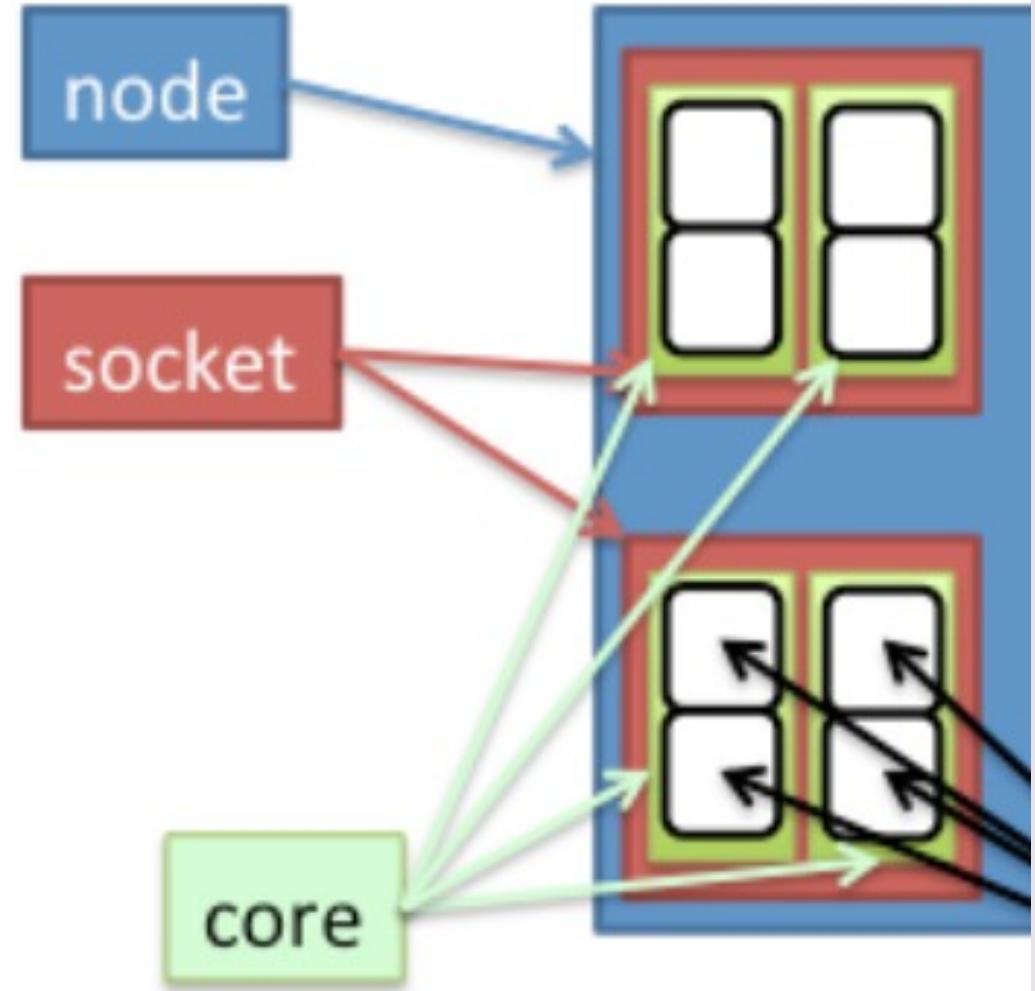


- A socket can be considered an entire computing chip in modern computers.
- Contains more than 1 core.
- On supercomputers a node usually contains more than 1 socket.

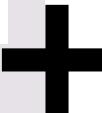
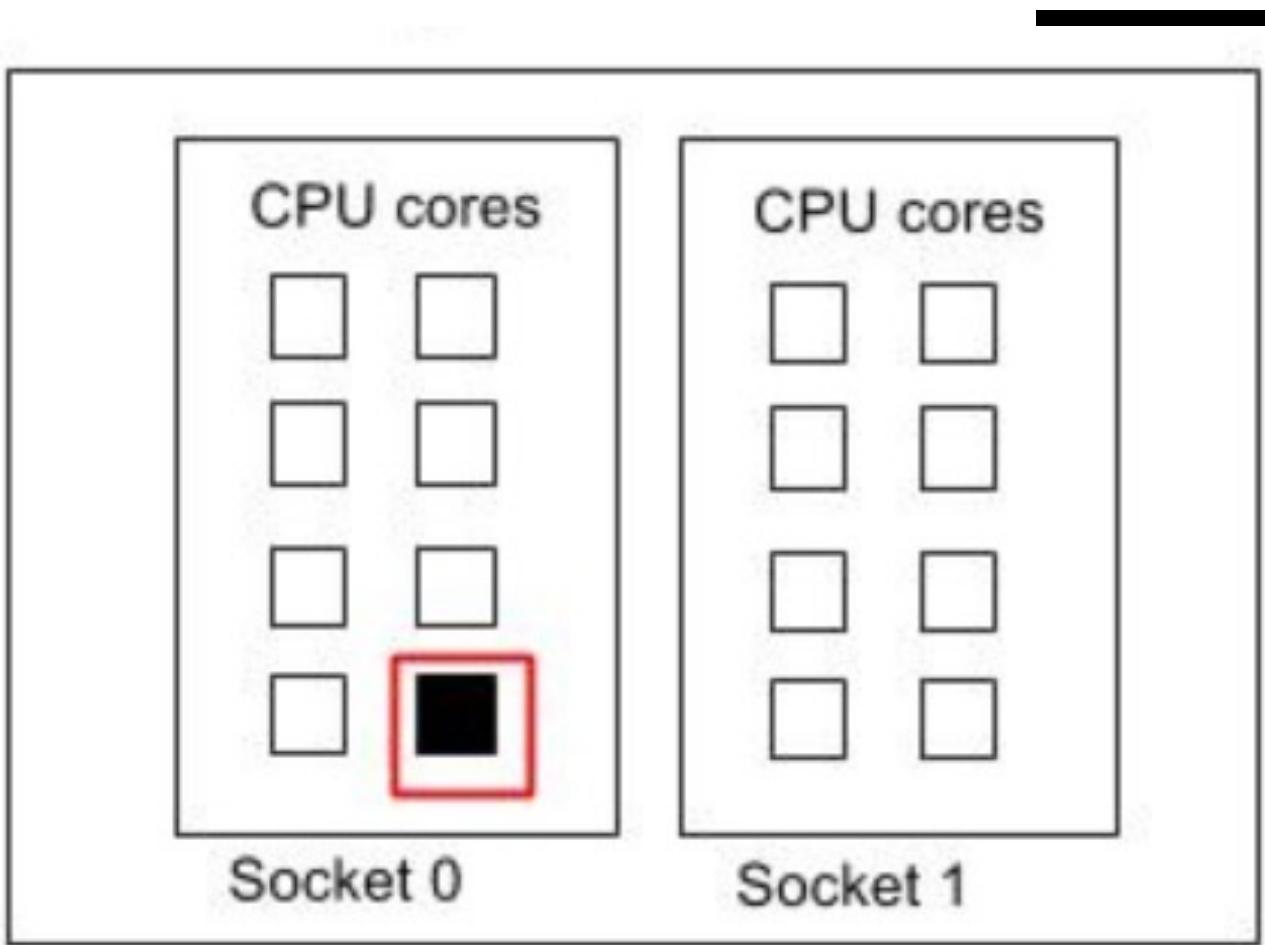


# Parts of the processing chip (2)

A core contains an Arithmetic and Logic Unit (**ALU**) + Registers

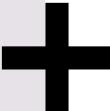
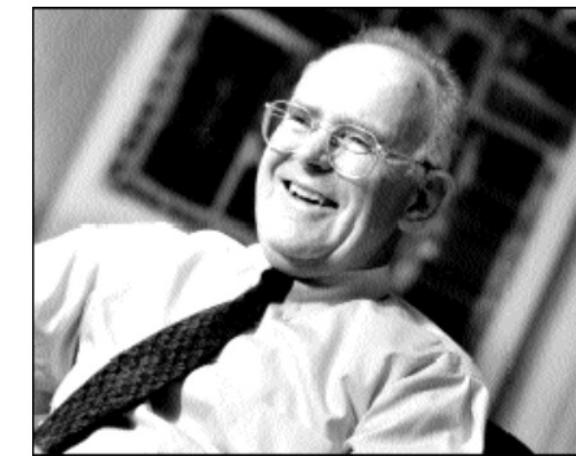
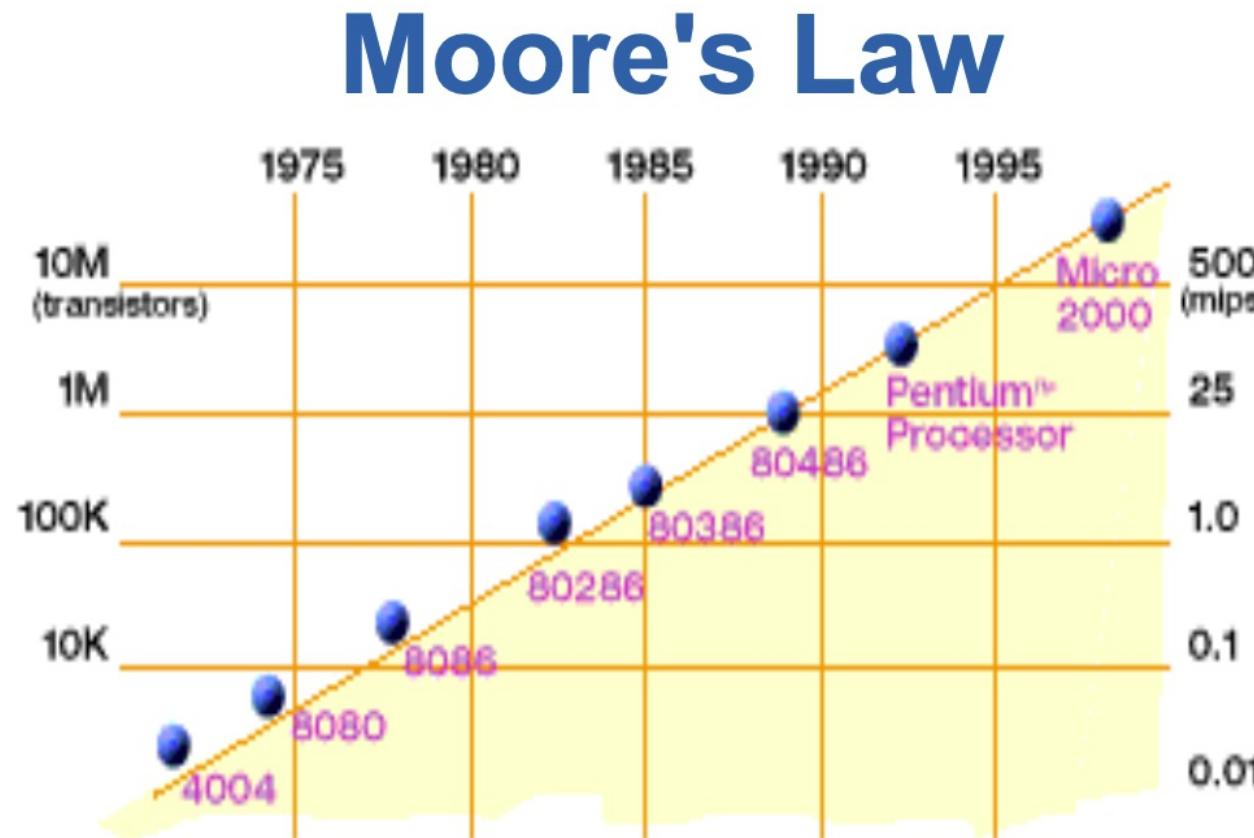


## Parts of the processing chip (2)



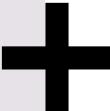
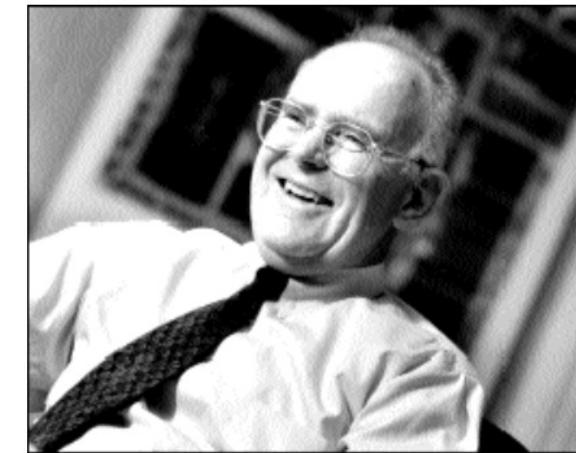
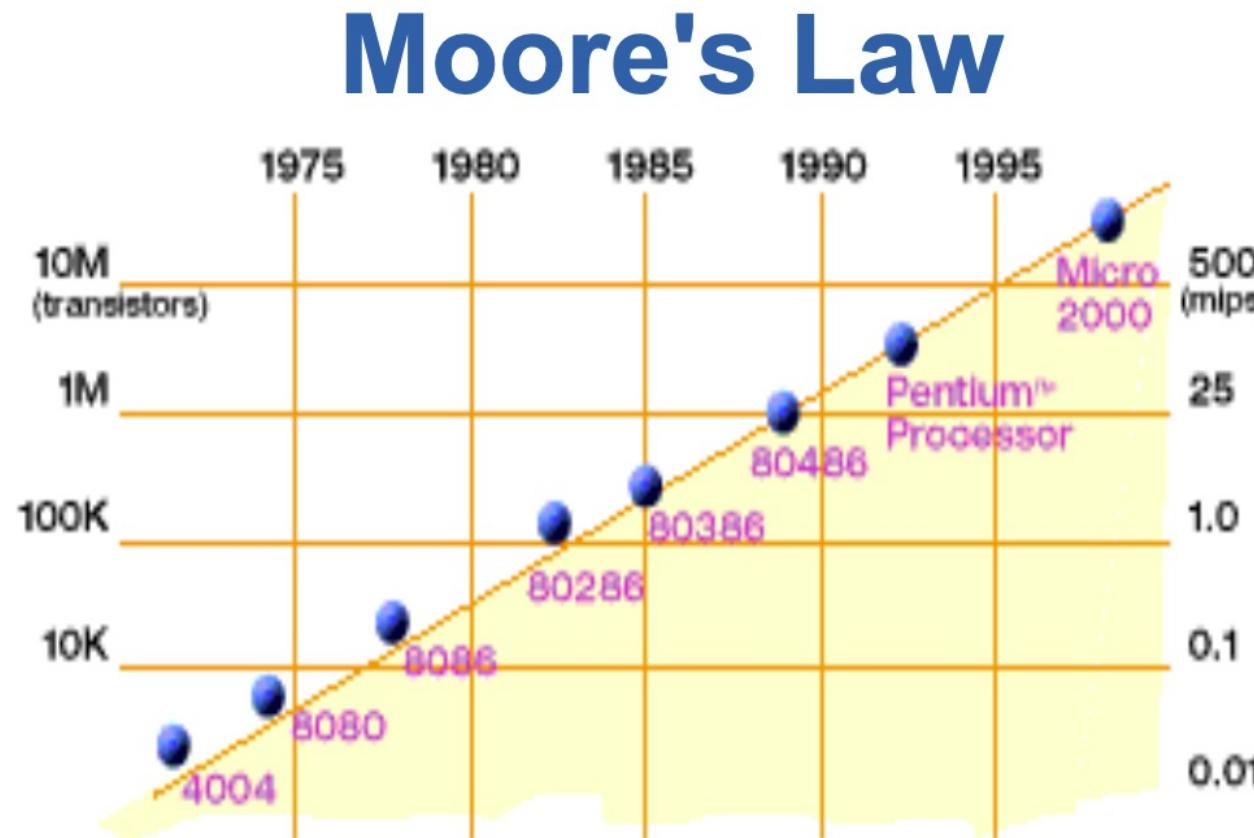
# Moore's law (1)

- Transistor count doubles in a chip every 18 months.



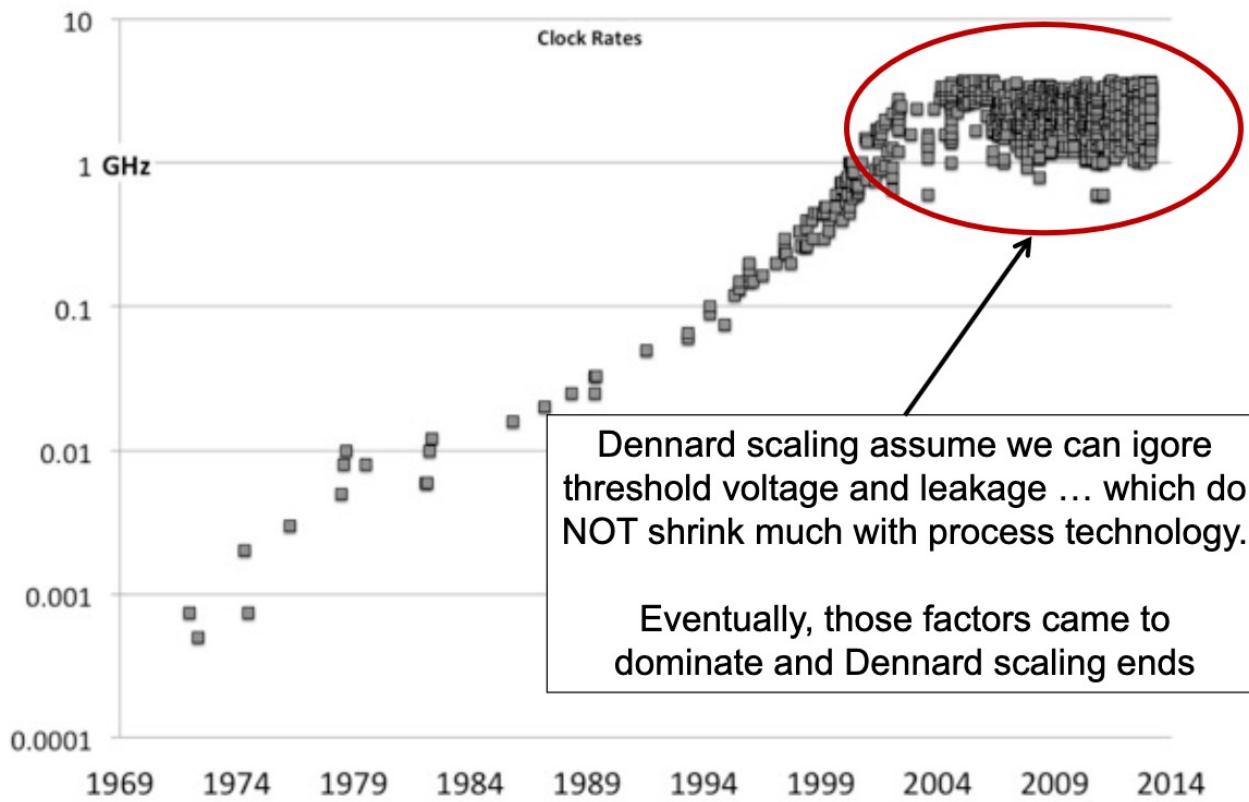
# Moore's law (2)

- Worked until around 2010 where it started to plateau

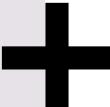


# Moore's law plateau

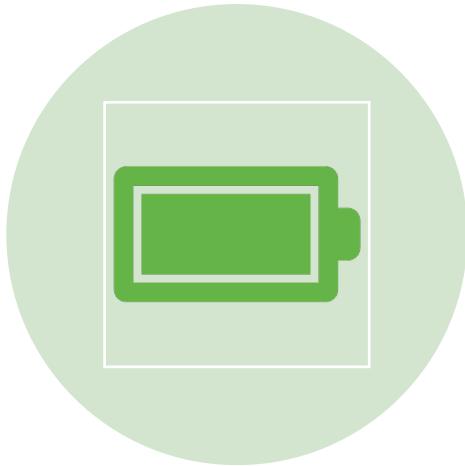
CPU Frequency (GHz) over time (years)



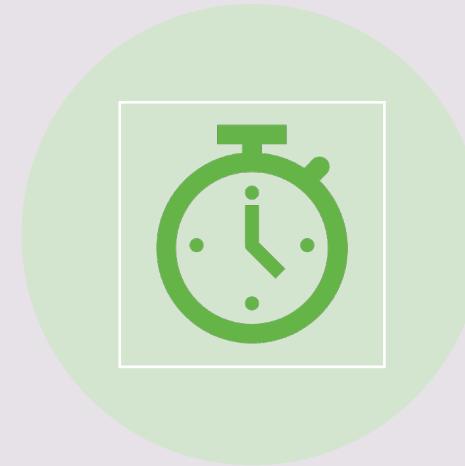
- Number of cycles per seconds plateaued in the 2010s
- Why?



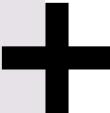
# Why?

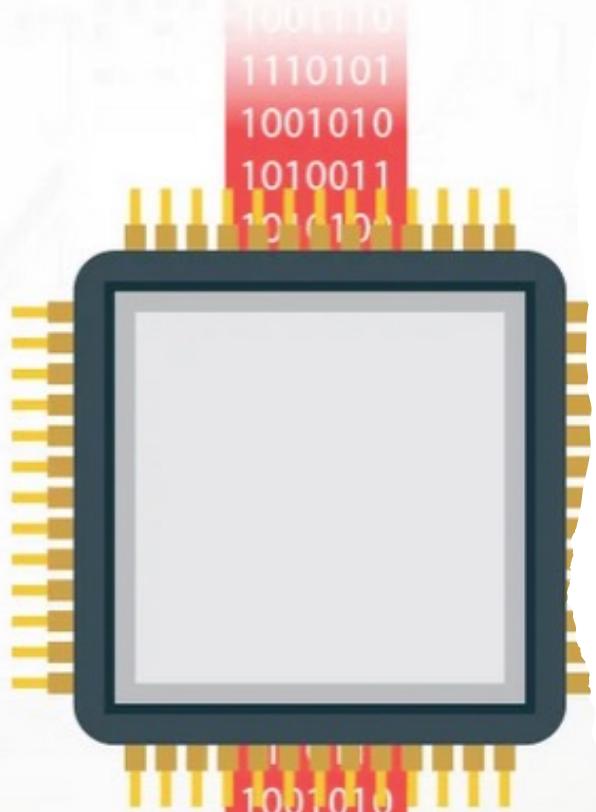


TRADITIONALLY, POWER DENSITY STAYS CONSTANT AS CHIPS FEATURE GOT SMALLER.



HOWEVER AROUND 2010, CLOCK FREQUENCY COULD NOT BE RAISED FURTHER BECAUSE THE HEAT DISSIPATION OR POWER WOULD HAVE GONE TOO FAR.



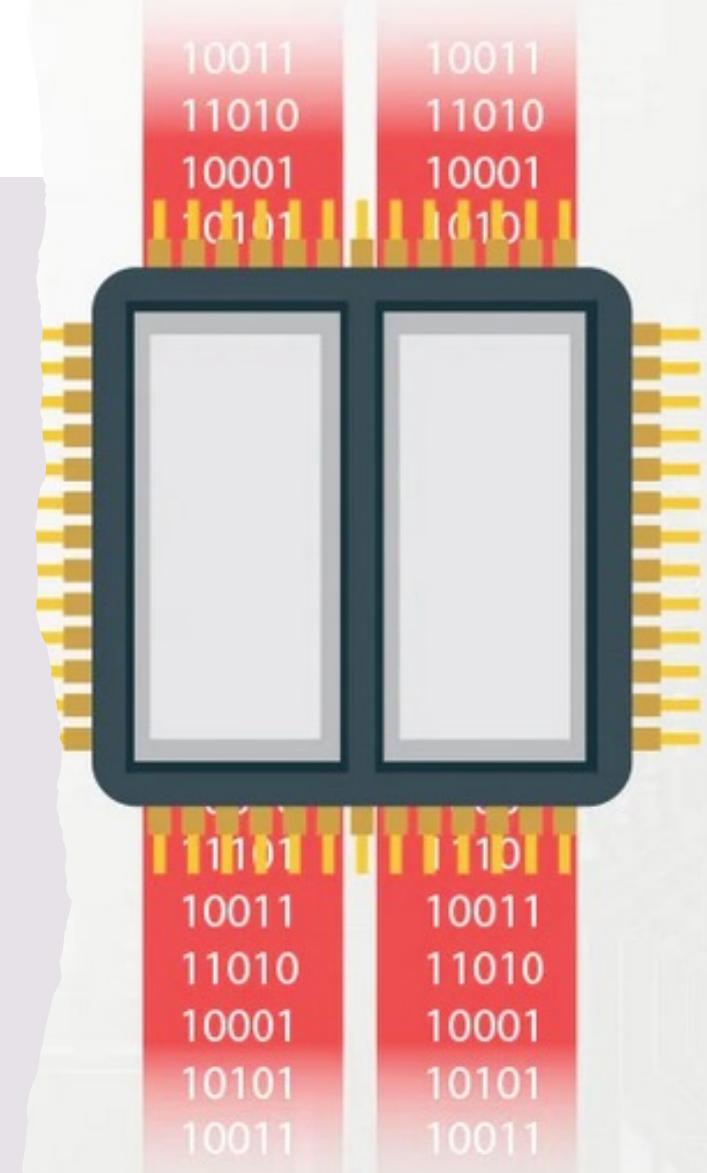


**Single-core**

# Increase need of parallelism

Reduce power by adding cores of small frequency

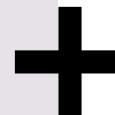
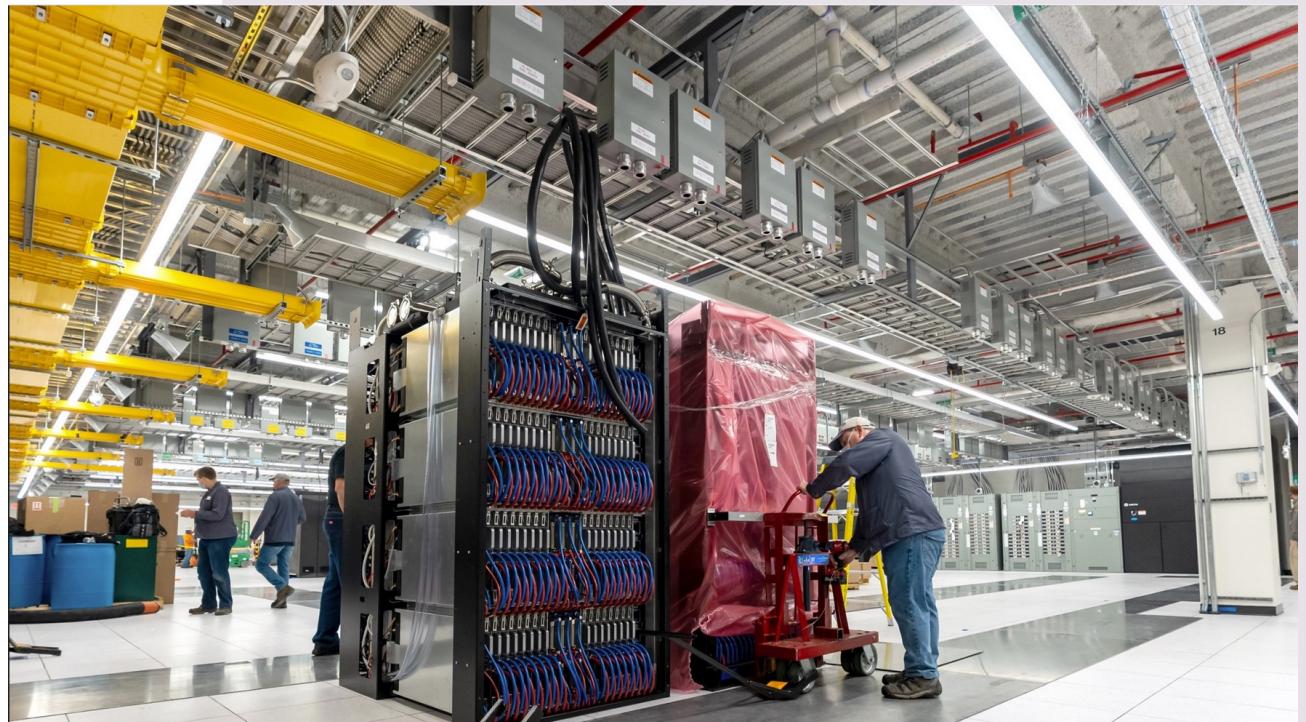
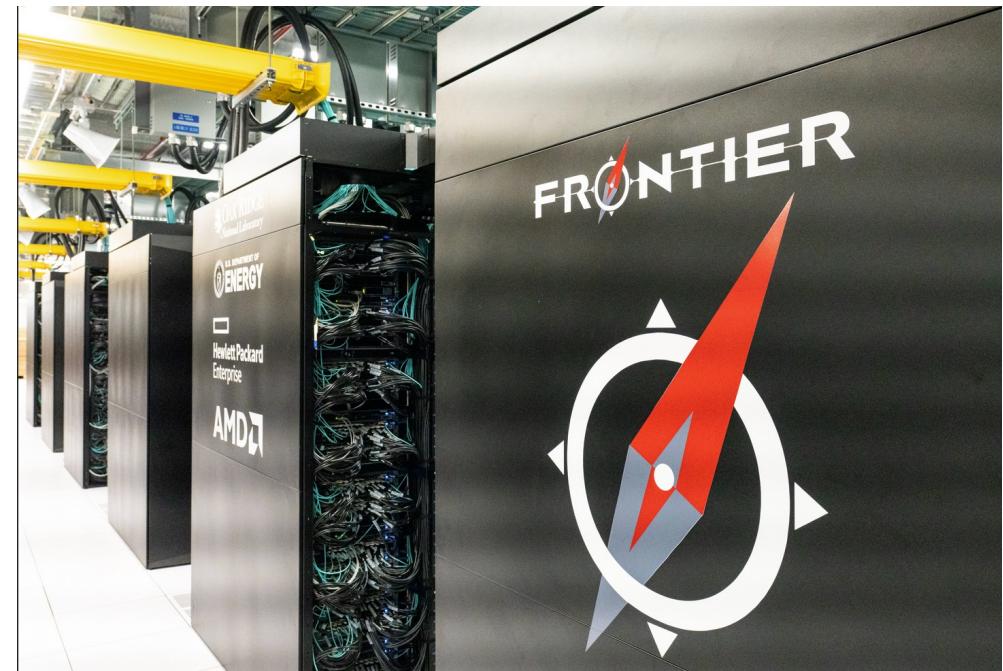
2 cores at half frequency  $f$  can achieve the same throughput as a single core of frequency  $f$



**Dual Core**

# A look at an HPC cluster(1)

- Frontier supercomputer OAK Ridge National lab.  
[AMD Instinct MI250X GPU; 8,699,904 Cores]
- Most powerful at the moment benchmarking with Linpack (e.g.  $Ax=b$ )



# A look at an HPC cluster(2)



June 2023: The TOP 10 Systems (52% of the Total Performance of Top500)

- Alpine has >18,000 cores!

Rank	Site	Computer	Country	Cores	Rmax [Pflops]	% of Peak	Power [MW]	GFlops/Watt
1	DOE / OS Oak Ridge Nat Lab	Frontier, HPE Cray Ex235a, AMD 3rd EPYC 64C, 2 GHz, <b>AMD Instinct MI250X</b> , Slingshot 10	USA	8,699,904	1,194	71	22.7	52.6
2	RIKEN Center for Computational Science	Fugaku, ARM A64FX (48C, 2.2 GHz), Tofu D Interconnect	Japan	7,299,072	442.	82	29.9	14.8
3	EuroHPC /CSC	LUMI, HPE Cray EX235a, AMD 3rd EPYC 64C, 2 GHz, <b>AMD Instinct MI250X</b> , Slingshot 10	Finland	1,268,736	304.	72	2.94	52.3
4	EuroHPC/CINECA	BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, <b>NVIDIA A100 (108C)</b> , Quad-rail NVIDIA HDR100	Italy	1,824,768	239.	78	7.4	32.2
5	DOE / OS Oak Ridge Nat Lab	Summit, IBM Power 9 (22C, 3.0 GHz), <b>NVIDIA GV100 (80C)</b> , Mellanox EDR	USA	2,397,824	149.	74	10.1	14.7
6	DOE / NNSA Livermore Nat Lab	Sierra, IBM Power 9 (22C, 3.1 GHz), <b>NVIDIA GV100 (80C)</b> , Mellanox EDR	USA	1,572,480	94.6	75	7.44	12.7
7	National Super Computer Center in Wuxi	Sunway TaihuLight, <b>SW26010 (260C)</b> , Custom Interconnect	China	10,649,000	93.0	74	15.4	6.05
8	DOE / OS NERSC - LBNL	Perlmutter HPE Cray EX235n, AMD EPYC 64C 2.45GHz, <b>NVIDIA A100</b> , Slingshot 10	USA	706,304	64.6	71	2.59	27.4
9	NVIDIA Corporation	Selene NVIDIA DGX A100, AMD EPYC 7742 (64C, 2.25GHz), <b>NVIDIA A100 (108C)</b> , Mellanox HDR	USA	555,520	63.4	80	2.64	23.9
10	National Super Computer Center in Guangzhou	Tianhe-2A NUDT, Xeon (12C), <b>MATRIX-2000 (128C)</b> + Custom Interconnect	China	4,981,760	61.4	61	18.5	3.32

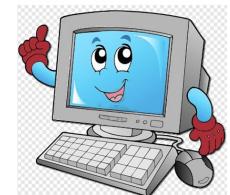
# Anatomy of supercomputing



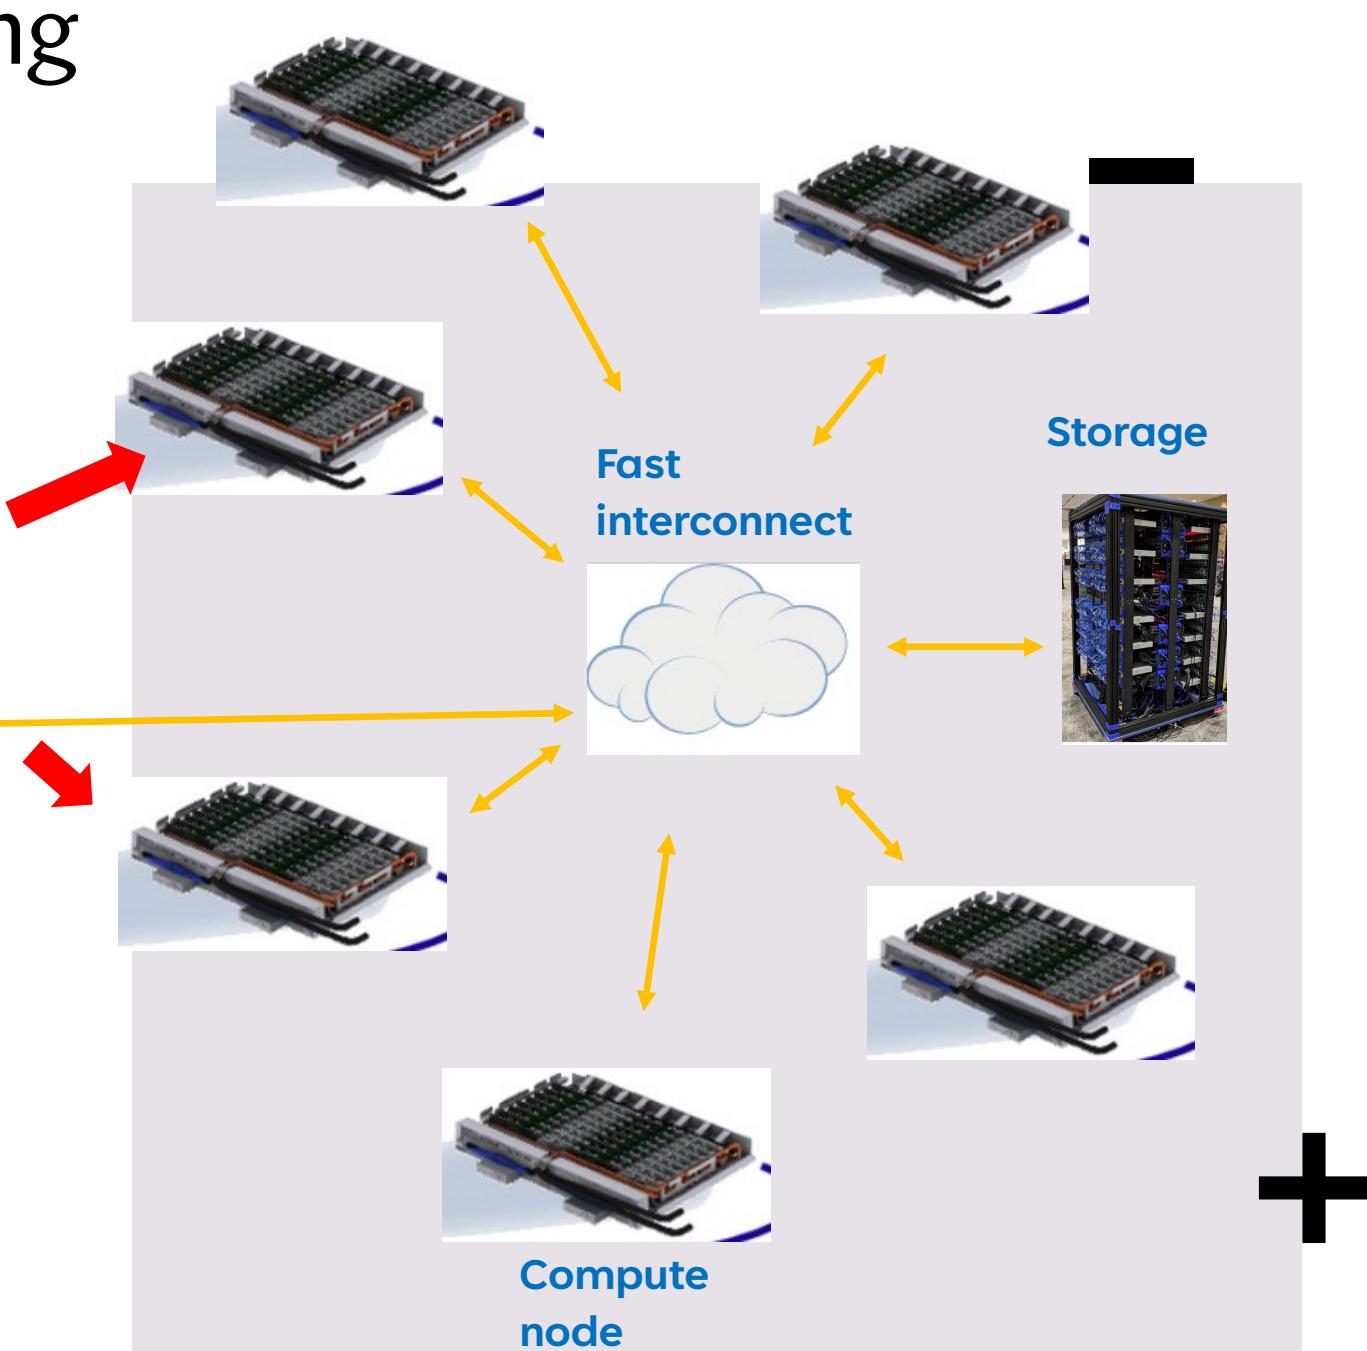
Login node



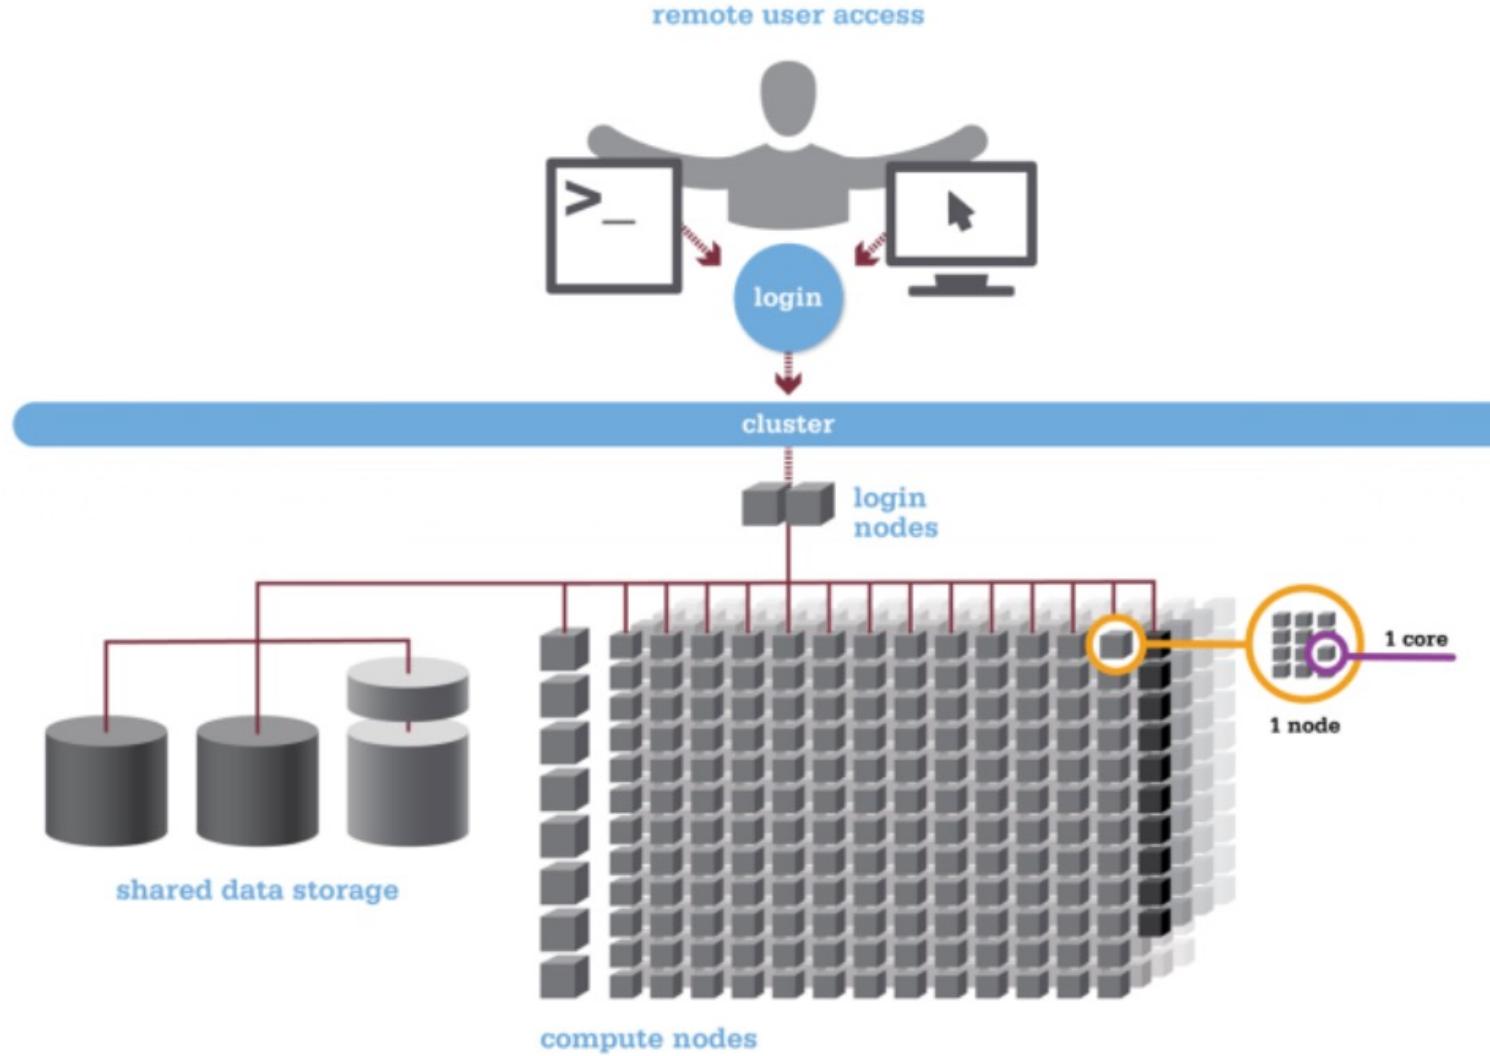
Slurm:  
scheduler



Laptop or  
PC



# Architecture of a supercomputer

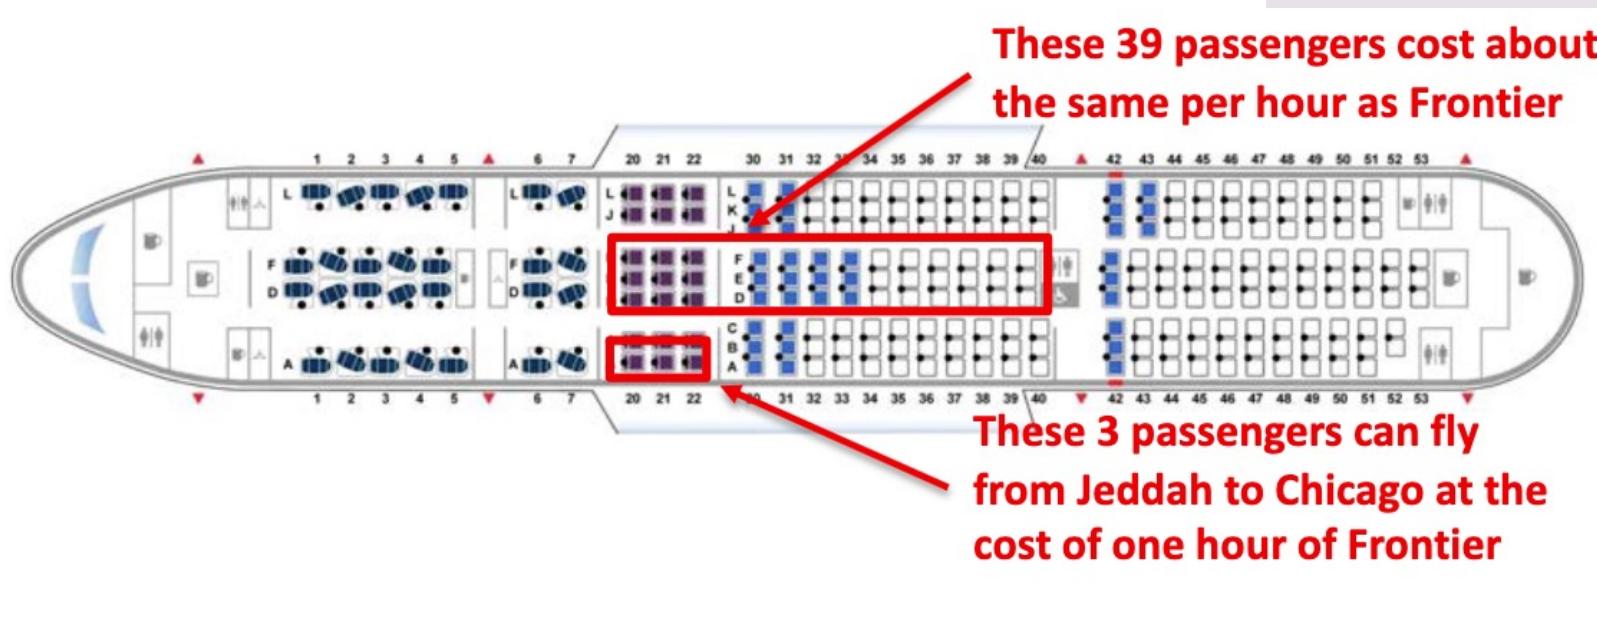


- Login nodes. To log into the system, cd into directories, look at files etc ...
- Compute node. Dedicated to do the computation
- The slurm scheduler controls access to the compute nodes to avoid a tragedy of the commons ...

# Why Slurm?



- Allows for better tracking of resource used per user
- Better management of resource across user base
- More sustainable and less tragedy of the commons
- Widely used and supported by applications in case of bugs.



# Why Slurm?

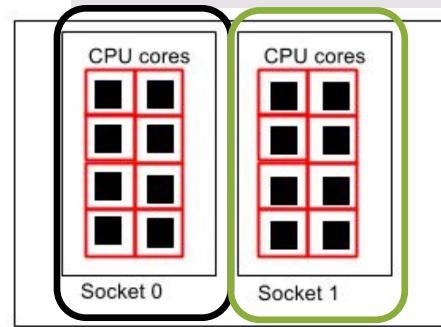
- Bulma request 8 cores from Node 1
- Dr. Brief requests 8 cores as well from Node 1
- Gohan requests 16 cores from Node 2
- No tragedy of the common!!!



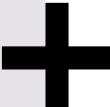
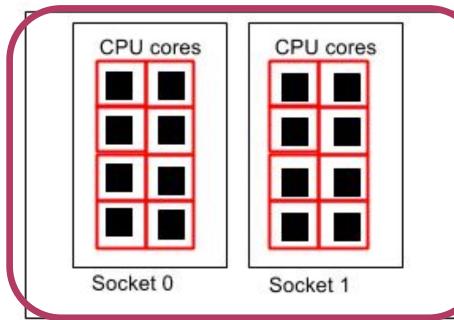
Slurm  
scheduler



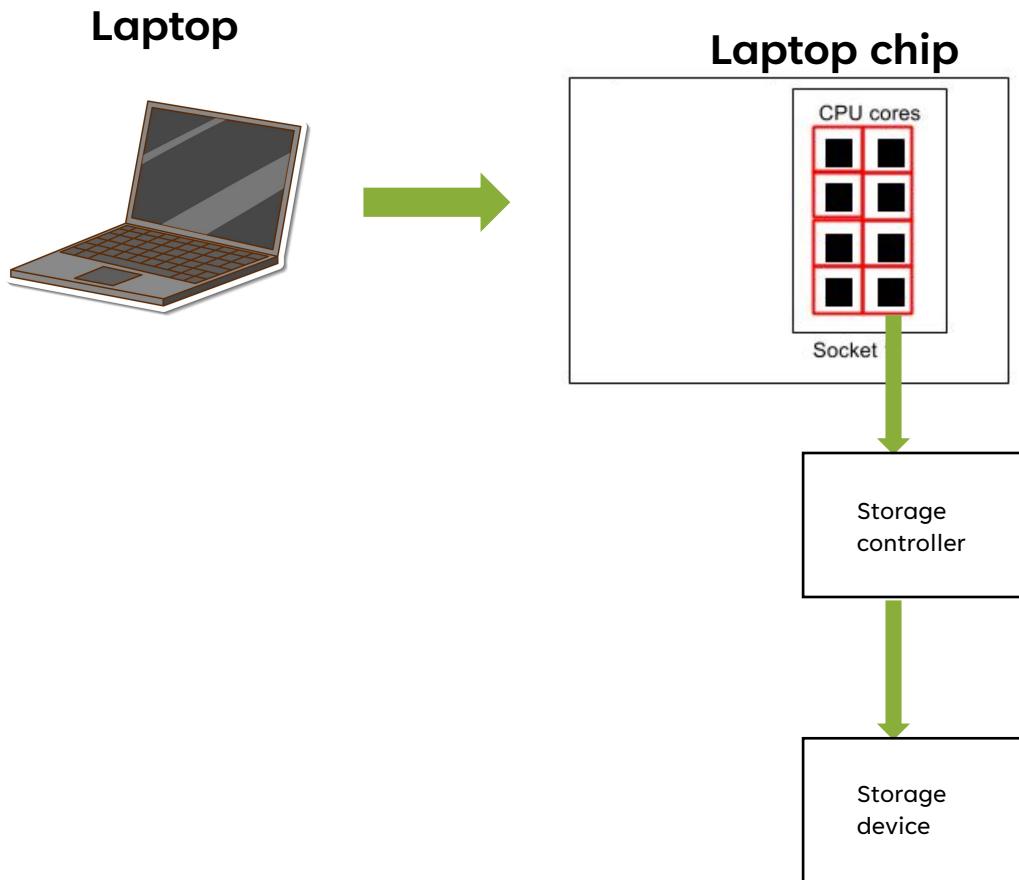
Node 1



Node 2



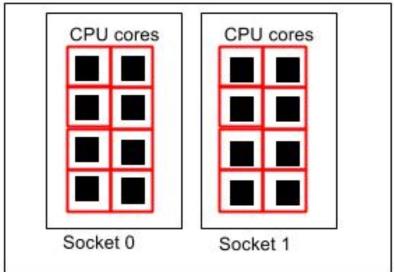
# Overview of storage [PC]



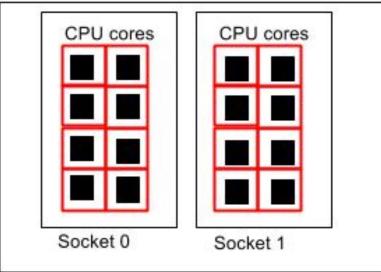
- Compute core talks to the storage controller which talks to the storage device to get the data
- Path between software to storage is short.
- Low latency and low bandwidth.

# Overview of storage [HPC]

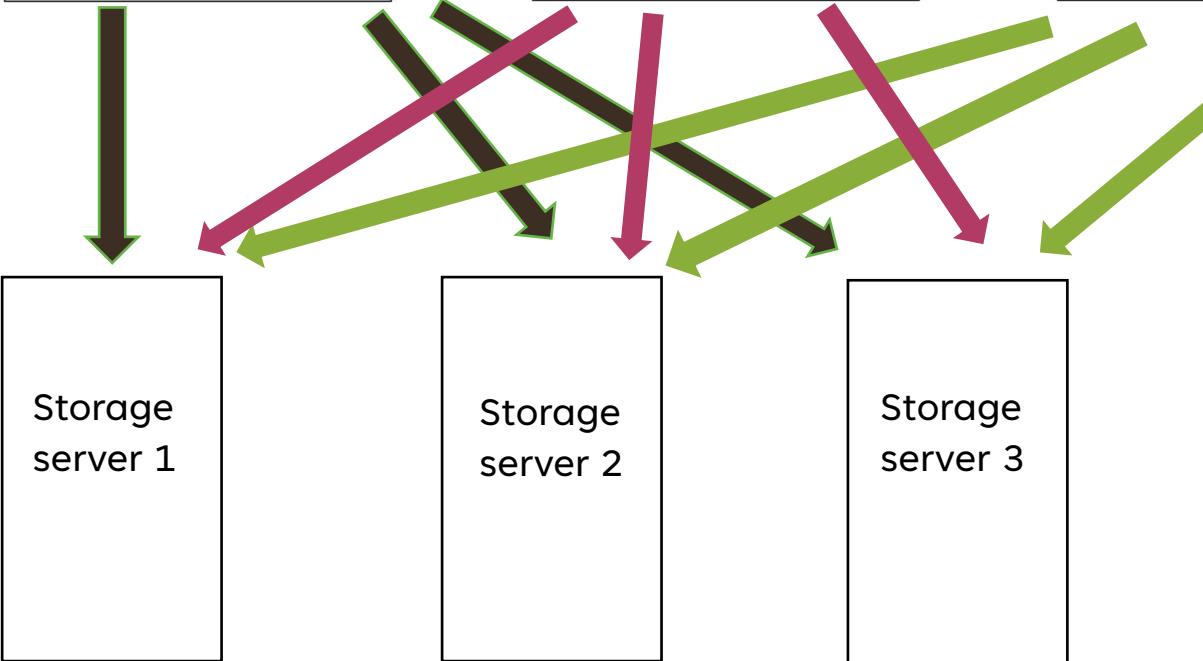
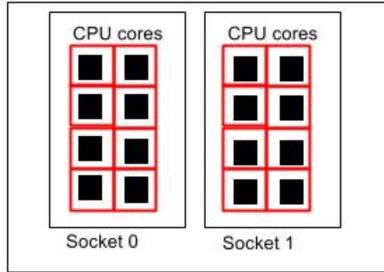
Node 1



Node 2



Node 3



- Files distributed across multiple servers.
- Much Higher bandwidth.
- In general higher latency.

# Filesystems on High Performance Computing

- Network filesystem (NFS):  
Allows users to access files and directories located on remote servers



# GPFS High performance filesystem

- Optimized for heavy I/O pipelines
- Optimized for heavy parallelized pipelines



# II- Part 2) Intro to Alpine



# Official Github pages:

- CU Anschutz HPC official Github page:

<https://github.com/kf-cuanschutz/CU-An documentation>.

- Note: You can submit pull requests as well.

The screenshot shows the GitHub repository page for kf-cuanschutz / CU-Anschutz-HPC-documentation. The repository is public and has the following details:

- Code: 63 commits
- Issues: 2
- Pull requests: 1
- Actions: 0
- Projects: 0
- Security: 0
- Insights: 0
- Settings: 0

The repository has 1 branch and 0 tags. The latest commit was made by kf-cuanschutz 2 weeks ago. The commits are listed below:

File	Description	Time Ago
Globus-local-entry-point-files	Delete Overview_of_Petalibrary.pdf	2 weeks ago
Office-hours-presentation-files	Introduction to Petalibrary	2 weeks ago
ALDEEx2-R-package-installation.md	Fixed the hyperlink that was pointing to globus	3 months ago
Alpine-cluster-maintenance.md	replaced Preventative maintenance with FAQ	3 months ago
Alpine-pipeline-opt-FAQ.md	Added Job arrays upper limit	2 months ago
EcholocateR.md	small typo fix	3 months ago
MATLAB-kernel-on-Jupyterlab.md	Create MATLAB-kernel-on-Jupyterlab.md	2 weeks ago
README.md	Update README.md	2 weeks ago
cellRangerRkit.md	Few changes added (e.g. indentation on Part 1)	2 months ago

The repository has 6 stars, 1 watching, and 0 forks. It also includes sections for About, Releases, and Packages.

- CU Boulder curc doc:

<https://curc.readthedocs.io/en/latest/access/logging-in.html>

- Tickets submitted by emailing: [rc-help@colorado.edu](mailto:rc-help@colorado.edu). You may email AMC HPC Support as well: [hpcsupport@cuanschutz.edu](mailto:hpcsupport@cuanschutz.edu).

# Logging into Alpine(1)

Alpine Office Hours\_.pdf  
422 KB

<User Name>,  
Congratulations on your new Alpine account! You are now ready to start using Alpine.

To help you navigate Alpine and get your workflows started quickly, we have compiled a few quick instructions and tips on how to get the most out of your experience. For a fully comprehensive overview of all documentation related to Alpine, please visit the following site: <https://curc.readthedocs.io/en/latest/index.html>. We recommend you bookmark this site and save this email message for future use.

**Please note, all functionality, commands, and directives can only be used after requesting an `acompile` node on the command line.** This can be done by typing `acompile` and pressing enter once you are logged into the terminal. From there, all SLURM and module commands will work as expected.

## Introduction to the SLURM job scheduler

Alpine utilizes the SLURM job scheduling software to manage job submissions to the compute nodes. The Boulder CURC team has provided tutorials and documentation for how to write your job scripts in order for the head node to allocate and assign your jobs with the appropriate resources you are requesting.

- Useful SLURM script examples: <https://curc.readthedocs.io/en/latest/clusters/alpine/examples.html>
- Commonly used SLURM commands on Alpine: <https://curc.readthedocs.io/en/latest/running-jobs/slurm-commands.html>
- SLURM flags and partitions on Alpine: <https://curc.readthedocs.io/en/latest/running-jobs/job-resources.html>
- How to write and submit a SLURM job script: <https://curc.readthedocs.io/en/latest/running-jobs/batch-jobs.html>
- Setting up an interactive job with SLURM: please use the `acompile` or `sinteractive` commands to spin up an interactive job; specify the `-h` flag to see all options for these commands.

All Anschutz users should specify `#SBATCH --account=amc-general` in their job submission scripts in order to get the Anschutz priority boost enabled on their jobs.

## Logging into Alpine

- For Anschutz users, you will need to login to Alpine by following the [set of instructions outlined here in this link](#) using the Open OnDemand portal.
- For now, directly accessing Alpine from the terminal using ssh is not yet enabled. We will notify you when this feature is available.

## Your directories on Alpine

The paths below use the built-in shell variable `$USER` to represent your username. You can access them directly using the paths shown below, since `$USER` will be replaced with your username by the shell. If you are unsure of your username, you can find it by typing `echo $USER` on the Alpine terminal. For Anschutz users, this may have the `@xsede.org` or `@access.org` extension.

`/home/$USER`

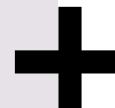
- 2 GB per user
- Not meant to be used for data storage or for reading and writing to/from compute nodes!
- It should be used for `.bashrc`, `.bash_profile`, `.ssh`, etc...

`/projects/$USER`

- 250 GB per user
- Used for very small data (<1 GB) and local software installs
- When compiling software, please use a compile node, not the head node. [For more information on nodes please visit node types here.](#)



- Once the user's EUA application is approved, they will receive an onboarding email similar to what is above.



# Logging into Alpine(2)

The screenshot shows an email from "Alpine Office Hours" with the subject "Alpine Office Hours\_pur". The email body contains the following text:

<User Name>,  
Congratulations on your new Alpine account! You are now ready to start using Alpine.

To help you navigate Alpine and get your workflows started quickly, we have compiled a few quick instructions and tips on how to get the most out of your experience. For a fully comprehensive overview of all documentation related to Alpine, please visit the following site: <https://curc.readthedocs.io/en/latest/index.html>. We recommend you bookmark this site and save this email message for future use.

**Please note, all functionality, commands, and directives can only be used after requesting an acompile node on the command line.** This can be done by typing acompile and pressing enter once you are logged into the terminal. From there, all SLURM and module commands will work as expected.

**Introduction to the SLURM job scheduler**  
Alpine utilizes the SLURM job scheduling software to manage job submissions to the compute nodes. The Boulder CURC team has provided tutorials and documentation for how to write your job scripts in order for the head node to allocate and assign your jobs with the appropriate resources you are requesting.

- Useful SLURM script examples: <https://curc.readthedocs.io/en/latest/clusters/alpine/examples.html>
- Commonly used SLURM commands on Alpine: <https://curc.readthedocs.io/en/latest/running-jobs.html>
- SLURM flags and partitions on Alpine: <https://curc.readthedocs.io/en/latest/running-jobs/job-reservation.html>
- How to write and submit a SLURM job script: <https://curc.readthedocs.io/en/latest/running-jobs/basic.html>
- Setting up an interactive job with SLURM: please use the acompile or sinteractive command options for these commands.

All Anschutz users should specify `#SBATCH --account=amc-general` in their job submission script.

**Logging into Alpine**

- For Anschutz users, you will need to login to Alpine by following the [set of instructions outlined here](#) using the Open OnDemand portal.
- For now, directly accessing Alpine from the terminal using ssh is not yet enabled. We will notify you when it is available.

**Your directories on Alpine**  
The paths below use the built-in shell variable \$USER to represent your username. You can access them directly using the paths shown below, since \$USER will be replaced with your username by the shell. If you are unsure of your username, you can find it by typing `echo $USER` on the Alpine terminal. For Anschutz users, this

A red oval highlights the link "set of instructions outlined here" in the "Logging into Alpine" section. A red arrow points from this oval to the right side of the slide, where a list of steps is provided.

- Then they will open the following link in listed in the email and follow the procedure.

# Logging into Alpine(3)

## Logging into Alpine from the Shell App

### Note

Make sure you already have your XSEDE/ACCESS user name and password set up before proceeding and Duo 2-factor authentication set up for your ACCESS/XSEDE account

1. Visit <https://ondemand-rmacc.rc.colorado.edu> You will be redirected to CILogon. From there, make sure you select the ACCESS CI (XSEDE) as your identity provider and then click the "Log On" button.

### CILogon

Consent to Attribute Release

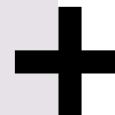
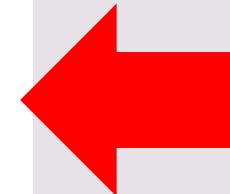
OOD RMACC requests access to the following information. If you do not approve this request, do not proceed.

- Your CILogon user identifier
- Your name
- Your email address
- Your username and affiliation from your identity provider

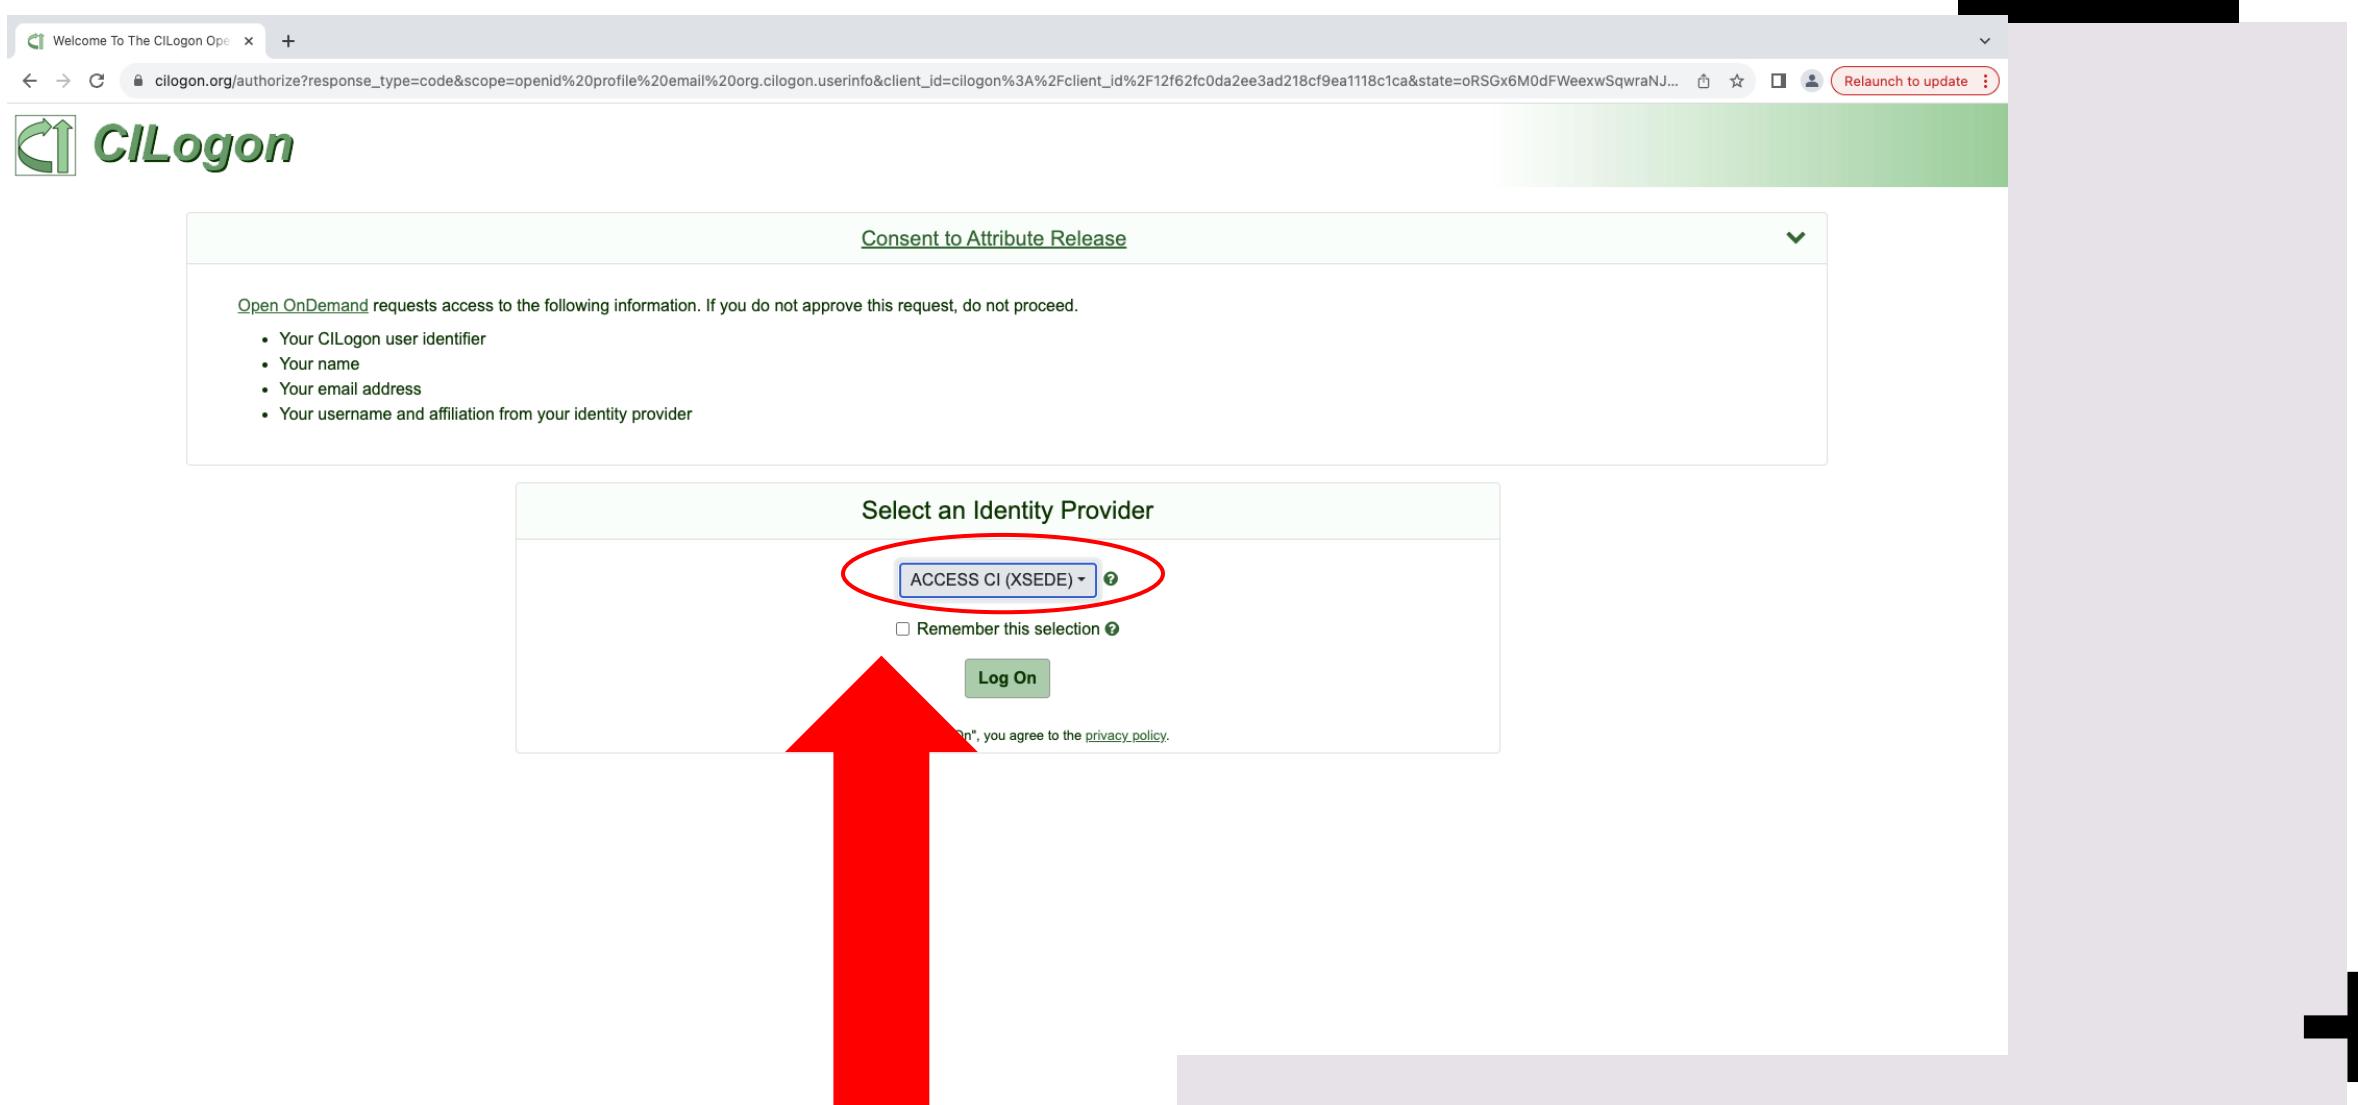
Select an Identity Provider

ACCESS CI (XSEDE) 

Remember this selection 

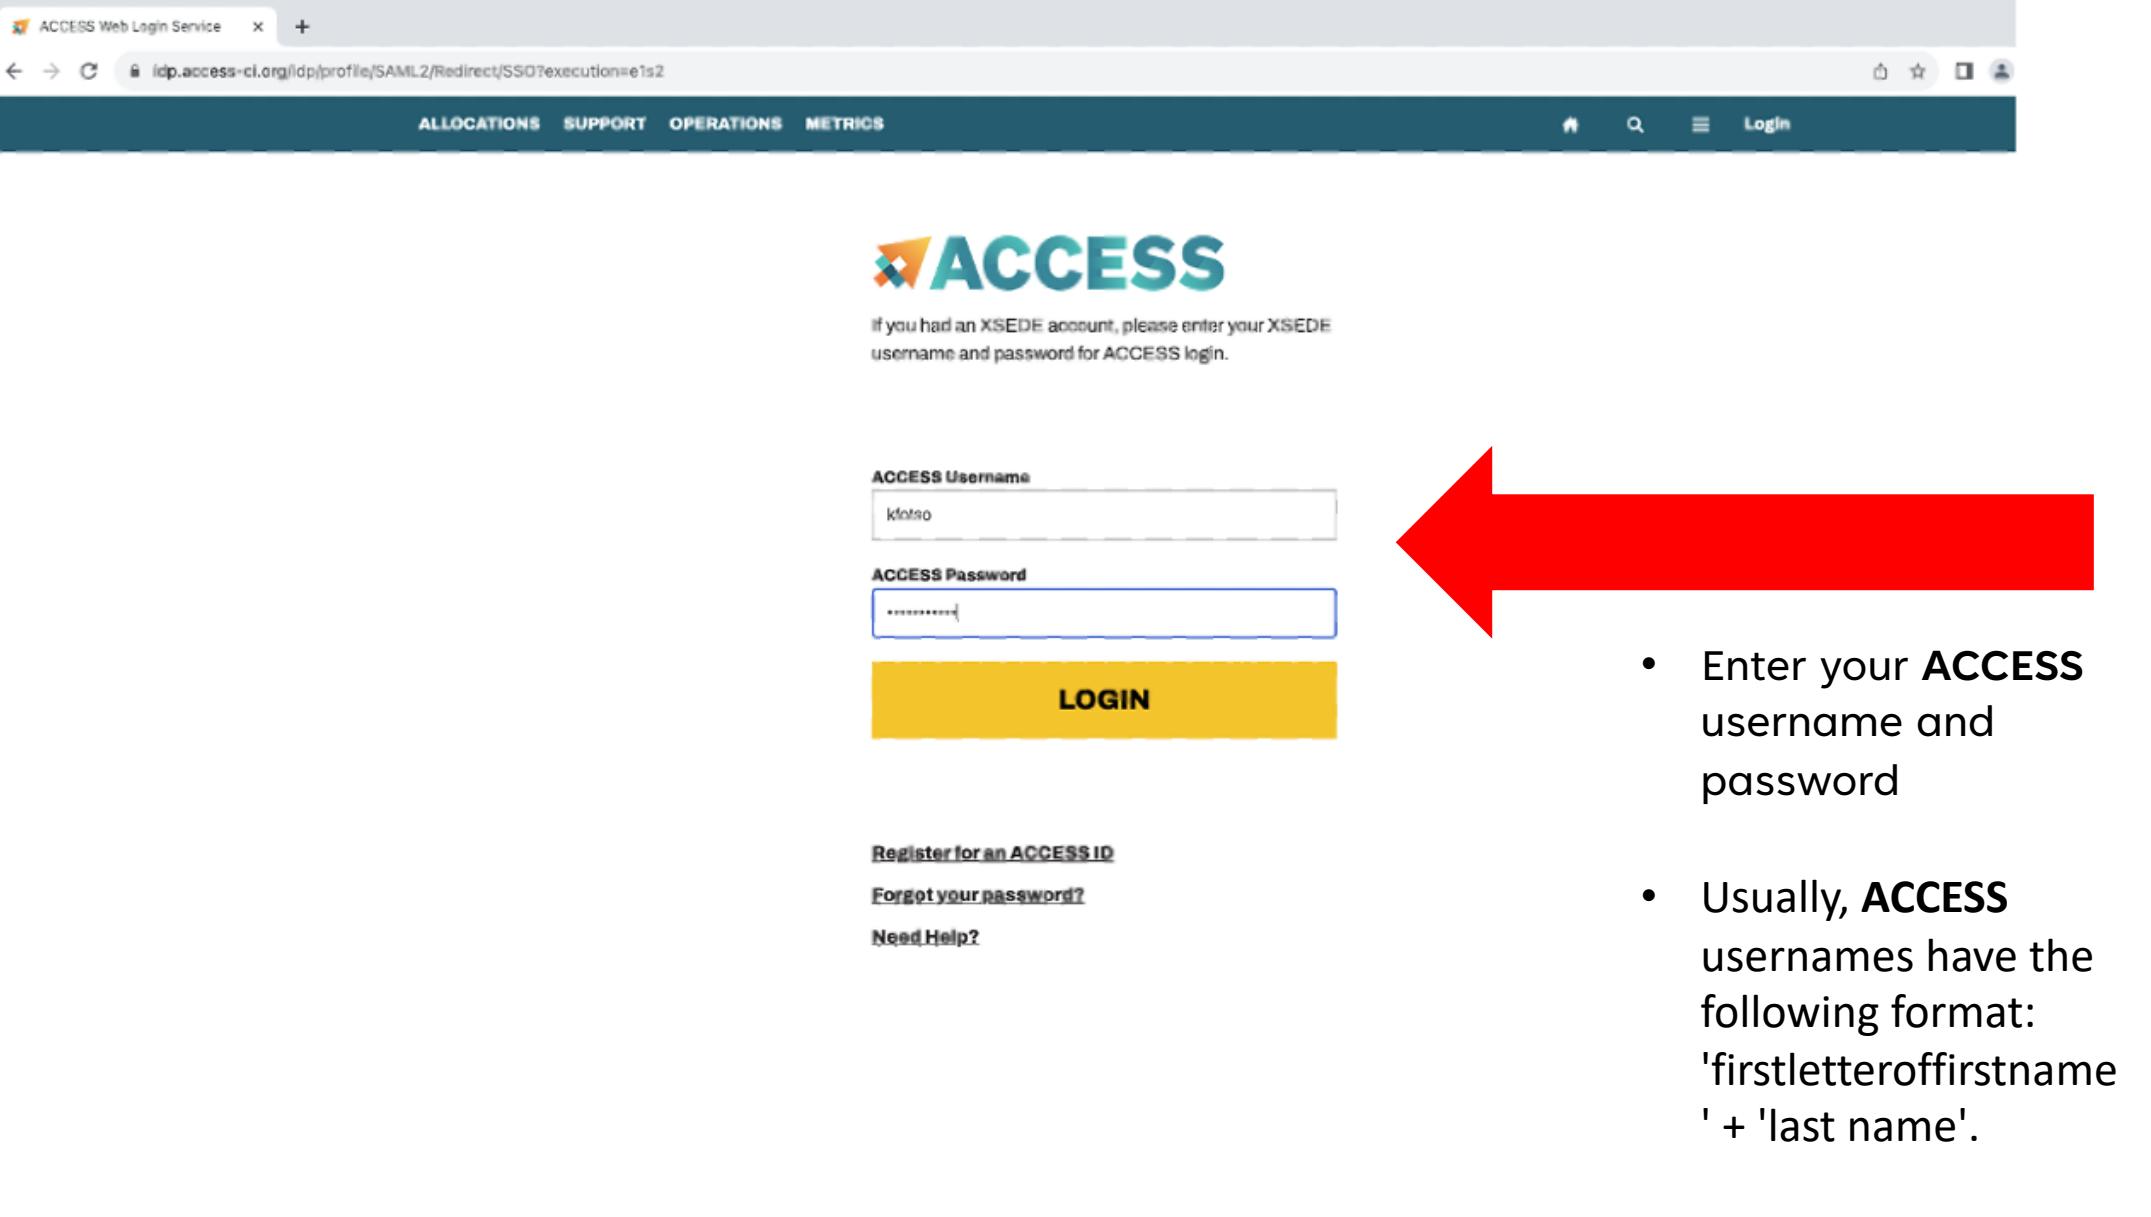


# Logging into Alpine(4)



- Make sure to select ACCESS CI (XSEDE) Identity Provider **and NOT CU Anschutz !!!**

# Logging into Alpine(5)



If you had an XSEDE account, please enter your XSEDE username and password for ACCESS login.

ACCESS Username

ACCESS Password

**LOGIN**

[Register for an ACCESS ID](#)  
[Forgot your password?](#)  
[Need Help?](#)

- Enter your **ACCESS** username and password
- Usually, **ACCESS** usernames have the following format:  
'firstletteroffirstname' + 'last name'.

# Logging into Alpine(6)

Duo Security

api-616d966d.duosecurity.com/frame/v4/auth/prompt?sid=frameless-db928ffc-8566-42c6-a72a-64e4f50995d5

ACCESS

Check for a Duo Push

Verify it's you by approving the notification...

Sent to "iOS" (\*\*\*\*\*-6519)

Other options

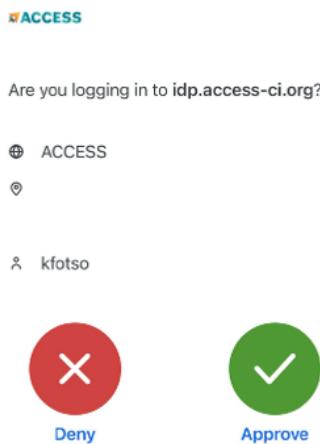
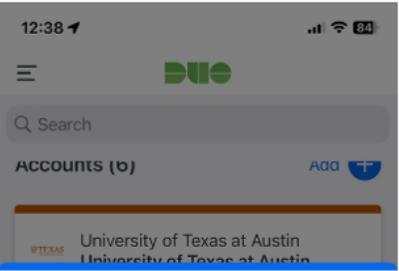
Need help?

Secured by Duo

A large red arrow points from the text on the right towards the 'Other options' link on the Duo push notification screen.

Make sure to check the DUO push notification on your phone.

# Logging into Alpine(7)



The authentication confirmation should look similar to the following on your phone.

# Logging into Alpine(2)

The screenshot shows the OnDemand web interface for the University of Colorado Research Computing system. At the top, there is a navigation bar with icons for Files, Jobs, Clusters (which is currently selected and highlighted in black), Interactive Apps, and My Interactive Sessions. A red circle highlights the 'Clusters' dropdown menu, which is open to show two options: '\_Alpine Shell' and '\_Blanca Shell'. Below the navigation bar is the CU Boulder Research Computing logo, featuring a gold 'CU' monogram and the text 'Research Computing UNIVERSITY OF COLORADO BOULDER'. A main message states: 'OnDemand provides an integrated, single access point for all of your HPC resources.' Below this is a 'Message of the Day' section with the text: 'Welcome to the University of Colorado Research Computing.' Underneath, there is a 'Quick Links' section with several blue hyperlinks: 'CU Boulder RC Status', 'Research Computing User Guide', 'Research Computing at CU Boulder', 'RMACC @ Ask.Cyberinfrastructure', and 'Need help? Email ([rc-help@colorado.edu](mailto:rc-help@colorado.edu))'. A large black plus sign icon is located in the bottom right corner of the page.

OnDemand provides an integrated, single access point for all of your HPC resources.

## Message of the Day

---

Welcome to the University of Colorado Research Computing.

### Quick Links

[CU Boulder RC Status](#)  
[Research Computing User Guide](#)  
[Research Computing at CU Boulder](#)  
[RMACC @ Ask.Cyberinfrastructure](#)  
Need help? Email ([rc-help@colorado.edu](mailto:rc-help@colorado.edu))

# On the login node (1)

```
Welcome to University of Colorado Boulder Research Computing!
```

```
Full documentation is available in our user guide at  
https://www.rc.colorado.edu/support/user-guide. If you have a question  
that's not answered there, contact us at rc-help@colorado.edu.
```

```
A number of directories have been created for you already:
```

```
* `/home/$USER`, your home directory  
* `/projects/$USER`, your project directory
```

```
Run the command `module avail` to see a list of available software.
```

```
To prevent this README from being displayed at login, edit your  
.bash_profile` or `.login` files.
```

```
bash: /usr/bin/python: No such file or directory
```

```
Welcome to CU-Boulder Research Computing.
```

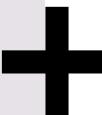
```
* Website http://colorado.edu/rc  
* Questions? rc-help@colorado.edu  
* Subscribe to system announcements: https://curc.statuspage.io/  
* Please type rc-help for the Acceptable Use Policy and a short help page.
```

```
You are using login node: login-ci1
```

```
bash: /usr/bin/python: No such file or directory  
[kfotso@xsede.org@login-ci1 ~]$ █
```



Name of the username



# On the login node (2)

```
Welcome to University of Colorado Boulder Research Computing!
```

```
Full documentation is available in our user guide at  
https://www.rc.colorado.edu/support/user-guide. If you have a question  
that's not answered there, contact us at rc-help@colorado.edu.
```

```
A number of directories have been created for you already:
```

```
* `/home/$USER`, your home directory  
* `/projects/$USER`, your project directory
```

```
Run the command `module avail` to see a list of available software.
```

```
To prevent this README from being displayed at login, edit your  
.bash_profile` or `.login` files.
```

```
bash: /usr/bin/python: No such file or directory
```

```
Welcome to CU-Boulder Research Computing.
```

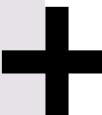
```
* Website http://colorado.edu/rc  
* Questions? rc-help@colorado.edu  
* Subscribe to system announcements: https://curc.statuspage.io/  
* Please type rc-help for the Acceptable Use Policy and a short help page.
```

```
You are using login node: login-ci1
```

```
bash: /usr/bin/python: No such file or directory  
[kfotso@xsede.org@login-ci1 ~]$
```



Name of the login node  
(hostname)



# Storage(1)

Home filesystem (2G). Backed up + for hosting config files.

Project filesystem (250G). Backed up -> for package installation

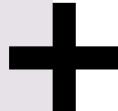
“cd .snapshot” to access those back ups

‘curc-quota’ or ‘du’ to check on space

# Storage(2)

```
[kfotso@xsede.org@login-ci1 ~]$ curc-quota
```

	Used	Avail	Quota	Limit
/home/kfotso@xsede.org	289M	1.8G	2.0G	
/projects/kfotso@xsede.org	144G	107G	250G	
/scratch/alpine1	7786G	20825G	28611G	



## Storage(3)

- Scratch space (10 TB)
- GPFS filesystem
- Very suitable for parallel application + heavy I/O
- Purged every 90 days

# Petalibrary

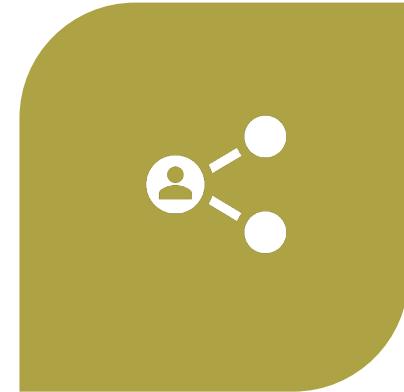
+

# Preliminary conditions (1)

---

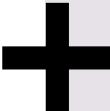


CREATION OF AN ACCESS  
GROUP

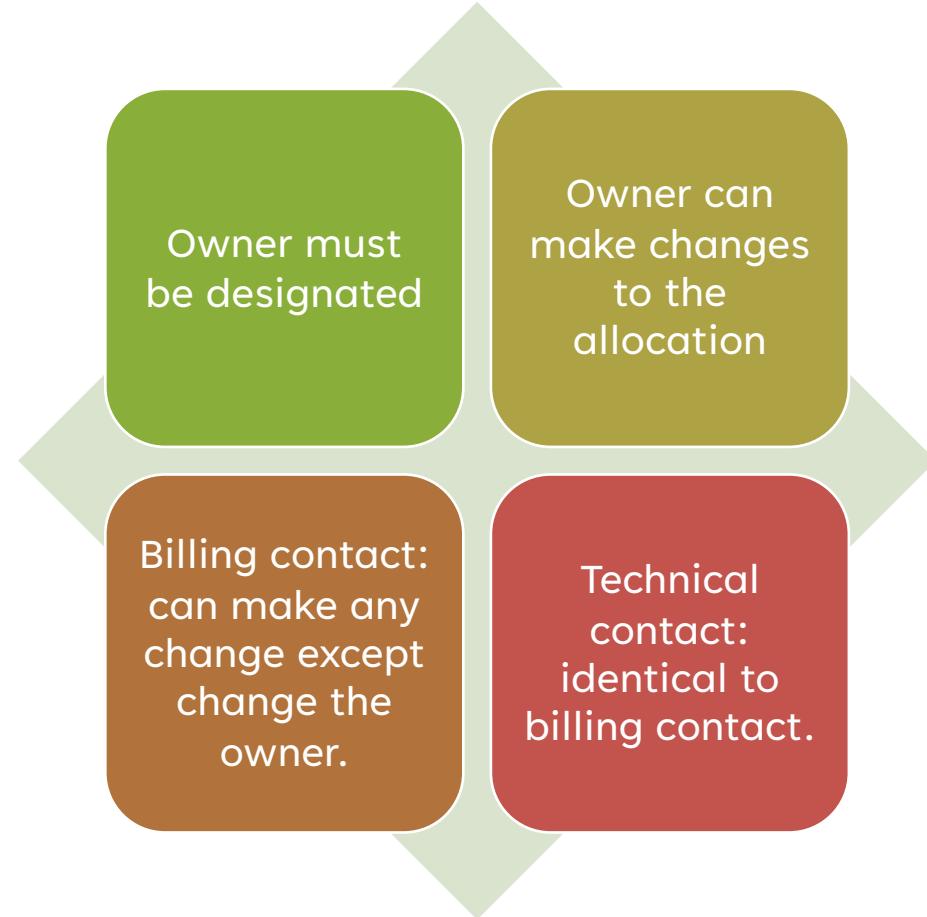


ALL MEMBER ADDED TO THE  
GROUP WILL NEED  
ACCESS/XSEDE ACCOUNTS.

/pl/active/<your\_allocation\_name>  
/pl/archive/<your\_allocation\_name>



# Preliminary conditions (2)





# Application

- Application submitted here:  
<https://www.colorado.edu/rcc/resources/petalibrary>
- The request form will need a speedtype: account# to which they plan to charge the allocation.

# Billing (1)

- On active storage: \$45/TB/yr.
- ZFS Raidz2 allow for frequent read/write + parity.
- **It is highly suggested that total TB size for the year is determined in advanced by the owner.**



## Billing (2)

- On archive storage: \$20/TB/yr.
- Tape-like storage for infrequently accessed data.
- Can be accessed from RC login node only and data transfer resources.
- Min size for any alloc is 1 TB



# Terms of usage

- Data will need to be in full compliance with term of service.
- No PHI data, no FERPA, no ITAR, no GDPR, no export control and data/software that comply with IRB requirements.
- More information here:  
<https://www.colorado.edu/rc/resources/petalibrary/tos>
- Data classification:  
<https://www.cu.edu/security/data-classification>.
- HIPAA identifiers:  
<https://www.dhcs.ca.gov/dataandstats/data/Pages>ListofHIPAAIdentifiers.aspx>



# Data redundancy (1)



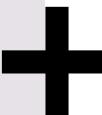
PL allocations are of single-copy nature.



All users should fill out the PL single copy acknowledgment.



Snapshots monitoring in place so that they are not missed unless on snapshot custom schedule



# Sort term back up solutions (1)

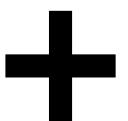
- Microsoft OneDrive: 5TB per person through A3 licensing.
- S3 buckets: customer paid & rates located here:  
<https://aws.amazon.com/s3/pricing/?p=pm&c=s3&z=4>
- Multiple PL allocations as though customer requested 2 allocations (X2 price).



# Sort+mid term back up solutions (2)

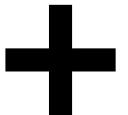
---

- [BETA]  
Replicated PetaLibrary active  
+archive allocation.
- Nothing is charged during  
beta period but will likely look  
like Multiple PL allocations.
- CURC **only takes data  
replication responsibility** for  
BETA



# Active+Archive replication

- Synchronization between active and archive happen every 15 min in theory but are not always guaranteed.
- Snapshots are also replicated from active to archive.
- The Boulder storage team maintains the second copy (e.g. archive) for you and you do not have direct access to it.
- \$65/TB/year total. **(NOT \$65 on top of the \$45/\$20 original rate!!)**





1

Creation of Globus Connect  
Personal endpoint.



2

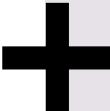
CU Boulder endpoint for  
Anschutz was CU Boulder  
Research Computing  
ACCESS.



3

Current ACCESS point  
(DTN23) is 10X faster than  
the legacy one  
“CU Boulder Research  
Computing”

# Files transfer (Globus)



# Globus (2)

Collections

Get Globus Connect Personal

CU Boulder Research Computing

X

QUICK FILTERS  RECENTLY USED  ADMINISTERED BY YOU  IN USE  SHAREABLE BY YOU  SHARED WITH YOU

COLLECTION	HA	MANAGED	STATUS	ROLE
<a href="#">CU Boulder Research Computing (internal - Legacy - Do Not Use) Managed GCSv4 Host</a>			requires activation	
<a href="#">CU Boulder Research Computing - OLD Managed GCSv4 Host</a>			requires activation	
<a href="#">CU Boulder Research Computing (Legacy - Do Not Use) Managed GCSv4 Host</a>			requires activation	
<a href="#">CU Boulder Research Computing Managed Mapped Collection (GCS) on CU Boulder Research Computing DTN23</a>			ready	

+

# Globus(3)

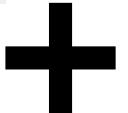
- More info here:

<https://curc.readthedocs.io/en/latest/compute/data>

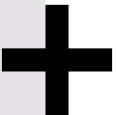
[transfer.html?highlight=globus#globus-transfers](#)

The screenshot shows a sidebar menu for "Research Computing University of Colorado Boulder". The sidebar includes sections for "Frequently Asked Questions", "ACCESSING RC RESOURCES" (with links to Logging In, Duo 2-factor Authentication, RMACC Access to Alpine, and SSH Key-Based Authentication for Anschutz Medical Campus), "THE COMPUTE ENVIRONMENT" (with links to Node types, Filesystems, and The Modules System), and "Data Transfer" (which is expanded to show "Globus transfers", "Guest Collections (Globus Shared Endpoints)", "Globus to AWS S3 Bucket connection", "Filezilla", "Secure Copy (scp)", "Rsync", and "Interactive file transfer with sftp").

The screenshot shows the main content area of the CURC Data Transfer page. It features a heading "Data Transfer", a paragraph about supported methods (Globus, scp, sftp, rsync), and a section about CURC data transfer nodes (DTN). Below this is a video player with a play button and a "Share" icon. The video player has a circular icon with a letter "C" and the text "Data Transfers". At the bottom, there is a "Watch on YouTube" button and the "Be Boulder." logo.



II) - Alpine computing



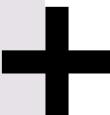
# Scheduler Slurm



**acompile --ntasks=1 --time=00:30:00** to build packages and do some testing. (max 12 hours)



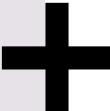
**sinteractive --ntasks-per-node=2 --nodes=2 --partition=atesting** to test pipelines. (max 1 hour)



```
[kfotso@xsede.org@login-ci1 kfotso@xsede.org]$ sinfo --partition=aa100
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
aa100          up 1-00:00:00      1  resv c3gpu-a9-u33-1
aa100          up 1-00:00:00      7  mix  c3gpu-a9-u31-1,c3gpu-a9-u35-1,c3gpu-c2-u
aa100          up 1-00:00:00      4  alloc c3gpu-a9-u29-1,c3gpu-c2-u[7,13,15]
```

# sinfo

Can be used to get  
information about a node.



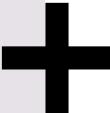
# Package availability (1)

Some packages that have been built and accessible through Imod.

Adding new packages through Imod takes a lot of round of approval so it is recommended to build them locally.

Solutions: (cmake+make), Anaconda, pip, containers, spack etc ...

Submit a ticket at [rc-help@Colorado.edu](mailto:rc-help@Colorado.edu) so that I can build it for you locally.



# LMOD (CONTEXT)

- Modern environment module system for HPC
- Initially introduced by Robert McLay from TACC (UT Austin) in 2011.



# LMOD (CONTEXT)

- Users have on HPC have different needs.
- Applications, compilers, libraries, versions being used might be different.



# LMOD (benefits)

- Users do not need to know where software is installed
- Environment variable to interface packages can be set (e.g. picard).
- Very useful for software reproducibility



# LMOD (benefits)

- Supports software hierarchy
- Users can set their environment modules at will.



# List of officially available bio modules

- “acompile”
- “module avail”

`alphafold/2.2.0`  
`alphafold/2.3.1 (D)`  
`bamtools/2.5.2`  
`bbtools/39.01`

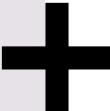
`bcftools/1.16`  
`bedtools/2.29.1`  
`bowtie2/2.5.0`  
`bwa/0.7.17`

`cellranger/7.1.0`  
`cutadapt/4.2`  
`fastqc/0.11.9`  
`gatk/4.3.0.0`

`htslib/1.16`  
`multiqc/1.14`  
`nextflow/22.10.6`  
`nextflow/23.04 (D)`

`picard/2.27.5`  
`plink2/2.00a2.3`  
`qiime2/2023.5`  
`samtools/1.16.1`

`sra-toolkit/3.0.0`  
`star/2.7.10b`  
`trimmmomatic/0.39`

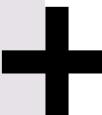


# Slurm example

```
#!/bin/bash ← Your shell is bash  
  
#SBATCH --partition=amilan ← Partition is amilan (CPU)  
#SBATCH --job-name=example-job ← Job name when it gets queued  
#SBATCH --output=example-job.%j.out ← Output file name  
#SBATCH --time=01:00:00 ← How long my job runs?  
#SBATCH --qos=normal  
#SBATCH --nodes=1  
#SBATCH --ntasks=4  
#SBATCH --mail-type=ALL  
#SBATCH --mail-user=youridentikey@colorado.edu
```

```
module purge  
module load anaconda  
conda activate custom-env
```

```
pythonmyscript.py
```



# Slurm example

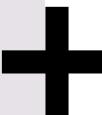
```
#!/bin/bash

#SBATCH --partition=amilan
#SBATCH --job-name=example-job
#SBATCH --output=example-job.%j.out
#SBATCH --time=01:00:00
#SBATCH --qos=normal
#SBATCH --nodes=1
#SBATCH --ntasks=4
#SBATCH --mail-type=ALL
#SBATCH --mail-user=youridentikey@colorado.edu
```

Quality of service  
How many nodes I want to run on?

```
module purge
module load anaconda
conda activate custom-env
```

```
python myscript.py
```



# Slurm example

```
#!/bin/bash

#SBATCH --partition=amilan
#SBATCH --job-name=example-job
#SBATCH --output=example-job.%j.out
#SBATCH --time=01:00:00
#SBATCH --qos=normal
#SBATCH --nodes=1
#SBATCH --ntasks=4
#SBATCH --mail-type=ALL
#SBATCH --mail-user=youridentikey@colorado.edu

module purge
module load anaconda
conda activate custom-env

python myscript.py
```

Total cores used

Email when job begins, ends and fails

Email  
address

# Slurm example

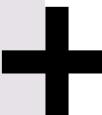
```
#!/bin/bash

#SBATCH --partition=amilan
#SBATCH --job-name=example-job
#SBATCH --output=example-job.%j.out
#SBATCH --time=01:00:00
#SBATCH --qos=normal
#SBATCH --nodes=1
#SBATCH --ntasks=4
#SBATCH --mail-type=ALL
#SBATCH --mail-user=youridentikey@colorado.edu
```

module purge ← Reinitialize module list  
module load anaconda ← Load anaconda  
conda activate custom-env ← Activate environment  
python myscript.py ← Running pipeline command

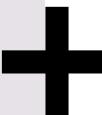
# Slurm cheatsheet (1)

Slurm script command	Description
<code>#!/bin/bash</code>	Sets the shell that the job will be executed on the compute node
<code>#SBATCH --ntasks=1</code> <code>#SBATCH --n1</code>	Requests for 1 processors on task, usually 1 cpu as 1 cpu per task is default.
<code>#SBATCH --time=0-05:00</code> <code>#SBATCH -t 0-05:00</code>	Sets the maximum runtime of 5 hours for your job
<code>#SBATCH --mail-user= &lt;email&gt;</code>	Sets the email address for sending notifications about your job state.
<code>#SBATCH --mail-type=BEGIN</code> <code>#SBATCH --mail-type=END</code> <code>#SBATCH --mail-type=FAIL</code> <code>#SBATCH --mail-type=REQUEUE</code> <code>#SBATCH --mail-type=ALL</code>	Sets the scheduling system to send you email when the job enters the following states: BEGIN,END,FAIL,REQUEUE,ALL
<code>#SBATCH --job-name=my-named-job</code>	Sets the Jobs name



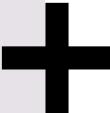
# Slurm cheatsheet(2)

Slurm script command	Description
#SBATCH --ntasks=X	Requests for X tasks. When cpus-per-task=1 (and this is the default) this requests X cores. When not otherwise constraint these CPUs may be running on any node
#SBATCH --nodes=X	Request that a minimum of X nodes be allocated to this job
#SBATCH --nodes=X-Y	Request that a minimum of X nodes and a maximum of Y nodes be allocated to this job
#SBATCH --cpus-per-task=X	Request that a minimum of X CPUs per task be allocated to this job
#SBATCH --tasks-per-node=X	Requests minimum of X task be allocated per node



# Slurm cheatsheet(3)

Slurm script commands	Description of effects
#SBATCH --ntasks=1 #SBATCH --cpus-per-task=1	Requests 1 CPU (Serial) cpus-per-task is set to 1 by default and may be omitted.
#SBATCH --cpus-per-task=X #SBATCH --ntasks=1 #SBATCH --nodes=1	Requests for X CPUs in 1 task on 1 node (OpenMP) Both ntasks and nodes are set to 1 by default and may be omitted
#SBATCH --ntasks=X #SBATCH --tasks-per-node=X #SBATCH --cpus-per-task=1	Requests for X CPUs and tasks on 1 node cpus-per-task is set to 1 by default and may be omitted.
#SBATCH --ntasks=X #SBATCH --nodes=1 #SBATCH --cpus-per-task=1	Requests for X CPUs and tasks on 1 node cpus-per-task is set to 1 by default and may be omitted.



# Get information about jobs

```
[kfotso@xsede.org@login-ci1 ~]$ squeue -l --me
```

JOBID	PARTITION	NAME	USER	STATE
2158225	acompile	acompile	kfotso@x	RUNNING

TIME	TIME_LIMI	NODES	NODELIST(REASON)
0:16	3:00	1	c3cpu-c11-u21-2

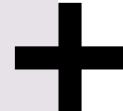
# Monitor resources

```
[kfotso@xsede.org@login-ci1 ~]$ module load slurmtools
[kfotso@xsede.org@login-ci1 ~]$ jobstats $USER 2
job stats for user kfotso@xsede.org over past 2 days
jobid      jobname  partition   qos       account    cpus state start-date-time elapsed    wait
-----
2064187    sinterac atesting_+ testing    amc-gener+  48  TIMEOUT 2023-06-20T23:32:52 01:00:04  0 hrs
2071952    vep_loft aa100        normal    amc-gener+  64  COMPLETE 2023-06-21T13:47:55 00:26:42  3 hrs
```

- Allows to get information about a past jobs

```
[kfotso@xsede.org@login-ci1 ~]$ seff 1451164
Job ID: 1451164
Cluster: alpine
User/Group: kfotso@xsede.org/kfotsopgrp@xsede.org
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 48
CPU Utilized: 26-03:21:39
CPU Efficiency: 94.06% of 27-19:00:48 core-walltime
Job Wall-clock time: 13:53:46
Memory Utilized: 412.77 GB
Memory Efficiency: 41.28% of 999.98 GB
```

- To get more computational information about the job efficiency



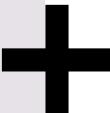
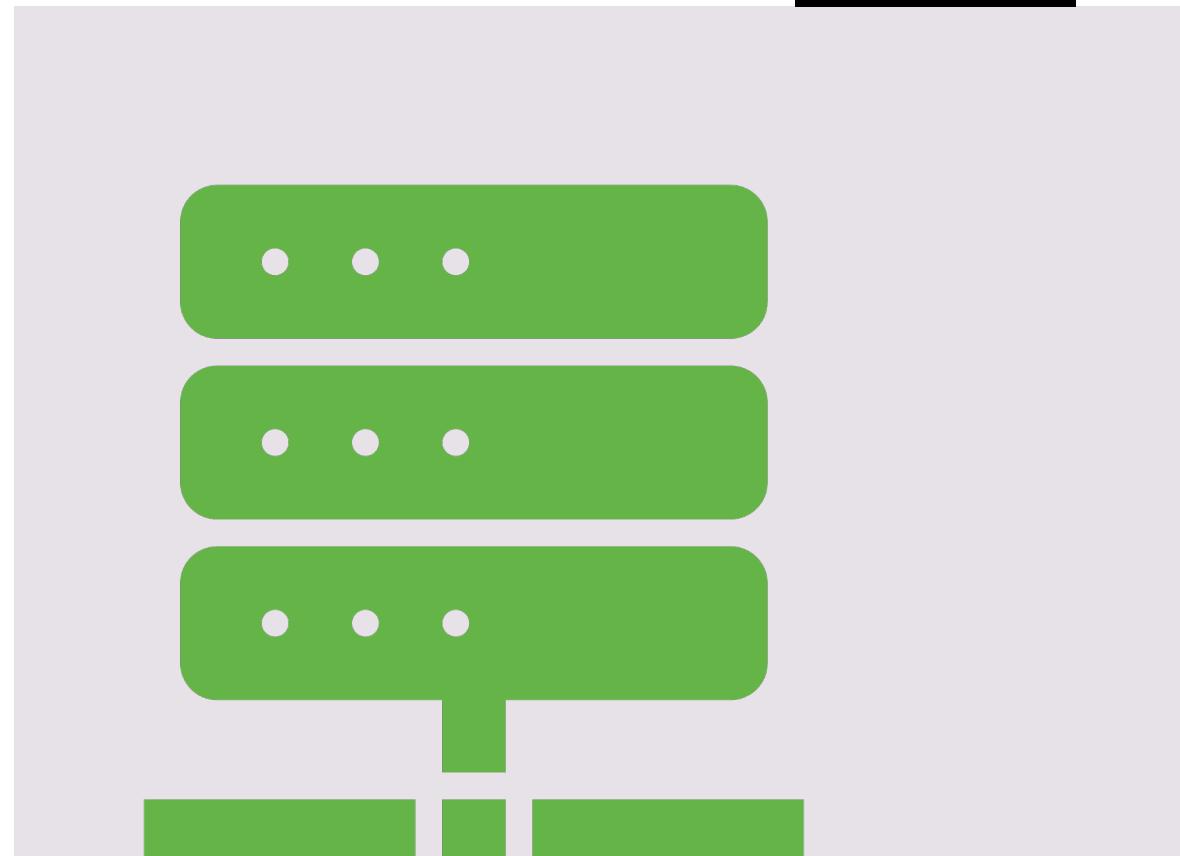
# Slurm Quality of service (qos)

- Used to modify or constrain characteristics that a job can have.
- **--qos=normal** corresponds to a walltime of 24 hours and is the default.
- **--qos=long** corresponds to a walltime of up to 7 days
- **--qos=mem** corresponds to high memory jobs only (up to 1TB)



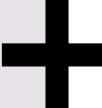
# Fairshare overview

- Difference between the portion of computing resource that has been promised and the amount of resources that has been consumed.
- Level fairshare of 1 indicates average priority compared to other users in that account (amc-general)
- **module load slurmtools; levels \$USER**



# Job priority calculation formula

```
Job_priority =  
    site_factor +  
    (PriorityWeightAge) * (age_factor) +  
    (PriorityWeightAssoc) * (assoc_factor) +  
    (PriorityWeightFairshare) * (fair-share_factor) +  
    (PriorityWeightJobSize) * (job_size_factor) +  
    (PriorityWeightPartition) * (partition_factor) +  
    (PriorityWeightQOS) * (QOS_factor) +  
    SUM(TRES_weight_cpu * TRES_factor_cpu,  
        TRES_weight_<type> * TRES_factor_<type>,  
        ...)  
    - nice_factor
```

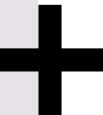


# Check fairshare

Host: login-ci1.rc.int.colorado.edu

```
[kfotso@xsede.org@login-ci1 ~]$ levelfs $USER
LevelFS for user kfotso@xsede.org and institution amc:
Account          LevelFS_User      LevelFS_Inst
-----
amc-general      0.194275        4.750220
[kfotso@xsede.org@login-ci1 ~]$
```

- 0.19 means that my priority will be low
- On the other hand 4.75 means that priority for the institution is high



# Service Units (SU)

- It is the number of core hours used.

```
[kfotso@xsede.org@login-ci1 ~]$ suuser $USER 10
SU used by user kfotso@xsede.org in the last 10 days:
Cluster|Account|Login|Proper Name|TRES Name|Used|
alpine|amc-general|kfotso@xsede.org|Kevin Fotso|billing|8393|
```

- suacct to get the number of core hours used by institution

```
Host: login-ci1.rc.int.colorado.edu
[kfotso@xsede.org@login-ci1 ~]$ suacct amc-general 180
SU used by account (allocation) amc-general in the last 180 days:
Cluster|Account|Login|Proper Name|TRES Name|Used
alpine|amc-general|||billing|1806360
alpine| amc-general|acozart@xsede.org|Abigail Cozart|billing|573
alpine| amc-general|agillen@xsede.org|Austin Gillen|billing|40320
alpine| amc-general|agray@xsede.org|Alyx Gray|billing|22
```

# Best practices on Alpine (1)

---

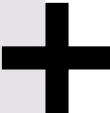
- /home or /tmp have a very small size. Please always include the following in your slurm script.

```
export TMP=/scratch/alpine/$USER
```

```
export TEMP=/scratch/alpine/$USER
```

```
export TMPDIR=/scratch/alpine/$USER
```

```
export TEMPDIR=/scratch/alpine/$USER
```



# Best practices on Alpine (2)

---

- /home or /tmp have a very small size. Please always include the following in your slurm script.

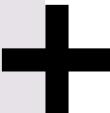
```
export TMP=/scratch/alpine/$USER
```

```
export TEMP=/scratch/alpine/$USER
```

```
export TMPDIR=/scratch/alpine/$USER
```

```
export TEMPDIR=/scratch/alpine/$USER
```

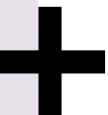
- In general, always make sure to check your .cache in \$HOME to make sure you are not running out of space.



# Best practices on Alpine (3) - Username Aliasing

---

- `/home/foo@xsede.org` → `/home/.xsede.org/foo/`
- `/projects/foo@xsede.org` → `/projects/.xsede.org/foo/`
- `/scratch/alpine/foo@xsede.org` → `/scratch/alpine/.xsede.org/foo/`



# Best practices on Alpine (4) – Anaconda config

---

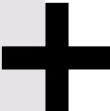
- Edit the .condarc file located in \$HOME/.condarc with the editor of choice (e.g. nano, vim etc ...)
- Paste the following 4 lines in the file and exit.

**pkgs\_dirs:**

- [/projects/.xsede.org/foo/.conda\\_pkgs](/projects/.xsede.org/foo/.conda_pkgs)

**envs\_dirs:**

- </projects/.xsede.org/foo/software/anaconda/envs>



# Brief mention of parallelism

- OpenMP (multithreading).
- MPI (Parallel distribution between multiple nodes)
- GNU parallel (embarrassingly parallel)
- Job arrays (max 1000 on Alpine)
- Load balancing (developed by CU Boulder)
- GPU computing



# Slurm example to see the number of cores(1)

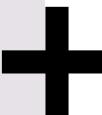
- Here, we are requesting 4 cores for a 2 minutes



```
[kfotso@login-ci2 ~]$ acompile --ntasks=4 --time=00:02:00
acompile: submitting job... salloc --nodes=1 --partition=acompile --ntasks=4 --time=00:02:00 --qos=compile --job-name=acompile --bell
--oversubscribe srun --pty /bin/bash
salloc: Granted job allocation 2998293
salloc: Nodes c3cpu-a5-u32-1 are ready for job
```



- We are getting node c3cpu-a5-u32-1 and our jobID is 2998293



# Slurm example to see the number of cores(2)

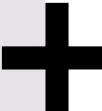
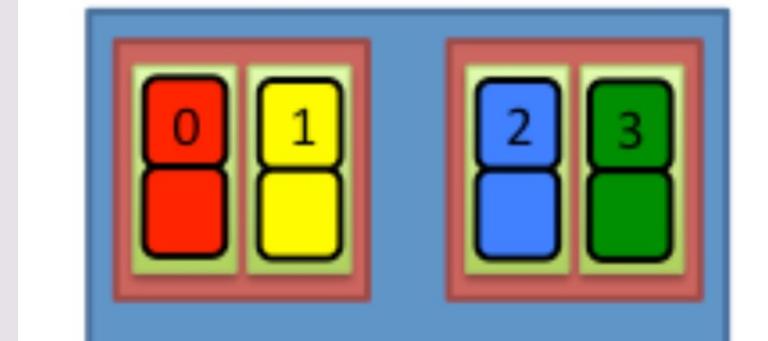
- We want load compiler gcc because openmpi depends on it



```
[kfotso@xsede.org@c3cpu-a5-u32-1 ~]$ module load gcc  
[kfotso@xsede.org@c3cpu-a5-u32-1 ~]$ module load openmpi
```



- We want to use openmpi because we want to demonstrate parallelism on the 4 cores we have reserved



# Slurm example to see the number of cores(3)

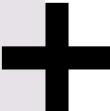
- hostname is just a program that prints the name of the node I am working with



```
[kfotso@xsede.org@c3cpu-a5-u32-1 ~]$ mpirun -n 4 hostname  
c3cpu-a5-u32-1.rc.int.colorado.edu  
c3cpu-a5-u32-1.rc.int.colorado.edu  
c3cpu-a5-u32-1.rc.int.colorado.edu  
c3cpu-a5-u32-1.rc.int.colorado.edu
```



- With mpirun, we schedule 4 tasks on the 4 cores, thus the name of the node gets printed 4 times !!!



# Slurm example to see the number of cores(4)

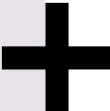
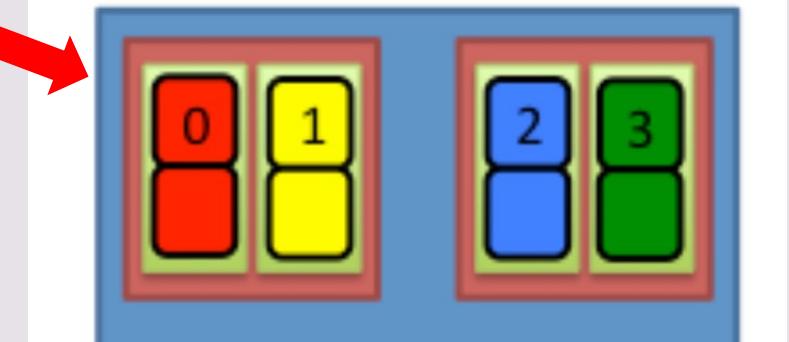
Recap of what we did

`mpirun -n 4 hostname`



tasks

0	hostname
1	hostname
2	hostname
3	hostname



# Slurm example to see the number of cores(4)

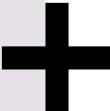
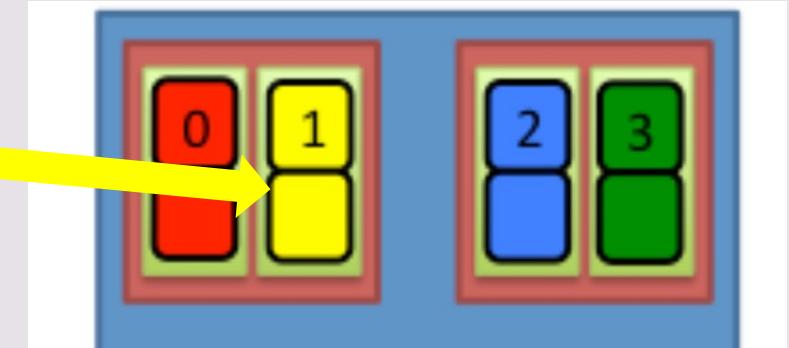
Recap of what we did

`mpirun -n 4 hostname`



tasks

0	hostname
1	hostname
2	hostname
3	hostname



# Slurm example to see the number of cores(4)

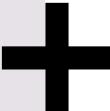
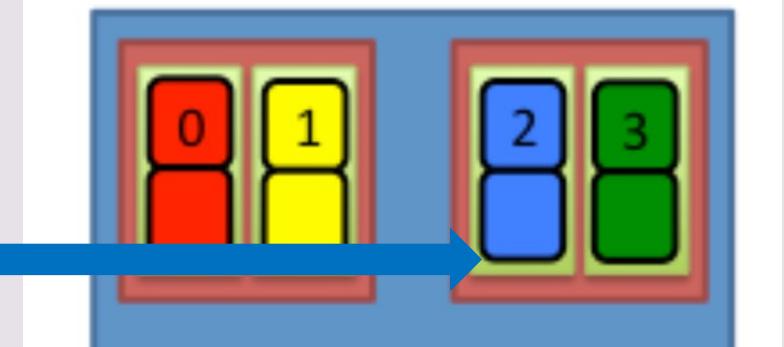
Recap of what we did

`mpirun -n 4 hostname`



tasks

0	hostname
1	hostname
2	hostname
3	hostname



# Slurm example to see the number of cores(4)

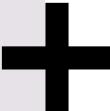
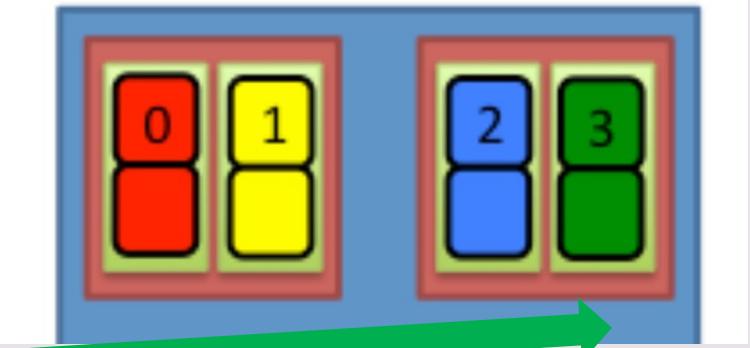
Recap of what we did

`mpirun -n 4 hostname`



tasks

0	hostname
1	hostname
2	hostname
3	hostname



# Slurm example to see the number of cores(4)

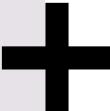
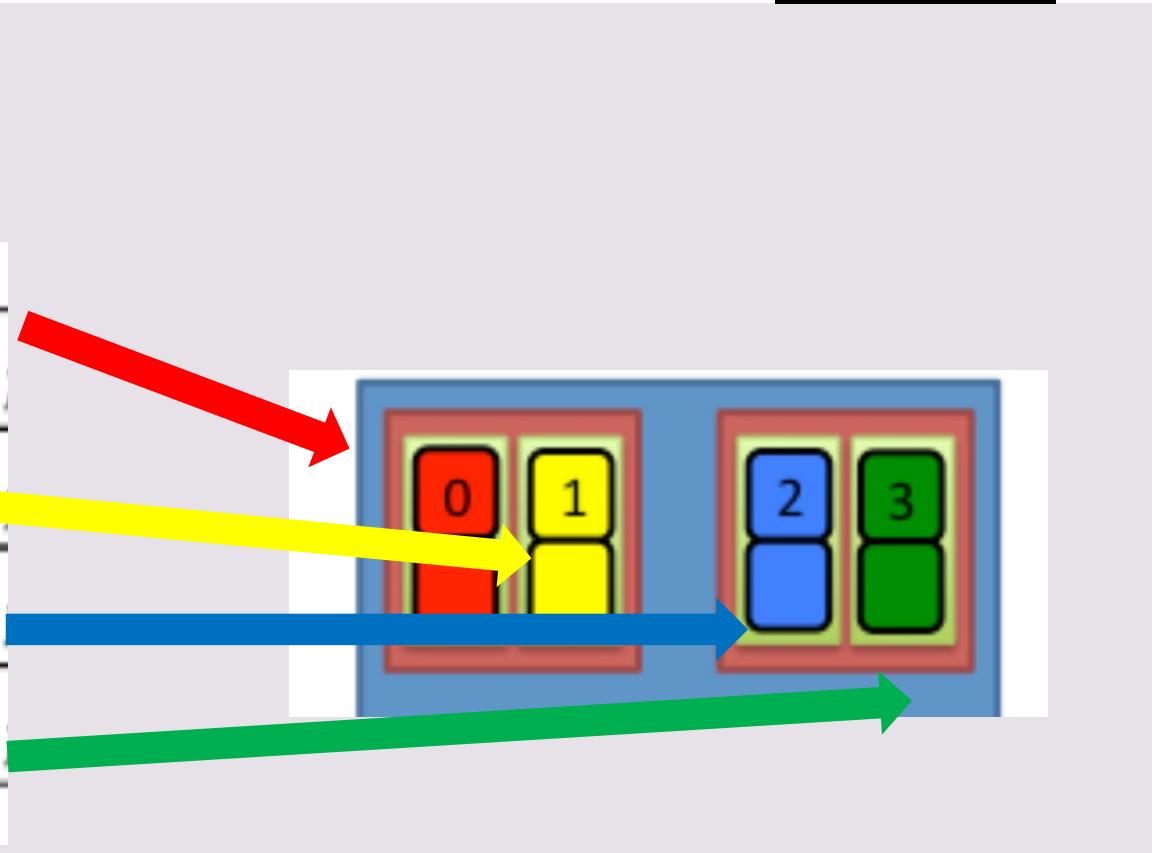
Recap of what we did

`mpirun -n 4 hostname`



tasks

0	hostname
1	hostname
2	hostname
3	hostname



# Slurm example to see the number of cores(5)

- Congratulations!!! You just wrote your first HPC parallel code program!

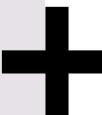


# Build the corresponding slurm script (1)

- Create the empty slurm script and modify it with your editor of choice (nano, vim, emacs, Ondemand etc ...)



```
[kfotso@xsede.org@login-ci2 openmpi_hostname]$ touch hello_mpi.sh  
[kfotso@xsede.org@login-ci2 openmpi_hostname]$
```



# Build the corresponding slurm script

```
#!/bin/bash

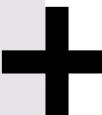
#SBATCH --partition=amilan
#SBATCH --job-name=first_mpirun
#SBATCH --output=first_mpirun-job.%j.out
#SBATCH --error=first_mpirun-job.%j.err
#SBATCH --account=amc-general
#SBATCH --qos=normal
#SBATCH --nodes=1
#SBATCH --ntasks=4
#SBATCH --mail-type=ALL
#SBATCH --mail-user=kfotsotagne@unm.edu
#SBATCH --time=00:00:01

module load gcc
module load openmpi

mpirun -n $SLURM_NTASKS hostname
```

Account name is  
amc-general

Job runs for 1  
second!!



# Build the corresponding slurm script

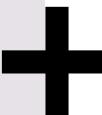
```
#!/bin/bash

#SBATCH --partition=amilan
#SBATCH --job-name=first_mpirun
#SBATCH --output=first_mpirun-job.%j.out
#SBATCH --error=first_mpirun-job.%j.err
#SBATCH --account=amc-general
#SBATCH --qos=normal
#SBATCH --nodes=1
#SBATCH --ntasks=4
#SBATCH --mail-type=ALL
#SBATCH --mail-user=kfotsotagne@unm.edu
#SBATCH --time=00:00:01

module load gcc
module load openmpi

mpirun -n $SLURM_NTASKS hostname
```

\$SLURM\_NTASKS is the slurm  
env variable for the 4 cores

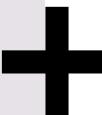


# Build the corresponding slurm script (4)

We use sbatch to submit the script



```
[kfotso@xsede.org@login-ci2 openmpi_hostname]$ sbatch hello_mpi.sh
Submitted batch job 2998325
```



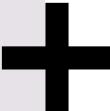
# Build the corresponding slurm script

(5)

The .out file shows the name of the nodes 4X.  
.err file is empty so no error.



```
[kfotso@xsede.org@login-ci2 openmpi_hostname]$ ls -latr
total 184
drwxrws---. 44 kfotso@xsede.org kfotsogrp@xsede.org 1672 Sep 15 01:45 ..
-rw-r--r--. 1 kfotso@xsede.org kfotsogrp@xsede.org 400 Sep 15 01:52 hello_mpi.sh
-rw-r--r--. 1 kfotso@xsede.org kfotsogrp@xsede.org 0 Sep 15 01:52 first_mpirun-job.2998325.err
drwxr-sr-x. 2 kfotso@xsede.org kfotsogrp@xsede.org 122 Sep 15 01:52 .
-rw-r--r--. 1 kfotso@xsede.org kfotsogrp@xsede.org 140 Sep 15 01:52 first_mpirun-job.2998325.out
[kfotso@xsede.org@login-ci2 openmpi_hostname]$ cat first_mpirun-job.2998325.out
c3cpu-c15-u1-2.rc.int.colorado.edu
c3cpu-c15-u1-2.rc.int.colorado.edu
c3cpu-c15-u1-2.rc.int.colorado.edu
c3cpu-c15-u1-2.rc.int.colorado.edu
[kfotso@xsede.org@login-ci2 openmpi_hostname]$
```



# Ondemand MATLAB

RC will conduct its monthly planned maintenance on Wednesday, September 6. Please visit [curc.statuspage.io](https://curc.statuspage.io) for details.

Notice: Users will be limited to a maximum of 8 cores per Core Desktop session through mid-September due to ongoing maintenance.

[Home](#) / [My Interactive Sessions](#) / [MATLAB \(Presets\)](#)

Interactive Apps

- Desktops
- Core Desktop (Presets)
- GUIs
- MATLAB (Presets)

Servers

- Jupyter Session (Custom)
- Jupyter Session (Presets)
- RStudio Server (Custom)
- RStudio Server (Presets)
- VS Code-Server (Custom)
- VS Code-Server (Presets)

## MATLAB (Presets)

This app will launch a MATLAB GUI on a CURC node. You will be able to interact with MATLAB through a VNC session. GPU based options are not meant for computationally intensive workflows. Additionally, please keep in mind that these GPU based options are a shared resource amongst all users. Thus, significant computation by one user can affect other users of this service.

MATLAB version

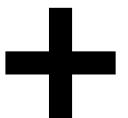
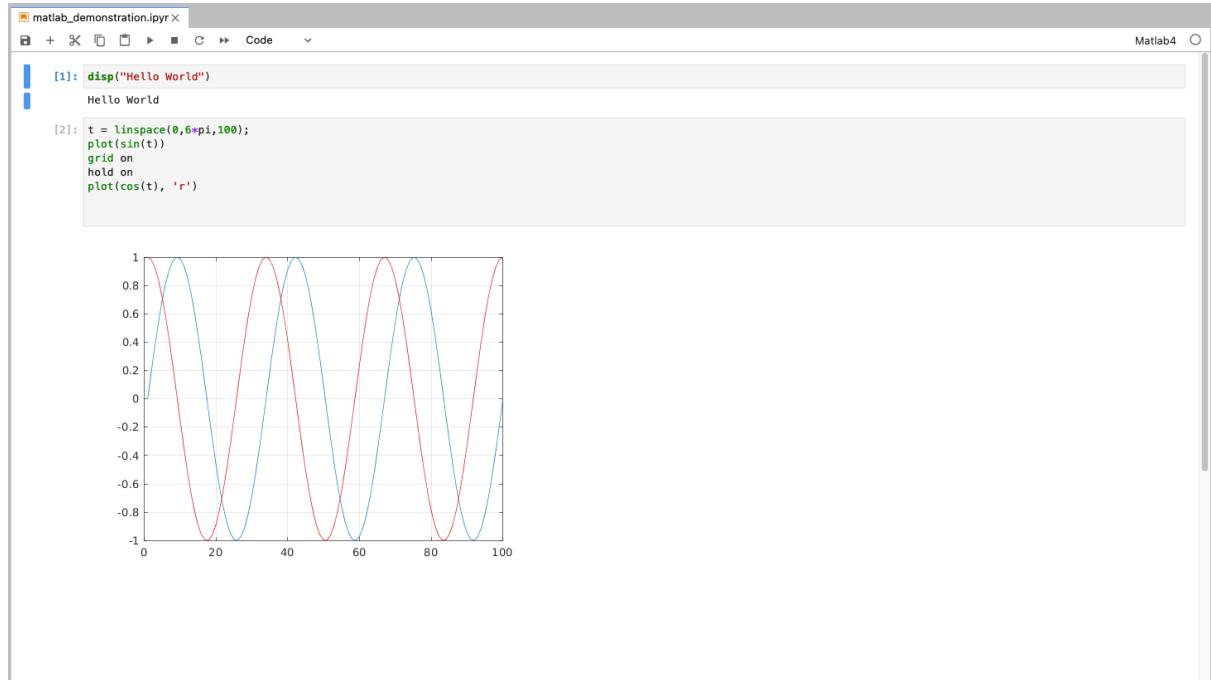
R2021b

Configuration

- 2 cores, 1 hour, K80 GPU
- 2 cores, 1 hour, RTX8000 GPU
- 2 cores, 12 hours, K80 GPU
- 2 cores, 12 hours, RTX8000 GPU
- 4 cores, 4 hours, K80 GPU
- 4 cores, 4 hours, RTX8000 GPU

# Ondemand Jupyterlab with MATLAB kernel

- Instructions here:  
<https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/blob/main/MATLAB-kernel-on-Jupyterlab.md>
- Possibility to use MATLAB with ssh (if registered)
- MATLAB Ondemand coming soon.



# Rstudio with Ondemand (preset)

Interactive Apps

Desktops

Core Desktop (Presets)

GUIs

MATLAB (Presets)

Servers

Jupyter Session (Custom)

Jupyter Session (Presets)

RStudio Server (Custom)

**RStudio Server (Presets)**

VS Code-Server (Custom)

VS Code-Server (Presets)

**RStudio Server (Presets)**

## RStudio Server (Presets)

This app will launch RStudio Server, an IDE for R on Alpine.

Before utilizing this application, please see the [RStudio section of the CURC documentation](#). This documentation includes important information regarding quitting an RStudio session. For more information on installing dependencies required by R packages, please see the [Installing dependencies for RStudio](#) section in our documentation.

Please note that the first time you launch RStudio it may take 15 minutes or more to create a [persistent overlay](#), if 1 core is used. For this reason, we recommend using 4 cores or more during your first launch of an RStudio session. Subsequent sessions will not require the creation of the persistent overlay; therefore 1 core may be sufficient.

RStudio Version

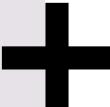
Rstudio 2023.03.0, R 4.2.2

Configuration

1 core, 1 hour, 0 GPUs (Alpine)

Launch

**Important:** It already interact with the Slurm scheduler!!!



# Rstudio with Ondemand (preset)

Interactive Apps

Desktops

Core Desktop (Presets)

GUIs

MATLAB (Presets)

Servers

Jupyter Session (Custom)

Jupyter Session (Presets)

RStudio Server (Custom)

**RStudio Server (Presets)**

VS Code-Server (Custom)

VS Code-Server (Presets)

**RStudio Server (Presets)**

RStudio Version

Rstudio 2023.03.0, R 4.2.2

Configuration

1 core, 1 hour, 0 GPUs (Alpine)

Launch

\* The RStudio Server (Presets) session data for this

## RStudio Server (Presets)

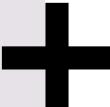
This app will launch [RStudio Server](#), an IDE for [R](#) on Alpine.

Before utilizing this application, please see the [RStudio section of the CURC documentation](#). This documentation includes important information regarding quitting an RStudio session. For more information on installing dependencies required by R packages, please see the [Installing dependencies for RStudio](#) section in our documentation.

Please note that the first time you launch RStudio it may take 15 minutes or more to create a [persistent overlay](#), if 1 core is used. For this reason, we recommend using 4 cores or more during your first launch of an RStudio session. Subsequent sessions will not require the creation of the persistent overlay; therefore 1 core may be sufficient.

**Important:** It already interact with the Slurm scheduler!!!

Preset option comes with already set parameters (e.g. 1 core or 4 cores)



# Rstudio with Ondemand (Custom)

of the persistent overlay; therefore 1 core may be sufficient.

RStudio Version

Rstudio 2023.03.0, R 4.2.2

Cluster

Alpine

Account

amc-general

Partition

amilan

Number of cores

64

Memory [GiB]

240

QoS Name

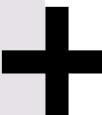
normal

Time

24



**64 cores max, 240G total  
(~3.8G per core), walltime  
24 hours**





# Rstudio on Ondemand

- We have a guide that show how to use dependencies in Rstudio:  
[https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio\\_related\\_scripts](https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio_related_scripts).



# Rstudio on Ondemand

- We have a guide that show how to use dependencies in Rstudio:  
[https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio\\_related\\_scripts](https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio_related_scripts).
- Use of apptainer, fakeroot and container overlay to modify the container on Alpine.



# Rstudio on Ondemand

- We have a guide that show how to use dependencies in Rstudio:  
[https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio\\_related\\_scripts](https://github.com/kf-cuanschutz/CU-Anschutz-HPC-documentation/tree/main/Rstudio_related_scripts).
- Use of apptainer, fakeroot and container overlay to modify the container on Alpine.
- 2 scripts that need to run

# Rstudio container

le

Blame

Executable File · 4 lines (3 loc) · 273 Bytes

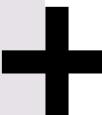
Raw



```
1 export APPTAINER_TMPDIR=/projects/$USER/.rstudioserver
2 export APPTAINER_CACHEDIR=/projects/$USER/.rstudioserver
3 aptainer shell --fakeroot --overlay /projects/$USER/.rstudioserver/rstudio-server-4.2.2_overlay.img /curc/sw/containers/open
```



We point the  
TMPDIR and CACHEDIR to  
either  
[`/scratch/alpine/\$USER`](#) or  
[`/projects/\$USER`](#) so that  
we do not fill /tmp



# Rstudio container

Blame Executable File · 4 lines (3 loc) · 273 Bytes

Raw   

```
1 export APPTAINER_TMPDIR=/projects/$USER/.rstudioserver
2 export APPTAINER_CACHEDIR=/projects/$USER/.rstudioserver
3 apttainer shell --fakeroot --overlay /projects/$USER/.rstudioserver/rstudio-server-4.2.2_overlay.img /curc/sw/containers/open
```



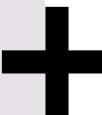
The overlay with .img  
Allows us to modify the .sif  
container image

# Rstudio container

```
1 apt-get -y update
2 apt install -y zlib1g-dev
3 apt install -y libglpk-dev
4 apt install -y libfftw3-3
5 apt install -y libcairo2-dev
6 apt install -y libxt-dev
7 apt install -y liblzma-dev
8 apt install -y libharfbuzz-dev
9 apt install -y libtiff-dev
10 apt install -y libpng-dev
11 apt install -y libfribidi-dev
12 apt install -y libcurl4-openssl-dev
13 apt install -y libcurl4-nss-dev
14 apt install -y libcurl4-gnutls-dev
15 apt install -y libfreetype-dev
16 apt install -y libbz2-dev
17 apt install -y libhdf5-dev
```



We install the necessary dependencies inside the container



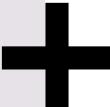
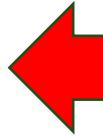
# Run Rstudio container in slurm script

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=24:00:00
#SBATCH --qos=normal
#SBATCH --partition=amilan
#SBATCH --ntasks=64
#SBATCH --account=amc-general
#SBATCH --job-name=R_job
#SBATCH --output=Rjob.%j.out
#SBATCH --error=Rjob.%j.err

# Exporting the apptainer tmp directories so that
# we do not fill /tmp
export ALPINE_SCRATCH=/gpfs/alpine1/scratch/$USER
export APPTAINER_TMPDIR=$ALPINE_SCRATCH/singularity/tmp
export APPTAINER_CACHEDIR=$ALPINE_SCRATCH/singularity/cache
mkdir -pv $APPTAINER_CACHEDIR $APPTAINER_TMPD

# Calling the Rstudio container to execute the Rscript
apptainer exec --bind /projects/$USER,/scratch/alpine/$USER --overlay /projects/
$USER/.rstudioserver/rstudio-server-4.2.2_overlay.img:ro /curc/sw/containers/
open_ondemand/rstudio-server-4.2.2.sif Rscript myscript.R
```

The slurm header looks similar to what we have seen before.



# Run Rstudio container in slurm script

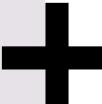
```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=24:00:00
#SBATCH --qos=normal
#SBATCH --partition=amilan
#SBATCH --ntasks=64
#SBATCH --account=amc-general
#SBATCH --job-name=R_job
#SBATCH --output=Rjob.%j.out
#SBATCH --error=Rjob.%j.err

# Exporting the apptainer tmp directories so that
# we do not fill /tmp
export ALPINE_SCRATCH=/gpfs/alpine1/scratch/$USER
export APPTAINER_TMPDIR=$ALPINE_SCRATCH/singularity/tmp
export APPTAINER_CACHEDIR=$ALPINE_SCRATCH/singularity/cache
mkdir -pv $APPTAINER_CACHEDIR $APPTAINER_TMPD

# Calling the Rstudio container to execute the Rscript
apptainer exec --bind /projects/$USER,/scratch/alpine/$USER --overlay /projects/
$USER/.rstudioserver/rstudio-server-4.2.2_overlay.img:ro /curc/sw/containers/
open_ondemand/rstudio-server-4.2.2.sif Rscript myscript.R
```



Export tmp related  
directories to the scratch.



# Run Rstudio container in slurm script

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=24:00:00
#SBATCH --qos=normal
#SBATCH --partition=amilan
#SBATCH --ntasks=64
#SBATCH --account=amc-general
#SBATCH --job-name=R_job
#SBATCH --output=Rjob.%j.out
#SBATCH --error=Rjob.%j.err

# Exporting the apptainer tmp directories so that
# we do not fill /tmp
export ALPINE_SCRATCH=/gpfs/alpine1/scratch/$USER
export APPTAINER_TMPDIR=$ALPINE_SCRATCH/singularity/tmp
export APPTAINER_CACHEDIR=$ALPINE_SCRATCH/singularity/cache
mkdir -pv $APPTAINER_CACHEDIR $APPTAINER_TMPD

# Calling the Rstudio container to execute the Rscript
apptainer exec --bind /projects/$USER,/scratch/alpine/$USER --overlay /projects/
$USER/.rstudioserver/rstudio-server-4.2.2_overlay.img:ro /curc/sw/containers/
open_ondemand/rstudio-server-4.2.2.sif Rscript myscript.R
```

Execute the container with  
R script.



# Core Desktop

Home / My Interactive Sessions / Core Desktop (Presets)

Interactive Apps
Desktops
<b>Core Desktop (Presets)</b>
GUIs
MATLAB (Presets)
Servers
Jupyter Session (Custom)
Jupyter Session (Presets)
RStudio Server

## Core Desktop (Presets)

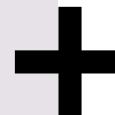
This app will launch an interactive desktop on a compute node. GPU based options are not meant for computationally intensive workflows. Additionally, please keep in mind that these GPU based options are a shared resource.

- 2 cores, 1 hour, K80 GPU
- 2 cores, 1 hour, RTX8000 GPU
- 2 cores, 12 hours, K80 GPU
- 2 cores, 12 hours, RTX8000 GPU
- ✓ 4 cores, 4 hours, K80 GPU**
- 4 cores, 4 hours, RTX8000 GPU

Launch

\* The Core Desktop (Presets) session data for this session can be accessed under the [data root directory](#).

Allows you to interact with a Desktop like env



# Core Desktop

Home / My Interactive Sessions / Core Desktop (Presets)

## Interactive Apps

Desktops

Core Desktop  
(Presets)

GUIs

MATLAB (Presets)

Servers

Jupyter Session  
(Custom)

Jupyter Session  
(Presets)

RStudio Server

## Core Desktop (Presets)

This app will launch an interactive desktop on a compute node. GPU based options are not meant for computationally intensive workflows. Additionally, please keep in mind that these GPU based options are a shared

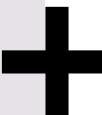
- 2 cores, 1 hour, K80 GPU
- 2 cores, 1 hour, RTX8000 GPU
- 2 cores, 12 hours, K80 GPU
- 2 cores, 12 hours, RTX8000 GPU
- ✓ 4 cores, 4 hours, K80 GPU**
- 4 cores, 4 hours, RTX8000 GPU

Launch



2 types of GPUs: RTX8000 and K80

\* The Core Desktop (Presets) session data for this session can be accessed under the [data root directory](#).



# Core Desktop

Home / My Interactive Sessions / Core Desktop (Presets)

## Interactive Apps

Desktops

Core Desktop  
(Presets)

GUIs

MATLAB (Presets)

Servers

Jupyter Session  
(Custom)

Jupyter Session  
(Presets)

RStudio Server

## Core Desktop (Presets)

This app will launch an interactive desktop on a compute node. GPU based options are not meant for computationally intensive workflows. Additionally, please keep in mind that these GPU based options are a shared

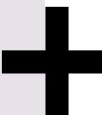
- 2 cores, 1 hour, K80 GPU
- 2 cores, 1 hour, RTX8000 GPU
- 2 cores, 12 hours, K80 GPU
- 2 cores, 12 hours, RTX8000 GPU
- ✓ 4 cores, 4 hours, K80 GPU**
- 4 cores, 4 hours, RTX8000 GPU

Launch



Max walltime 12 hours.

\* The Core Desktop (Presets) session data for this session can be accessed under the [data root directory](#).



# VS Code on Ondemand-rmacc

Home / My Interactive Sessions / VS Code-Server (Custom)

Interactive Apps
Desktops
Core Desktop (Presets)
GUIs
MATLAB (Presets)
Servers
Jupyter Session (Custom)
Jupyter Session (Presets)
RStudio Server (Custom)
RStudio Server (Presets)
<b>VS Code-Server (Custom)</b>
VS Code-Server (Presets)

## VS Code-Server (Custom)

This app will launch a [VS Code](#) server using [Code-Server](#). For more information on possible settings for this application, see [Running Custom Interactive applications](#) in our documentation. Additionally, for more information on installing VS Code extensions, please see our [Installing VS Code-Server Extensions](#) section of the documentation.

### Cluster

Alpine

### Code-Server version

✓ 4.16.1

4.14.1

### Account

amc-general

### Partition

ahub

### QoS Name

interactive

### Time

4

### Number of cores

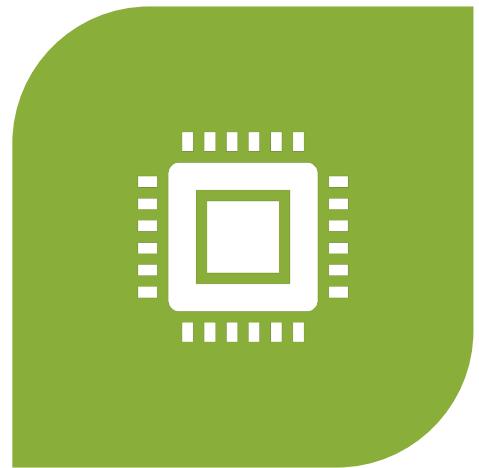
1

Launch

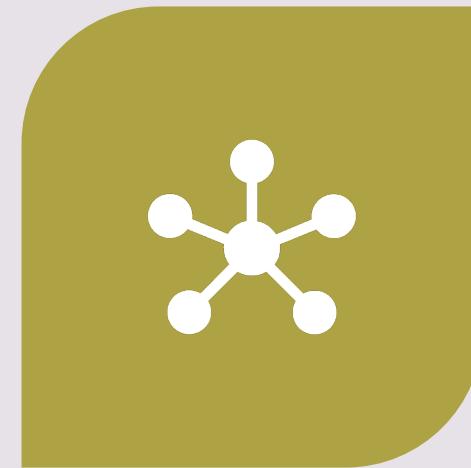
\* The VS Code-Server (Custom) session data for this session can be accessed under the [data root directory](#).

# GPU nodes on Alpine

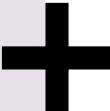
---



12 NVIDIA GPU NODES  
(3 GPU EACH)



8 AMD GPU NODES (3  
GPU EACH)



# GPU partition on Alpine

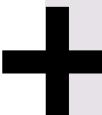
aa100 -> NVIDIA  
RELATED PARTITION;  
MAX WALLTIME 24  
HOURS

ami100 -> AMD  
RELATED PARTITION;  
MAX WALLTIME 24  
HOURS

atesting\_a100 ->  
NVIDIA TESTING  
PARTITION; MAX  
WALLTIME 1 HOUR

atesting\_mi100 ->  
AMD TESTING  
PARTITION; MAX  
WALLTIME 1 HOUR

Important: to request  
more than 24 hours  
walltime use:  
`--qos=long`

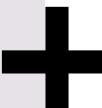


# Access to NVIDIA Job partition

```
#!/bin/bash

#SBATCH --job-name=full_genome_nn_hp_opt_gpu
#SBATCH --output=full_genome_nn_hp_opt_gpu.out
#SBATCH --error=full_genome_nn_hp_opt_gpu.err
#SBATCH --partition=aa100
#SBATCH --nodes=1
#SBATCH --ntasks=45
#SBATCH --gres=gpu:1
#SBATCH --account=amc-general
#SBATCH --time=24:00:00
```

- To access an AMD partition add: **--partition=ami100**
- Here, the number of gpus are allocated with: **--gres=gpu:1**

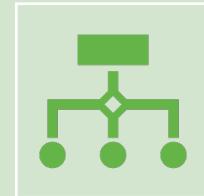


# Debug nodes for NVIDIA A100 and AMD MI100



[NVIDIA] sinteractive

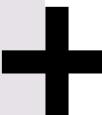
```
--partition=atesting_a100 --qos=testing  
--time=00:05:00 --gres=gpu:1 --ntasks=2
```



[AMD] sinteractive

```
--partition=atesting_mi100 --qos=testing  
--time=00:05:00 --gres=gpu:1 --ntasks=2
```

Users are now limited to up to 2/3  
of the GPU partition (not per node)

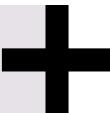


# Package availability for ML (2)

- Cuda versions 11.2, 11.3, 11.4, 11.8 and 12.1.1 on Alpine.
- Only cudnn 8.1, 8.2 and 8.6 on Alpine.

GPU

Version	Python version	Compiler	Build tools	cuDNN	CUDA
tensorflow-2.13.0	3.8-3.11	Clang 16.0.0	Bazel 5.3.0	8.6	11.8
tensorflow-2.12.0	3.8-3.11	GCC 9.3.1	Bazel 5.3.0	8.6	11.8
tensorflow-2.11.0	3.7-3.10	GCC 9.3.1	Bazel 5.3.0	8.1	11.2
tensorflow-2.10.0	3.7-3.10	GCC 9.3.1	Bazel 5.1.1	8.1	11.2
tensorflow-2.9.0	3.7-3.10	GCC 9.3.1	Bazel 5.0.0	8.1	11.2
tensorflow-2.8.0	3.7-3.10	GCC 7.3.1	Bazel 4.2.1	8.1	11.2
tensorflow-2.7.0	3.7-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.6.0	3.6-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.5.0	3.6-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.4.0	3.6-3.8	GCC 7.3.1	Bazel 3.1.0	8.0	11.0
tensorflow-2.3.0	3.5-3.8	GCC 7.3.1	Bazel 3.1.0	7.6	10.1



# Pytorch installation on Alpine [NVIDIA] (1)

---

1. After you log into the Alpine cluster, please load the slurm modules and request allocation so that you can install the packages:

```
module load slurm/alpine  
acompile --ntasks=4
```



2. Load anaconda, create your environment with python 3.10 and activate it.

```
module load anaconda  
conda create -n pytorch_env python=3.10  
conda activate pytorch_env
```



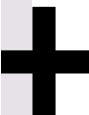
3. Install pytorch, pytorch-cuda. You could also install torchvision and torchaudio if needed for your workflow.

```
conda install pytorch==2.0.0 torchvision==0.15.0 torchaudio==2.0.0 pytorch-cuda=11.8 -c pytorch -c nvidia
```



4. Install cuda-toolkit 11.8.0

```
conda install -c "nvidia/label/cuda-11.8.0" cuda-toolkit
```



# Pytorch installation on Alpine [NVIDIA] (2)

5. Install nvidia-cudnn 8.6.0

```
pip install nvidia-cudnn-cu11==8.6.0.163
```

6. To test that your installation is working you will need to exit "acompile" first and load on the NVIDIA gpu debug partition on Alpine"

```
conda deactivate
exit
sinteractive --partition=atesting_a100 --qos=testing --time=00:05:00 --gres=gpu:1 --ntasks=2
module load anaconda
conda activate pytorch_env
python
```

7. Finally, run the following:

```
>>>import torch
>>>print(torch.cuda.is_available())
True
```

8. Make sure to exit the GPU debug node partition after testing the installation.

```
$ exit
exit
```

# Tensorflow installation on Alpine [NVIDIA] (1)

1. After you log into the Alpine cluster, please load the Slurm modules and request allocation so that you can install the packages:

```
module load slurm/alpine  
acompile --ntasks=4 --time=01:30:00
```

2. Load anaconda, create your environment with python 3.9 and activate it.

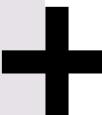
```
module load anaconda  
conda create -n tf_env python=3.9  
conda activate tf_env
```

3. Install cudnn 8.6.0

```
pip install nvidia-cudnn-cu11==8.6.0.163
```

4. Install cuda-toolkit 11.8.0

```
conda install -c "nvidia/label/cuda-11.8.0" cuda-toolkit
```



# Tensorflow installation on Alpine [NVIDIA] (2)

---

5. Install Tensorflow 2.12.0. Also note that tensorflow, cuda and cudnn have to follow a strict versioning: <https://www.tensorflow.org/install/source#gpu>

```
python3 -m pip install tensorflow==2.12.0
```

6. Export the correct paths by following this guide here: <https://www.tensorflow.org/install/pip>

```
mkdir -p $CONDA_PREFIX/etc/conda/activate.d
echo 'CUDNN_PATH=$(dirname $(python -c "import nvidia.cudnn;print(nvidia.cudnn.__file__)"))' >> $CONDA_PREFIX/etc/conda/activate.d/env_vars.sh
echo 'export LD_LIBRARY_PATH=$CONDA_PREFIX/lib/:$CUDNN_PATH/lib:$LD_LIBRARY_PATH' >> $CONDA_PREFIX/etc/conda/activate.d/env_vars.sh
source $CONDA_PREFIX/etc/conda/activate.d/env_vars.sh
export LD_LIBRARY_PATH=$CONDA_PREFIX/lib/python3.9/site-packages/nvidia/cudnn/lib:$LD_LIBRARY_PATH
export LD_LIBRARY_PATH=$CONDA_PREFIX/lib:$LD_LIBRARY_PATH
export PATH=$CONDA_PREFIX/bin:$PATH
export XLA_FLAGS=--xla_gpu_cuda_data_dir=$CONDA_PREFIX
```

7. Install Tensorrt and export PATH:

```
pip install nvidia-tensorrt==8.4.1.5
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$CONDA_PREFIX/lib/python3.9/site-packages/tensorrt
```

# Tensorflow installation on Alpine [NVIDIA] (3)

---

8. We link libnvinfer.so.8 to libnvinfer.so.7 :

```
ln -s $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer.so.8 $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer.so.7  
ln -s $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer_plugin.so.8 $CONDA_PREFIX/lib/python3.9/site-packages/tensorrt/libnvinfer_plugin.so.7
```

9. To test that your installation is working you will need to exit "acompile" first and load on the NVIDIA GPU debug partition on Alpine"

```
conda deactivate  
exit  
sinteractive --partition=atesting_a100 --qos=testing --time=00:05:00 --gres=gpu:1 --ntasks=2  
module load anaconda  
conda activate tf_env  
python3 -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

10. Make sure to exit the GPU debug node partition after testing the installation.

```
$ exit  
exit
```

# Slurm script for AMD GPU partition

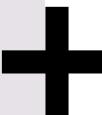
```
#!/bin/bash
#SBATCH --job-name=run_amd_gpu_standard
#SBATCH --output=run_amd_gpu_standard.%j.out
#SBATCH --error=run_amd_gpu_standard.%j.err
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH --account=amc-general
#SBATCH --reservation=amc_workshop
#SBATCH --time=00:10:00

# We make sure that the user can run mambaforge
cp ~/.condarc ~/.mambarc

echo "***** Loading ROCM and pytorch + exporting useful paths *****"
# We load ROCM which is the collection of drivers, libraries, tools and API to run
# on an AMD GPU : https://rocm.docs.amd.com/en/latest/what-is-rocm.html
# For NVIDIA it is NVIDIA CUDA
module load rocm/5.2.3

# We load pytorch 1.13.0 that has been built for the ROCM,
# which is compatible with ROCM5.2.X : https://pytorch.org/get-started/previous-versions/
module load pytorch/1.13.0
```

We give  
job name,  
output  
and error



# Slurm script for AMD GPU partition

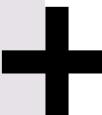
```
#!/bin/bash
#SBATCH --job-name=run_amd_gpu_standard
#SBATCH --output=run_amd_gpu_standard.%j.out
#SBATCH --error=run_amd_gpu_standard.%j.err
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH --account=amc-general
#SBATCH --reservation=amc_workshop
#SBATCH --time=00:10:00

# We make sure that the user can run mambaforge
cp ~/.condarc ~/.mambarc

echo "***** Loading ROCM and pytorch + exporting useful paths *****"
# We load ROCM which is the collection of drivers, libraries, tools and API to run
# on an AMD GPU : https://rocm.docs.amd.com/en/latest/what-is-rocm.html
# For NVIDIA it is NVIDIA CUDA
module load rocm/5.2.3

# We load pytorch 1.13.0 that has been built for the ROCM,
# which is compatible with ROCM5.2.X : https://pytorch.org/get-started/previous-versions/
module load pytorch/1.13.0
```

Name of  
the AMD  
GPU  
partition



# Slurm script for AMD GPU partition

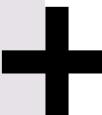
```
#!/bin/bash
#SBATCH --job-name=run_amd_gpu_standard
#SBATCH --output=run_amd_gpu_standard.%j.out
#SBATCH --error=run_amd_gpu_standard.%j.err
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH --account=amc-general
#SBATCH --reservation=amc_workshop
#SBATCH --time=00:10:00

# We make sure that the user can run mambaforge
cp ~/.condarc ~/.mambarc

echo "***** Loading ROCM and pytorch + exporting useful paths *****"
# We load ROCM which is the collection of drivers, libraries, tools and API to run
# on an AMD GPU : https://rocm.docs.amd.com/en/latest/what-is-rocm.html
# For NVIDIA it is NVIDIA CUDA
module load rocm/5.2.3

# We load pytorch 1.13.0 that has been built for the ROCM,
# which is compatible with ROCM5.2.X : https://pytorch.org/get-started/previous-versions/
module load pytorch/1.13.0
```

We ask for  
1 node  
and a few  
cores  
since most  
computation will be  
on the gpu



# Slurm script for AMD GPU partition

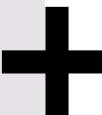
```
#!/bin/bash
#SBATCH --job-name=run_amd_gpu_standard
#SBATCH --output=run_amd_gpu_standard.%j.out
#SBATCH --error=run_amd_gpu_standard.%j.err
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH --account=amc-gener...
#SBATCH --reservation=amc_workshop
#SBATCH --time=00:10:00

# We make sure that the user can run mambaforge
cp ~/.condarc ~/.mambarc

echo "***** Loading ROCM and pytorch + exporting useful paths *****"
# We load ROCM which is the collection of drivers, libraries, tools and API to run
# on an AMD GPU : https://rocm.docs.amd.com/en/latest/what-is-rocm.html
# For NVIDIA it is NVIDIA CUDA
module load rocm/5.2.3

# We load pytorch 1.13.0 that has been built for the ROCM,
# which is compatible with ROCM5.2.X : https://pytorch.org/get-started/previous-versions/
module load pytorch/1.13.0
```

We ask for  
1 GPU



# Slurm script for AMD GPU partition

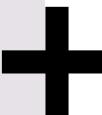
```
#!/bin/bash
#SBATCH --job-name=run_amd_gpu_standard
#SBATCH --output=run_amd_gpu_standard.%j.out
#SBATCH --error=run_amd_gpu_standard.%j.err
#SBATCH --partition=ami100
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=6
#SBATCH --gres=gpu:1
#SBATCH --account=amc-general
#SBATCH --reservation=amc_workshop
#SBATCH --time=00:10:00

# We make sure that the user can run mambaforge
cp ~/.condarc ~/.mambarc

echo "***** Loading ROCM and pytorch + exporting useful paths *****"
# We load ROCM which is the collection of drivers, libraries, tools and API to run
# on an AMD GPU : https://rocm.docs.amd.com/en/latest/what-is-rocm.html
# For NVIDIA it is NVIDIA CUDA
module load rocm/5.2.3

# We load pytorch 1.13.0 that has been built for the ROCM,
# which is compatible with ROCM5.2.X : https://pytorch.org/get-started/previous-versions/
module load pytorch/1.13.0
```

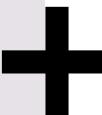
We load  
the  
modules



# Slurm script for AMD GPU partition

```
echo "***** Running the models *****"  
# We want to make sure that we can profile what is going on  
# We want to set AMD_LOG_LEVEL for profiling and debugging  
#CODE 0 means No log  
#CODE 1 means log the error  
#CODE 2 means log warnings  
#CODE 3 means log information  
#CODE 4 means debug mode  
# More information here: https://rocm.docs.amd.com/projects/HIP/en/develop/developer\_guide/logging.html  
  
HIP_VISIBLE_DEVICES=0 ROCR_VISIBLE_DEVICES=0 AMD_LOG_LEVEL=1 HSA_ENABLE_SDMA=0 python tensors.py
```

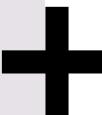
Here is a way we can debug our code or even profile what is going on



# Slurm script for AMD GPU partition

```
echo "***** Running the models *****"  
# We want to make sure that we can profile what is going on  
# We want to set AMD_LOG_LEVEL for profiling and debugging  
#CODE 0 means No log  
#CODE 1 means log the error  
#CODE 2 means log warnings  
#CODE 3 means log information  
#CODE 4 means debug mode  
# More information here: https://rocm.docs.amd.com/projects/HIP/en/develop/developer\_guide/logging.html  
  
HIP_VISIBLE_DEVICES=0 ROCR_VISIBLE_DEVICES=0 AMD_LOG_LEVEL=1 HSA_ENABLE_SDMA=0 python tensors.py
```

AMD\_LOG  
\_LEVEL  
variable is  
very  
useful



# Jupyterhub with NVIDIA GPUs on Ondemand

Home / My Interactive Sessions / Jupyter Session (Custom)

**Interactive Apps**

- Desktops
- Core Desktop (Presets)
- GUIs
- MATLAB (Presets)
- Servers
- Jupyter Session (Custom)**
- Jupyter Session (Presets)
- RStudio Server (Custom)
- RStudio Server (Presets)
- VS Code-Server (Custom)
- VS Code-Server (Presets)

**Jupyter Session (Custom)**

This app will launch a Jupyter Notebook or JupyterLab session. For more information on possible settings for this application, see [Running Custom Interactive applications](#) in our documentation.

Cluster: Alpine

Account: amc-general

Partition: atesting\_a100 

Time: 1

Number of cores: 8

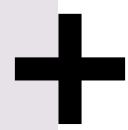
QoS Name: normal

Use JupyterLab instead of Jupyter Notebook?

**Launch**

\* The Jupyter Session (Custom) session data for this session can be accessed under the [data root directory](#).

OnDemand version: 2.0.32

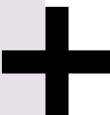


# Containers (1)

---



- Apptainer/Singularity can now be built directly on the cluster
- Can be built either from a definition file or converted from a docker image.
- e.g: `apptainer build -f R_spack_env.sif R_spack_env.def`



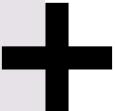
# Containers (2)



- acompile --ntasks=4 -time=04:00:00
- export ALPINE\_SCRATCH=/gpfs/alpine1/scratch/\$USER
- export APPTAINER\_TMPDIR=\$ALPINE\_SCRATCH/singularity/tmp
- export APPTAINER\_CACHEDIR=\$ALPINE\_SCRATCH/singularity/cache
- mkdir -pv \$APPTAINER\_CACHEDIR \$APPTAINER\_TMPD



Exporting the TMP  
related directories



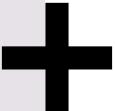
# Containers (2)



- acompile --ntasks=4 -time=04:00:00
- export ALPINE\_SCRATCH=/gpfs/alpine1/scratch/\$USER
- export APPTAINER\_TMPDIR=\$ALPINE\_SCRATCH/singularity/tmp
- export APPTAINER\_CACHEDIR=\$ALPINE\_SCRATCH/singularity/cache
- mkdir -pv \$APPTAINER\_CACHEDIR \$APPTAINER\_TMPD
  
- apptainer build deepvariant.sif  
**docker://google/deepvariant:"\${BIN\_VERSION}"**



Always add docker://  
when building from a  
docker container



# Questions?



Let's a go!!

